



Unlock more value from your data with Microsoft Fabric

Discovery Days – April 17th 2024

Why is your Data Strategy & Foundation a necessity in the Era of AI?

Unstructured data will double in 2024¹

Representing only ~1/3 of managed data today, unstructured data is an untapped potential for businesses. Within the Era of AI, more insights & intelligence will arise (customer feedback, social posts, conversational interactions,...), coming from several angles within your organization^{1&2}.

Language models can surface game-changing trends from unstructured sources¹.



Data Strategy



Data Platform



**Data Governance
(incl. Data Quality & Relevance)**



Data Culture



Sources:

1. Predictions 2024: Data Analytics Set The Stage For Generative AI (forrester.com)

2. Data Analytics | Microsoft Fabric

Learn & Connect

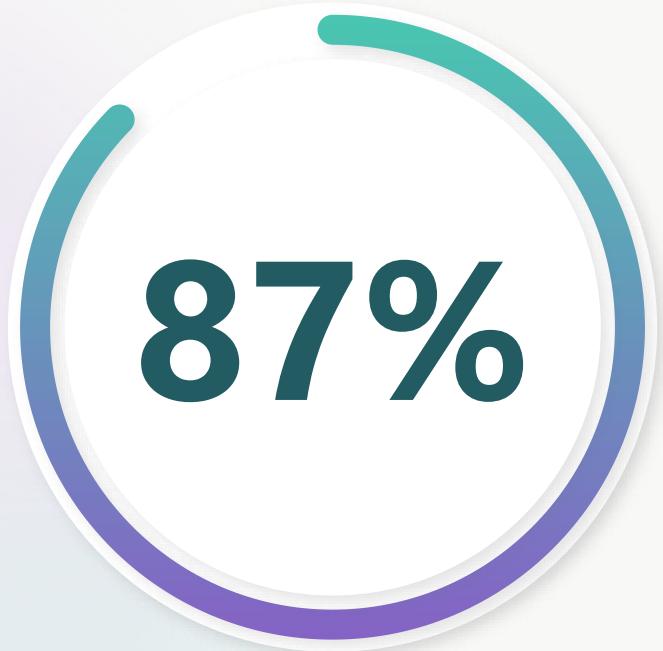
Morning	Afternoon	
09:15 - 10:50	14:15 - 15:50	MS Fabric Intro and Demo
10:50 - 11:00	15:50 - 16:00	Break and Networking
11:00 - 11:50	16:00 - 16:50	Deep Dive
12:00 - 12:30	17:00 - 17:30	Morning: Lunch and Networking Afternoon: Drinks and Networking



Microsoft

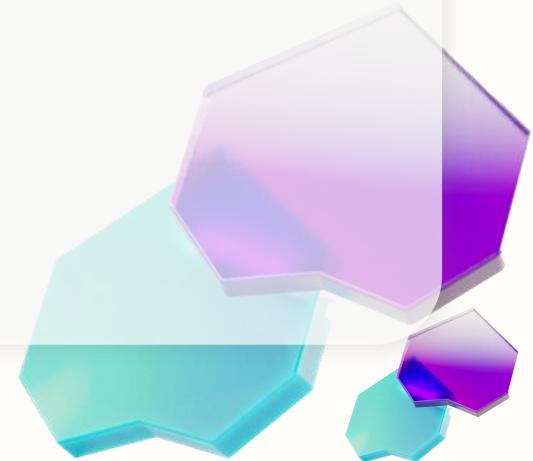
Microsoft Fabric

Microsoft Fabric Overview



of organizations believe
AI will give them a
competitive edge

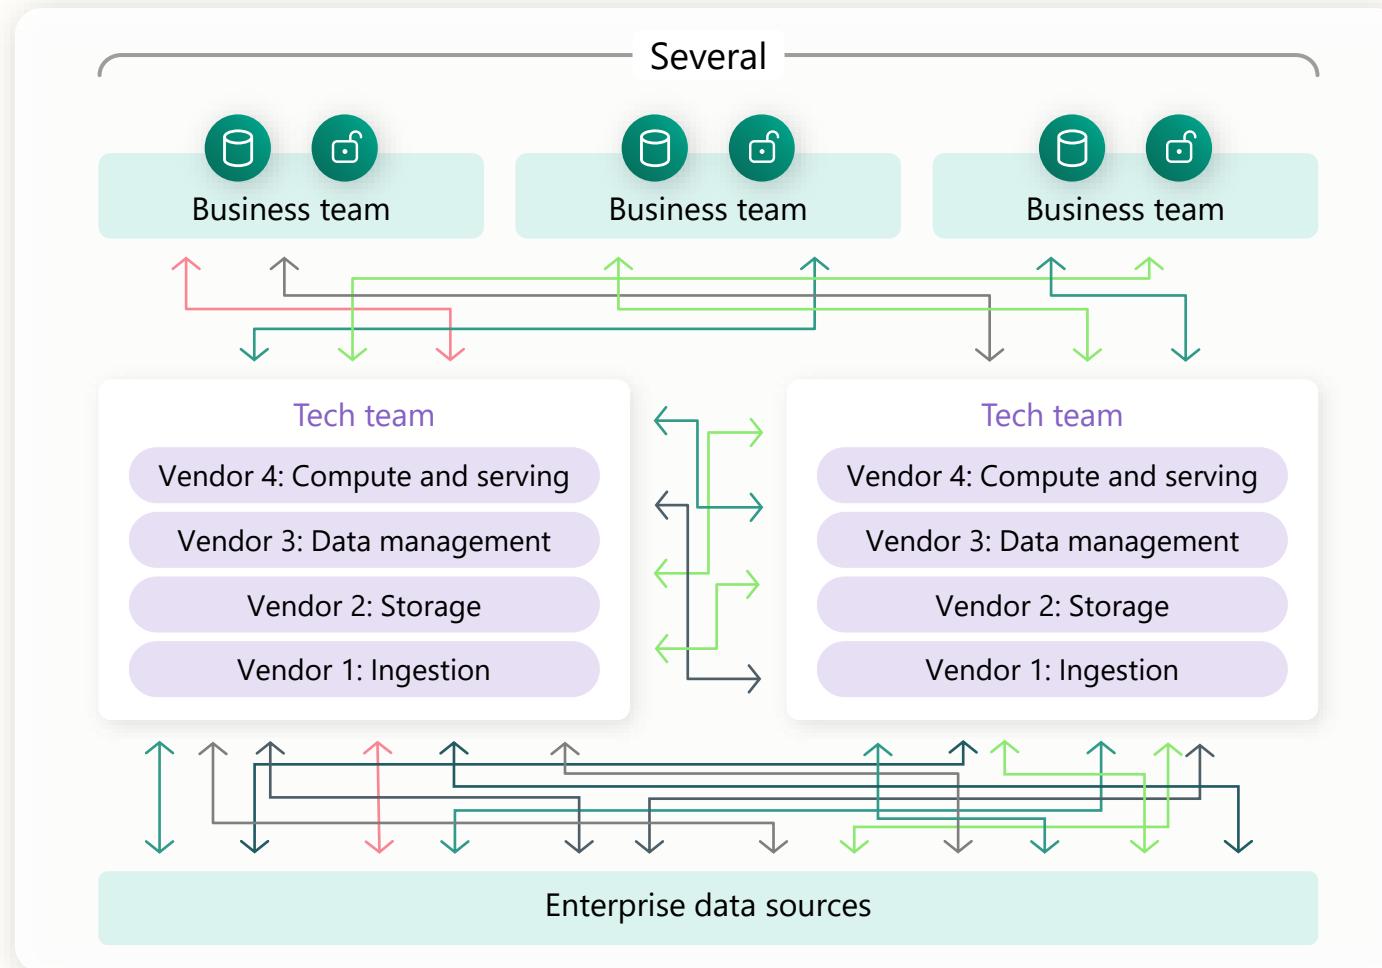
Source: [MIT Sloan Management Review](#)



CIOs are accelerating their efforts to bring AI to their data estate

The starting line

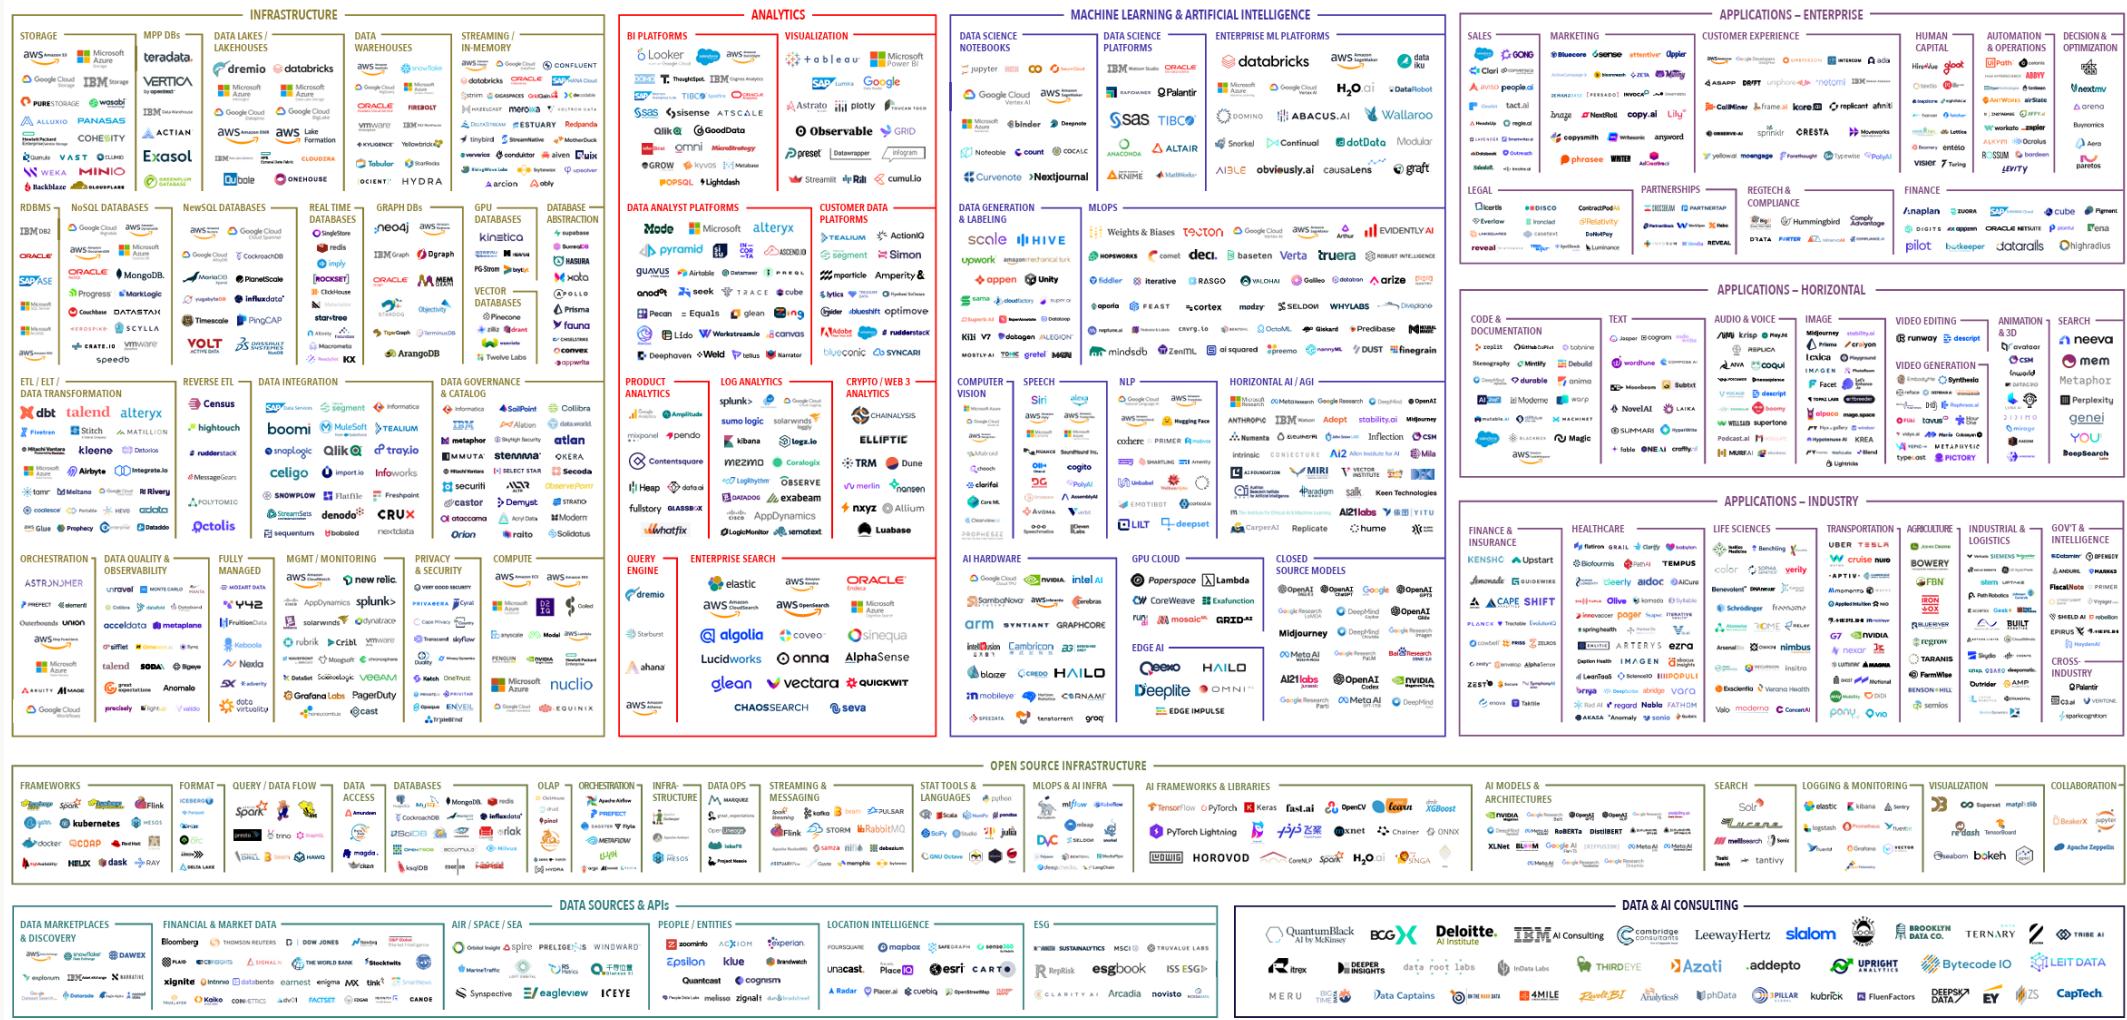
A complex, organically evolved data estate



- 1 Data copies and infrastructure inefficiencies
- 2 Limited interoperability between vendor services
- 3 Data exposure risks



Customers enhancing their data estate face immense complexity



Version 1.0 - Feb 2023

© Matt Turck (@mattturck), Kevin Zhang (@kevinzhang) & FirstMark (@firstmarkcap)

Blog post: mattturck.com/MAD2023

Interactive version: MAD.firstmarkcap.com

Comments? Email MAD2023@firstmarkcap.com

FIRSTMARK
EARLY STAGE VENTURE CAPITAL



Microsoft Fabric

The data platform for the era of AI

From

To

Multiple analytics services



Unified stack

Disconnected data sources



All the data in one place

Isolated application

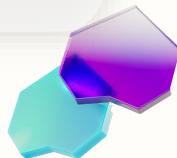


Entire estate

Gen AI bolt on



Gen AI built in





Microsoft Fabric

Intelligent data foundation



Data
Factory



Data
Engineering



Data
Warehouse



Data
Science



Real Time
Analytics



Power BI



Data Activator



Powered by AI with Copilot in Microsoft Fabric



Single data foundation
OneLake



UNIFIED

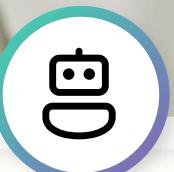
SaaS product experience

Security and governance

Compute and storage

Business model

Microsoft Fabric Capabilities



Empower every business user

Empower everyone to uncover insights accessible data, easy-to-use tools, and visuals embedded apps they use everyday

Unify your analytics on a complete, governed platform

Reduce the cost and effort of integration with a unified, secure, and governed platform

Establish a trusted data foundation

Connect to any data from a single, multi-cloud data lake and use a single copy of data across engines



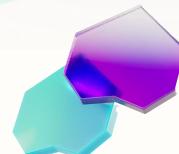
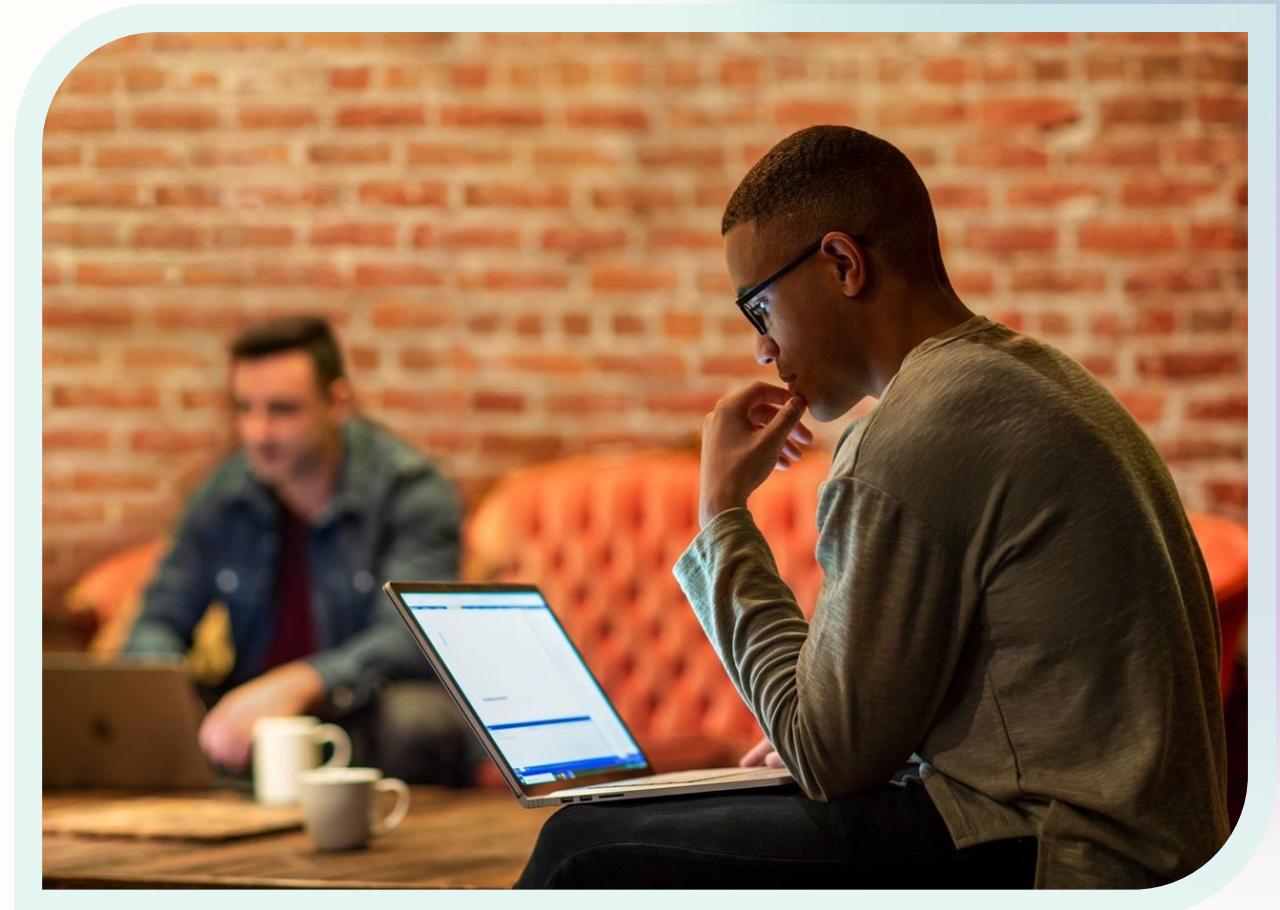
Fuel your AI innovation

Adopt a data platform infused with AI at every layer to help you get more done, faster

Unify your analytics on a complete platform

Give your data teams all the tools they need in a unified, governed, and secure experience that reduces the cost and effort of integration

- Empower data engineers, data scientists, analysts, and business users with role-specific tools in a SaaS platform built for collaboration
- Gain industry-leading, end-to-end security, governance, compliance, and visibility across the unified platform
- Simplify billing and reduce costs with a single pool of capacity and storage that can be used for every workload

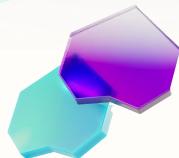


Establish a trusted data foundation



Integrate data from anywhere into a single, multi-cloud data lake for the entire organization, and work from the same copy of data across analytics engines

- Easily connect to data across clouds using “Shortcuts” to virtualize data in OneLake without having to move or duplicate the data
- Create, integrate, manage, and operate data lakes standardized on Delta Parquet format; the same open data format as Azure Databricks
- Intuitively organize your data in Microsoft Fabric’s data lake—OneLake—for central data discovery, sharing, governance, and compliance

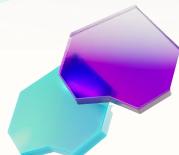


Empower every business user



Empower everyone to uncover insights with the data they need, easy-to-use tools, and visuals embedded in the Microsoft 365 apps they use everyday

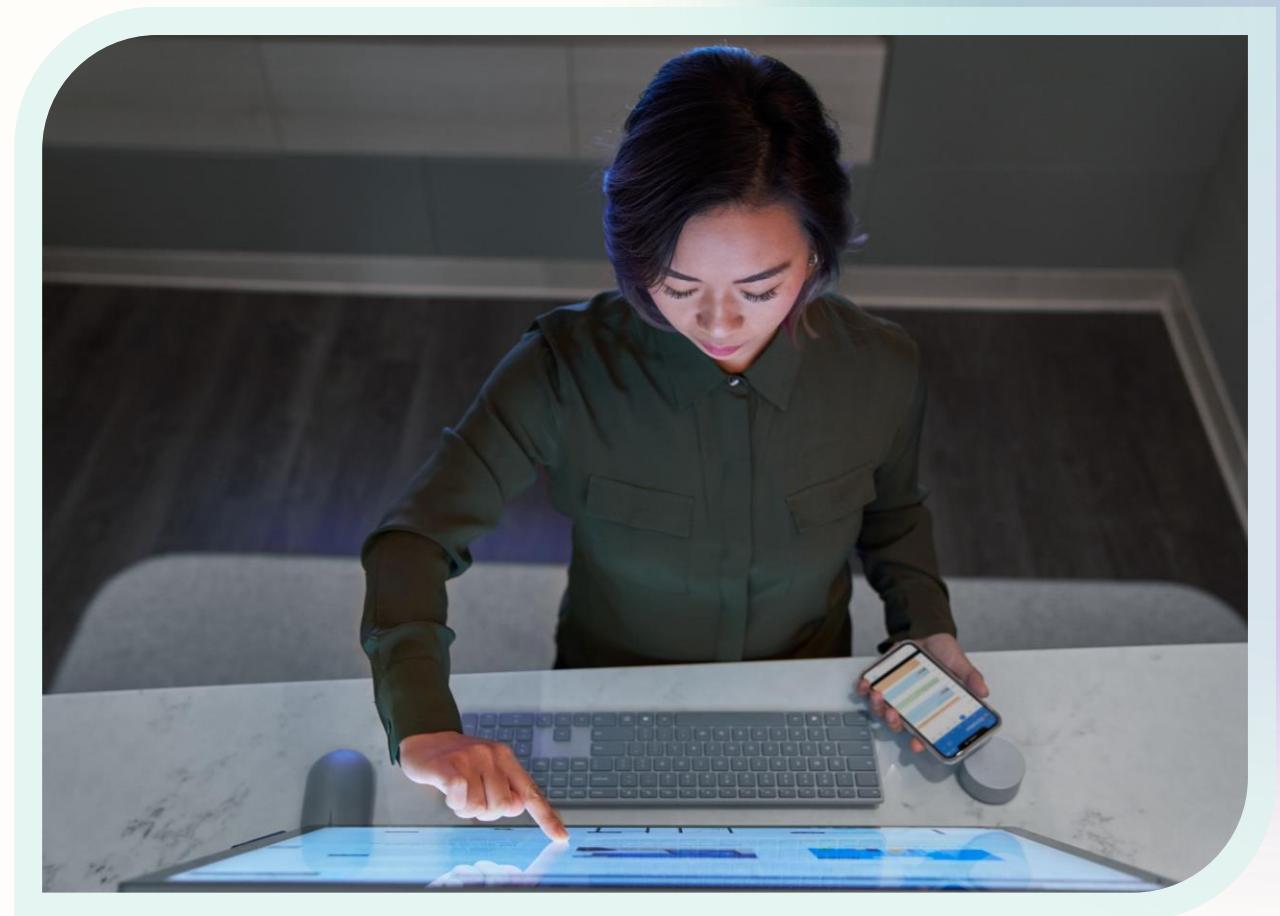
- Quickly go from data in a Lakehouse to insights in the hands of your business users
- Save time for analysts and provide up-to-date insights with Direct Lake mode, a blazing fast, real-time connection to your data in OneLake
- Foster a data-driven culture by seamlessly and securely embedded insights into Teams, Excel, PowerPoint, Outlook, and more with native integration



Fuel your AI innovation

Adopt a data platform that's infused with AI at every layer to help you get more done, faster

- Use conversational language with Copilot in Fabric to create dataflows and pipelines, write SQL statements, or even build machine learning models
- Simply describe what you need—including reports, summaries, and calculations—or ask a question, and Copilot in Power BI does the rest
- Use LangChain and Semantic Kernel to develop and scale custom generative AI models—right from your Fabric Notebook
- Deliver custom generative AI experiences like tailored Q&A on your data*



*Features coming soon



Build your own powerful Gen AI experiences



Microsoft Fabric

Your unified
solution for
analytics

Enhance data insights

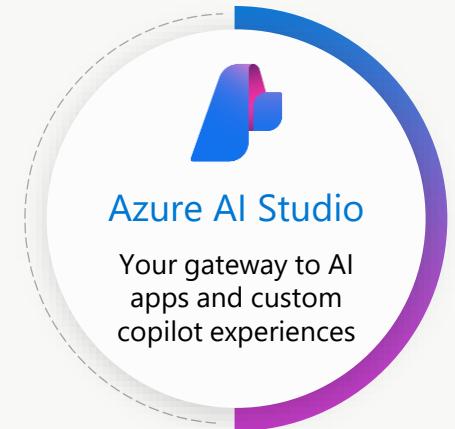
Build enterprise chat applications to uncover insights using natural language from structured, unstructured, and real-time data

Analyze customer interactions

Build speech analytics applications to enhance customer service, tailor support responses, and make data-driven decisions



OneLake serves as the connective tissue to build generative AI apps powered by your data

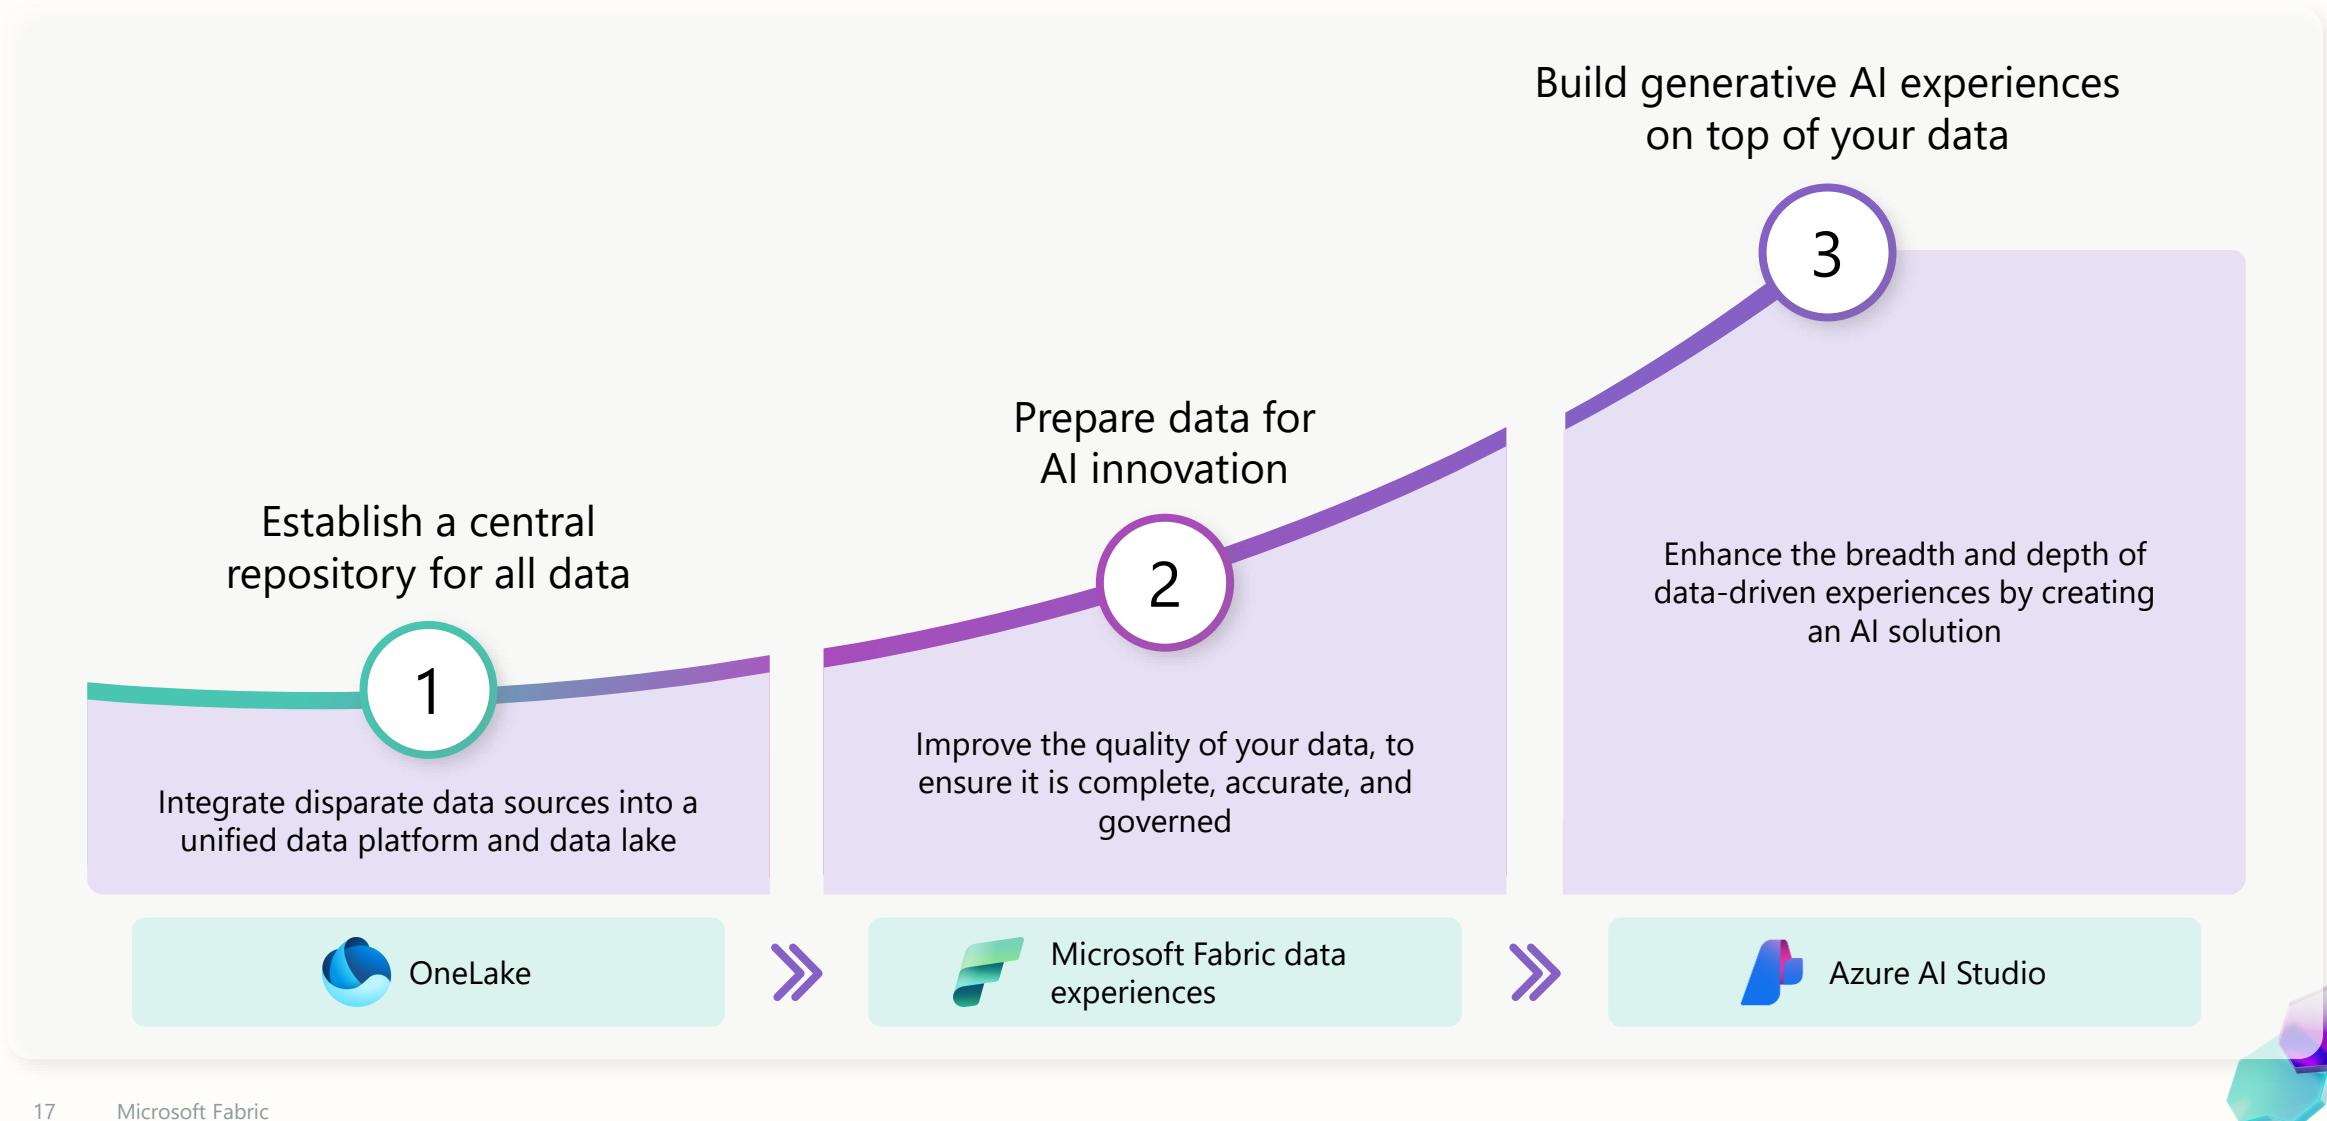


Azure AI Studio

Your gateway to AI
apps and custom
copilot experiences



Fuel your journey to AI innovation by getting your data AI ready



Seven key workloads for end-to-end analytics

Workloads are designed to target specific personas and tasks, yet work together seamlessly in a unified platform via OneLake to enable creators to collaboratively do their best work



Data Factory: Unify your data estate with a data integration experience and 300+ data transformations to easily solve the most complex ETL scenarios



Data Engineering: Enable data engineers to design, build, and maintain infrastructures at scale using World-class Spark platform with great authoring experiences to



Data Warehouse: Provide industry-leading SQL performance and scale, fully separating compute from storage for independently scaling and natively storing data in open Parquet/Delta Lake



Data Science: Empower data scientists and analysts to quickly build, deploy, and operationalize sophisticated AI directly within Fabric



Real Time Analytics: Create actionable insights with best-in-class engine for exploring data and performing real-time analysis with high performance and low latency

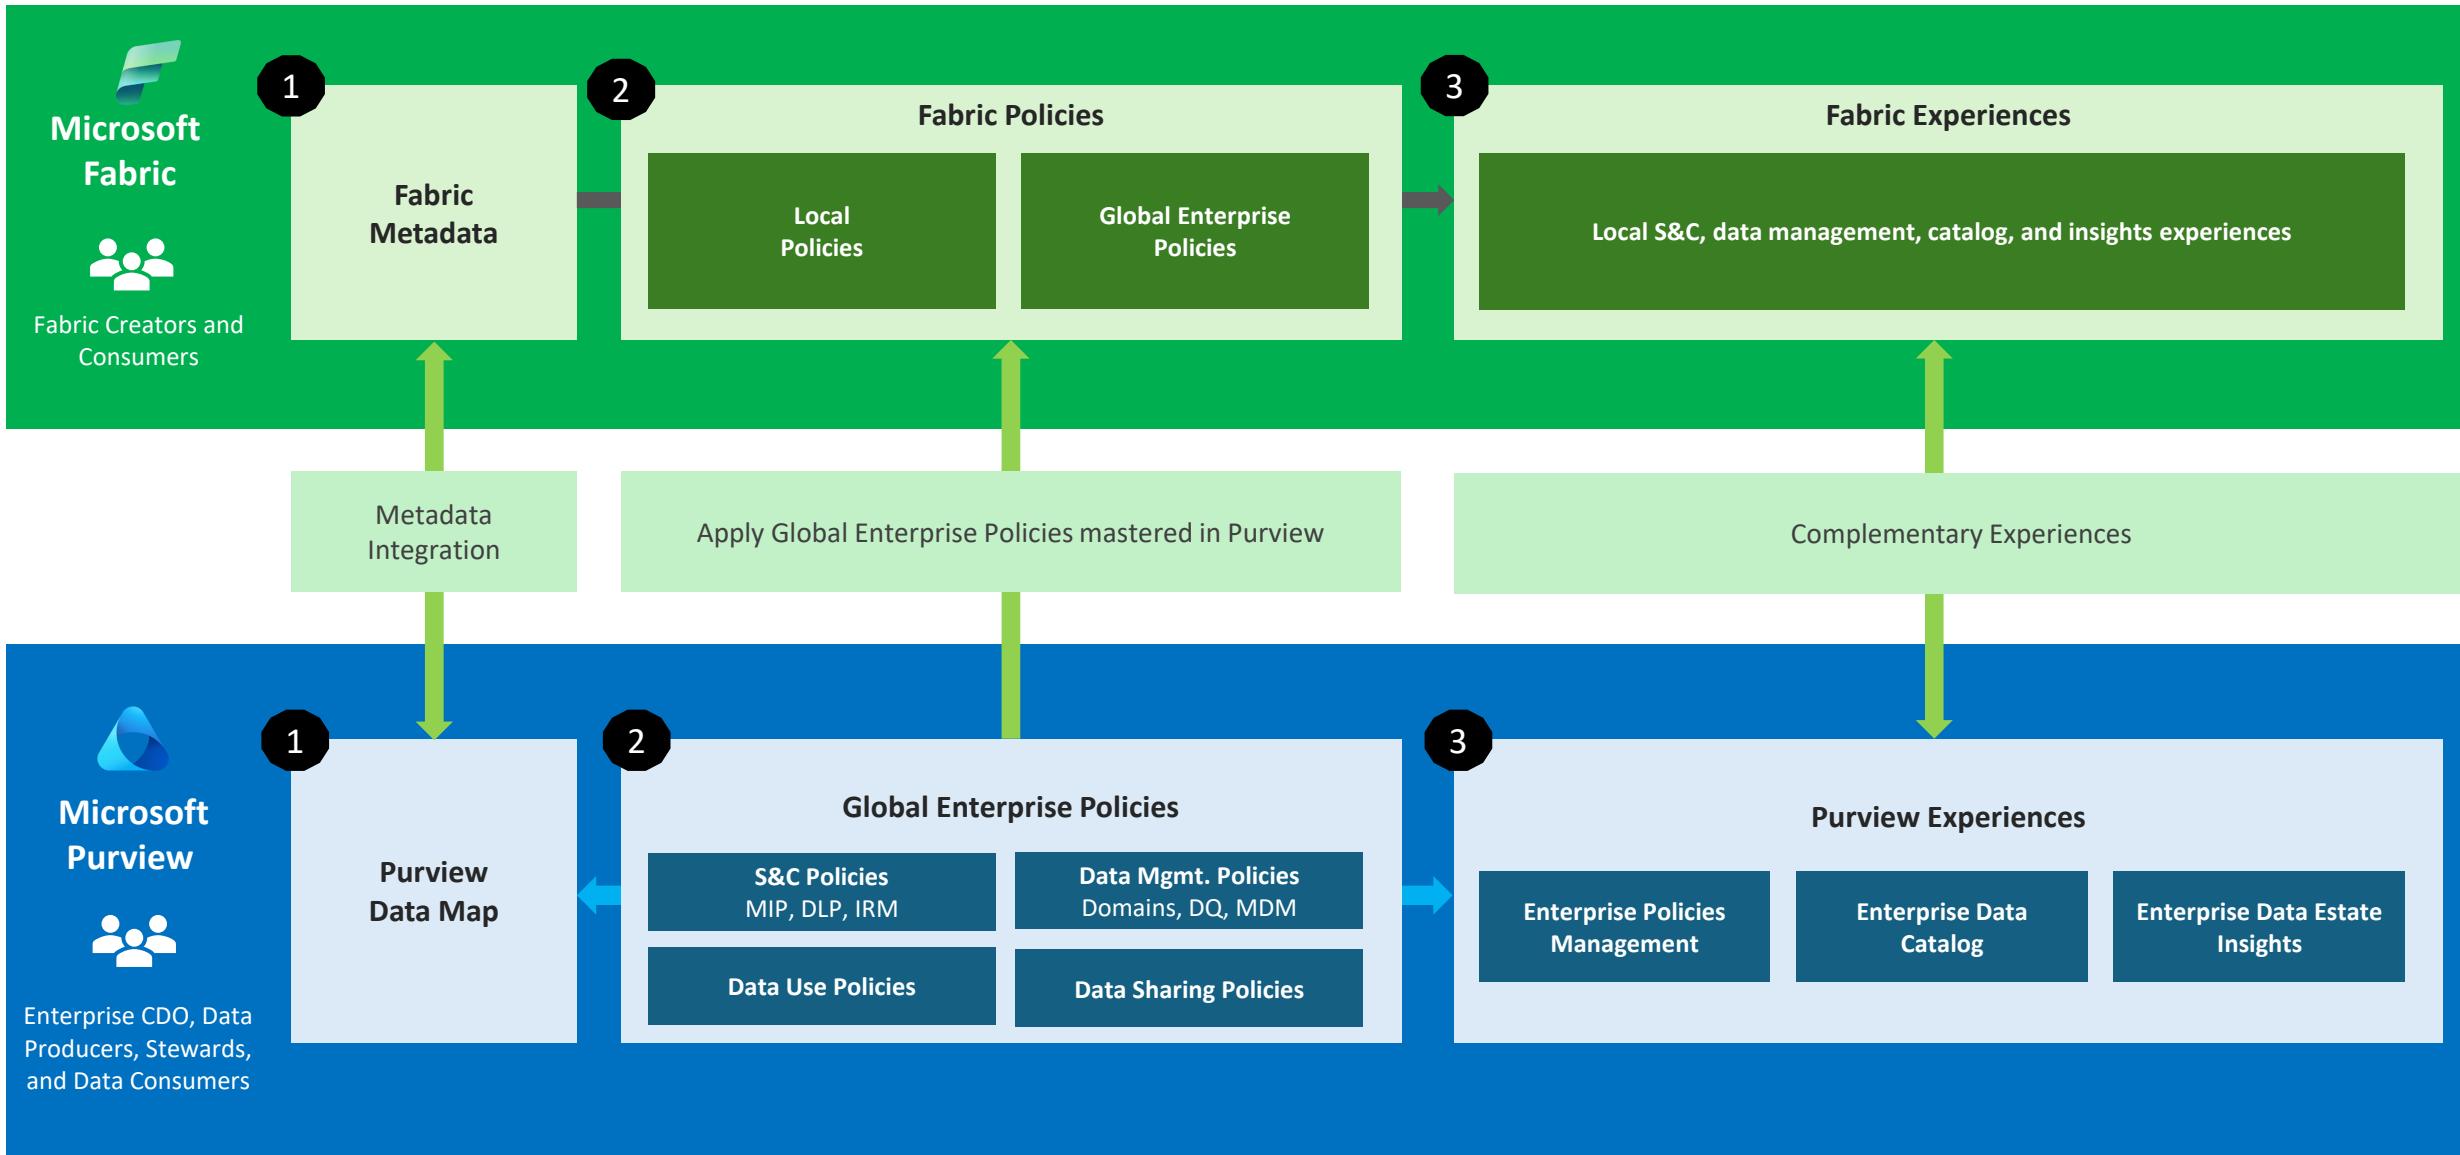


Power BI: Make better, data-driven decisions with the world's leading business intelligence platform that turns unrelated sources of data into coherent, interactive insights

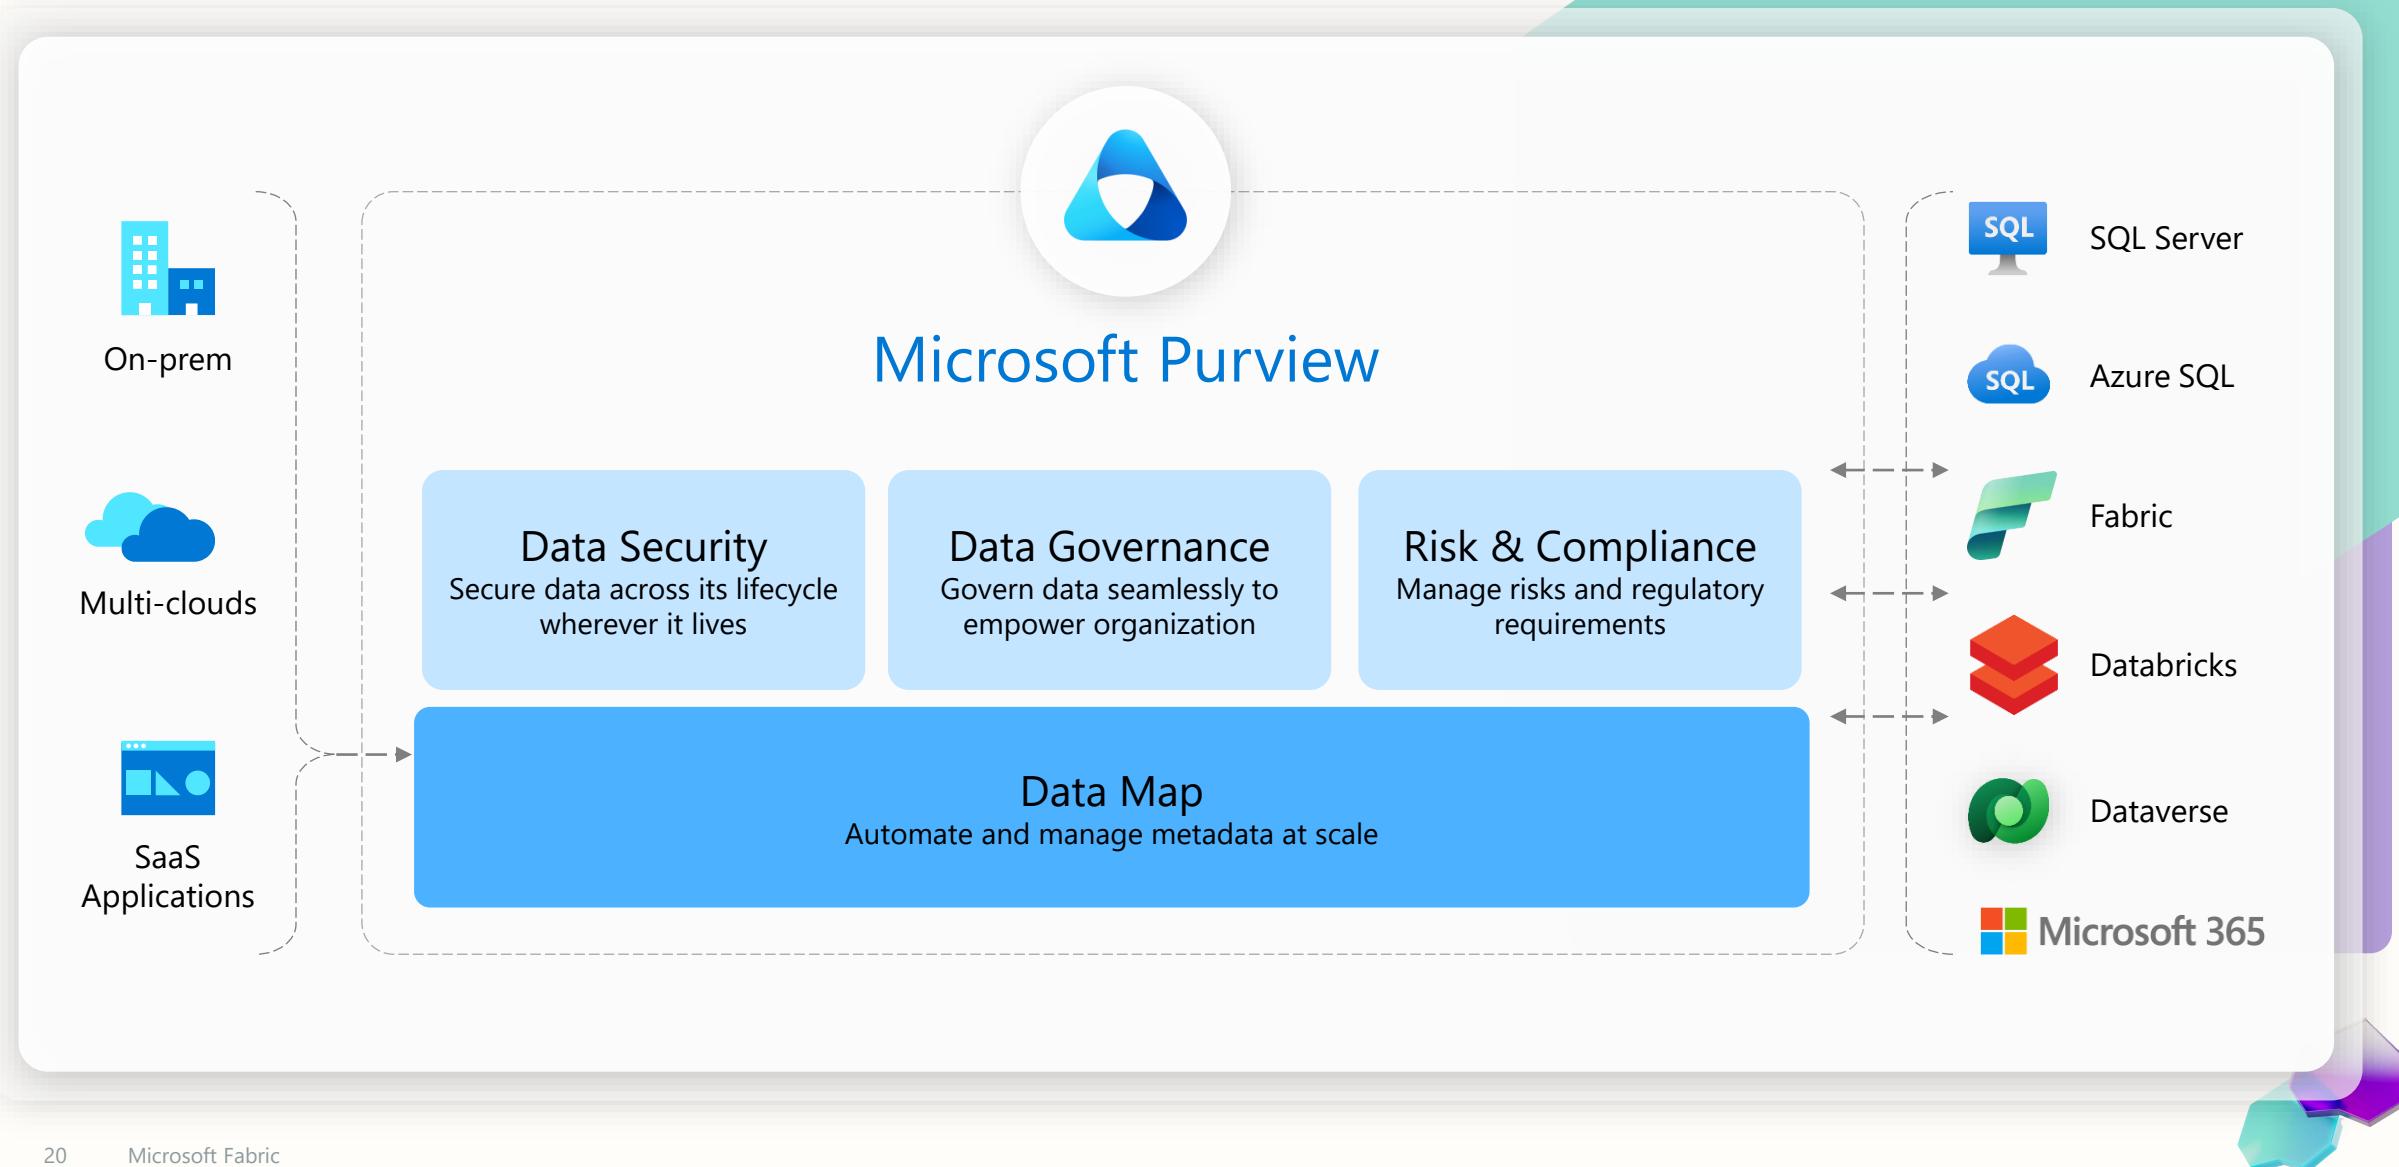


Data Activator: Drive actions automatically on your data without writing code using a system of detection that monitors analytics

Microsoft Fabric and Microsoft Purview



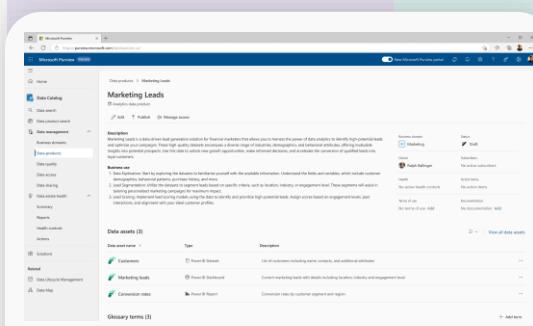
Govern and secure your data estate



Govern entire data estate with Purview Data Governance

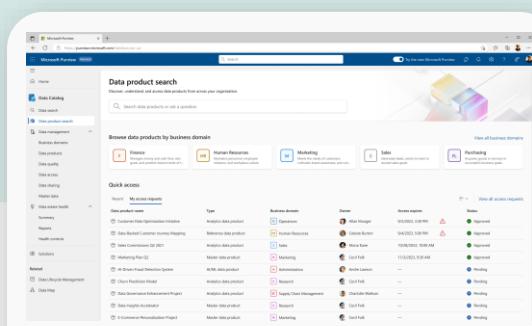
In public preview

- Europe north today
- Europe west planned next month



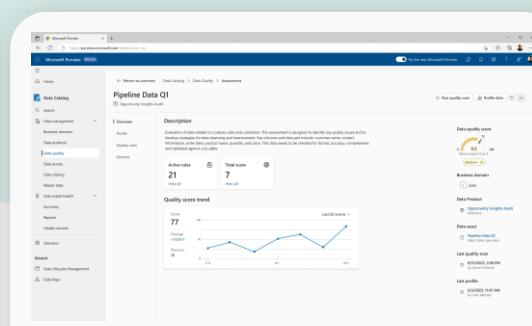
Data Catalog

Curate data across entire data estate (including Fabric data) into **business domains** and **data products** for business-friendly data discovery and understanding



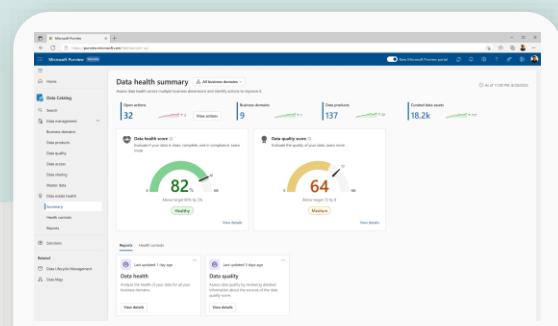
Data Discovery*

Discover, understand, and get access to data products containing Fabric data in the Purview data catalog



Data Management*

Manage access and data quality on data products containing Fabric data; **Master data management** through partners



Data Estate Health*

Monitor and improve data estate health for data products across your organization including Fabric data



Additional Microsoft Purview purchase required
Currently onboarding preview by invitation



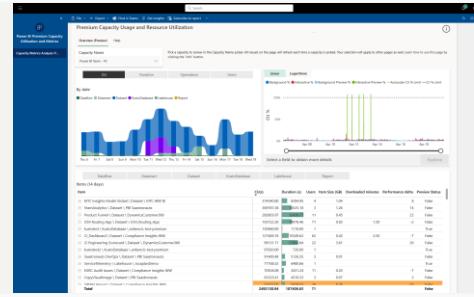
*Fabric integration coming soon



Create a logical data mesh to simplify management and increase usage

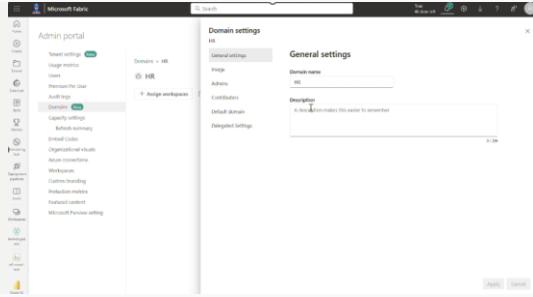
In public preview

- Europe north today
- Europe west planned next month



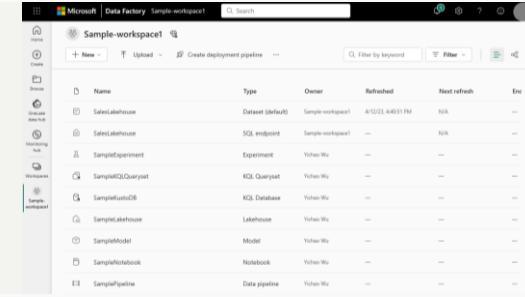
Capacities

Easily manage multiple capacities in a single tenant to **better manage cost**



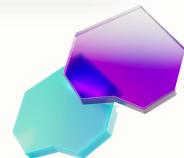
Domains and subdomains

Group data and workspaces together into a domain and sub domains to establish a **data mesh architecture**

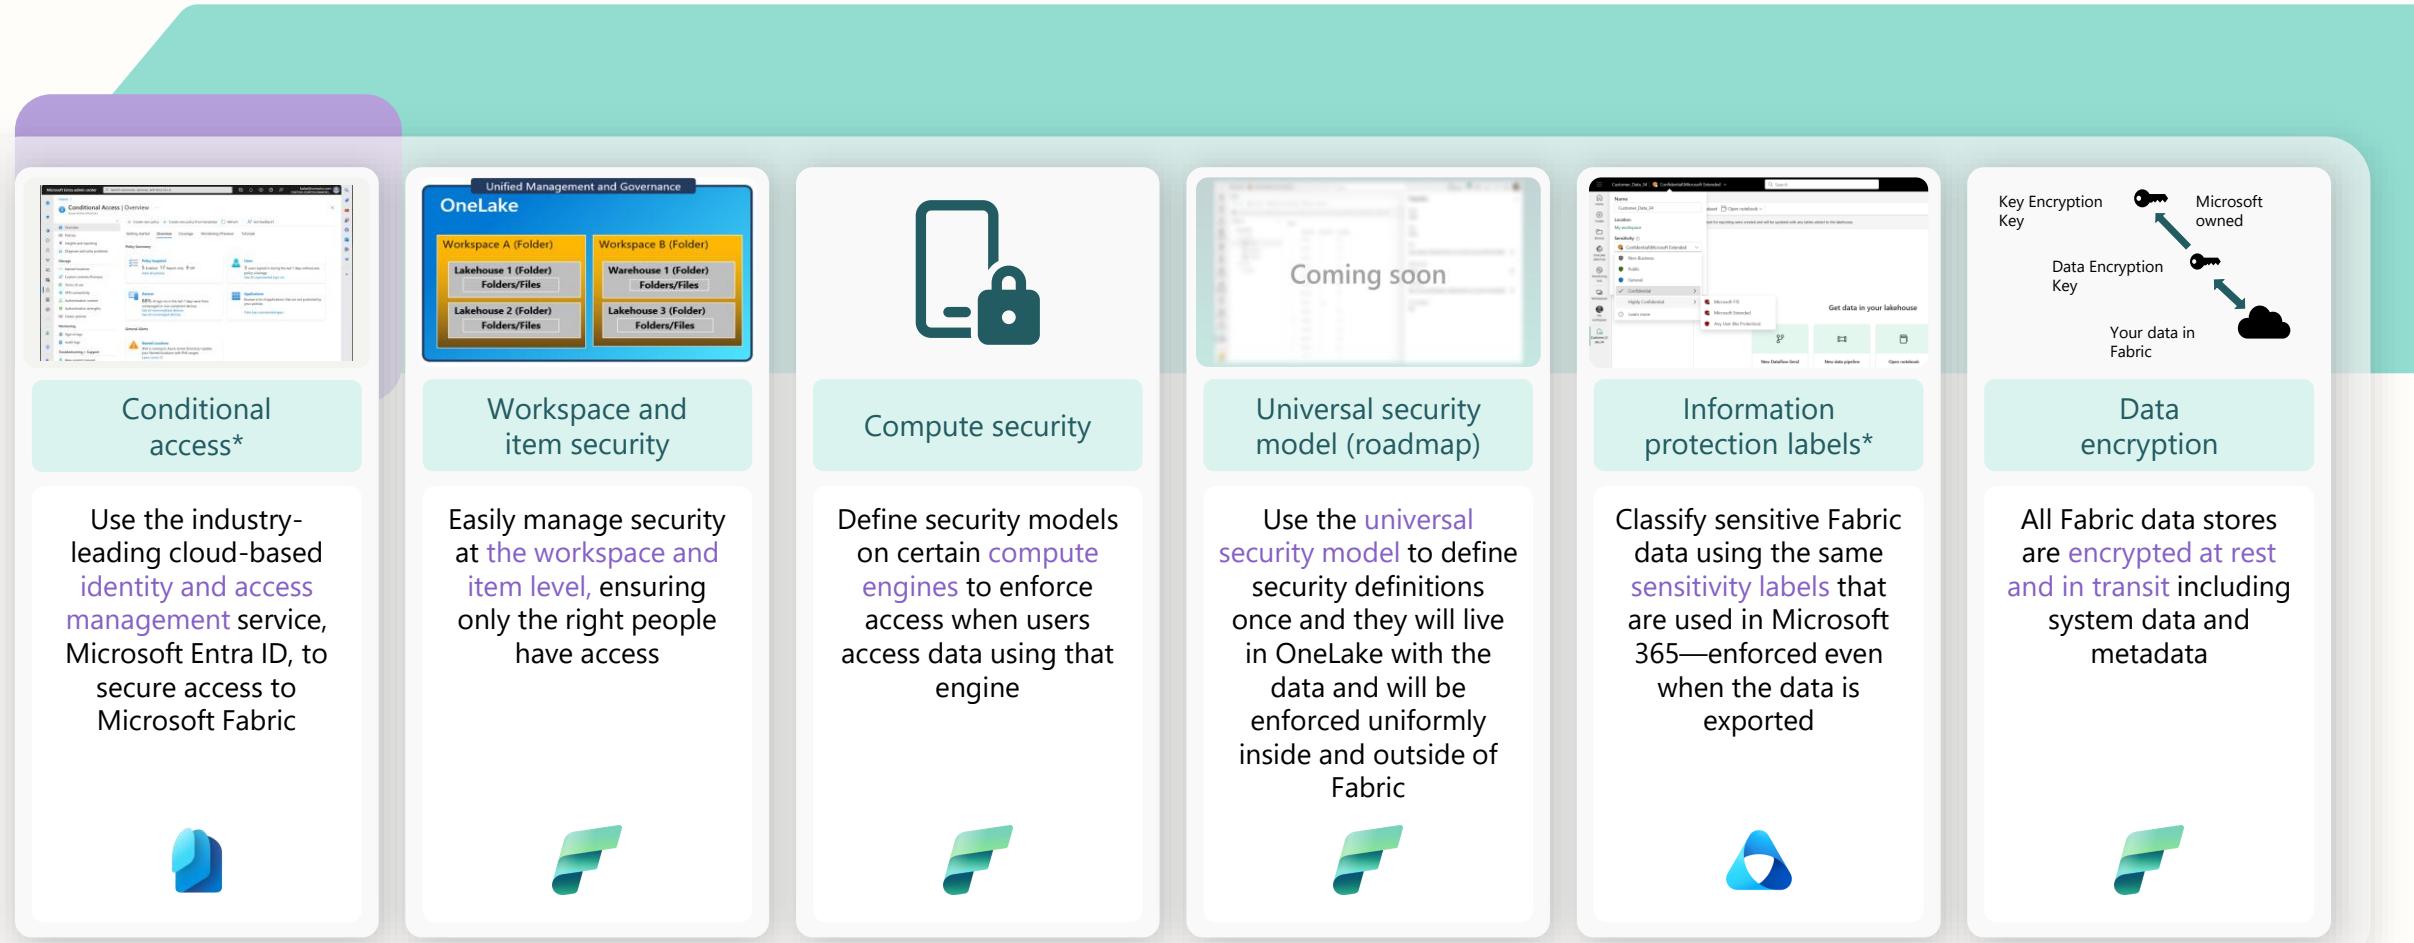


Workspaces

Group items like lakehouses, warehouses, and reports into workspaces to organize and more easily manage items together



Ensure only the right people have access to the right data



*Additional purchase required

Remain in compliance with even the strictest requirements

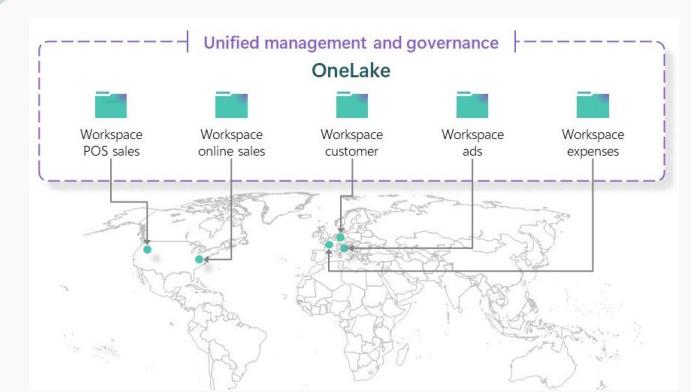


Certifications and standards

Microsoft Fabric, as a core online service, supports a wide range of compliance standards including GDPR, EUDB, HIPAA, ISO certifications and more



*Additional Microsoft Purview purchase required



Data residency

Control where your data is stored and meet data residency requirements with over 54 data centers world-wide



A screenshot of the Microsoft Purview Audit interface. It shows a search bar at the top with filters for 'New Search', 'Classic Search', and 'Active unified searches'. Below the search bar is a table with columns: 'Search name', 'Job status', 'Progress...', 'Search #', 'Total results', and 'Creation time...'. There are several rows of audit logs, each with a checkbox, a timestamp, a status, and a progress percentage. At the bottom right, it says '31 items'.

Microsoft Purview Audit*

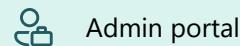
Log user activities from Microsoft Fabric to support security, forensic, and internal investigations



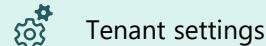
Gain full visibility and control with industry-leading features

Governance and security features built into Fabric

Manage your data estate



Admin portal



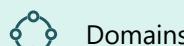
Tenant settings



Workspace settings



Capacities

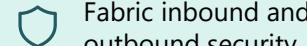


Domains

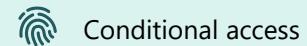


Workspaces

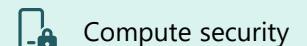
Secure, protect, and comply



Fabric inbound and outbound security



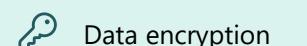
Conditional access



Compute security



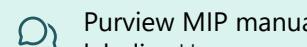
Workspace & item security



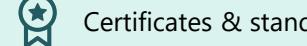
Data encryption



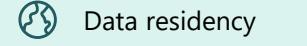
Universal security model*



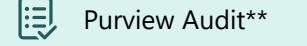
Purview MIP manual labeling**



Certificates & standards

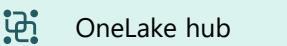


Data residency



Purview Audit**

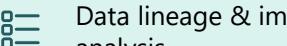
Encourage data discovery and use



OneLake hub



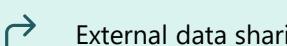
Endorsements



Data lineage & impact analysis



Metadata scanning



External data sharing*



Metadata curation



Discover & Reuse

Uncover insights and act



Admin monitoring



Capacity metrics



Monitoring hub



Purview Hub

Fabric metadata



Microsoft Purview



Comprehensive data security¹



Federated data governance¹



Risk & compliance posture¹

Fabric auto-attaches to Microsoft Purview

*Features coming soon



**Additional Microsoft Purview purchase required



Get started with Fabric and Purview for free



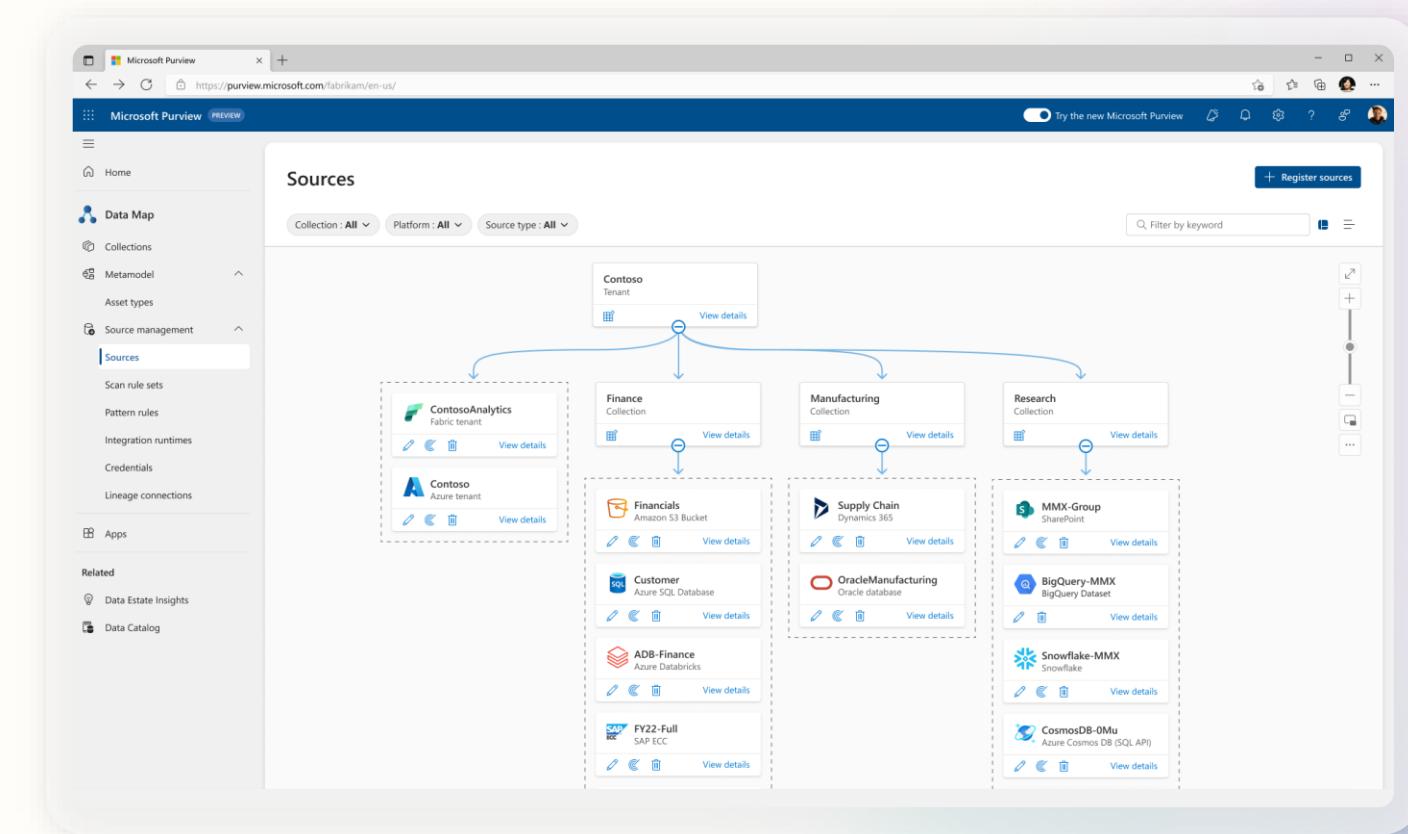
Try out Purview easily with **no set-up required** with Purview Data Map automatically provisioned and attached to every Fabric instance by default



Scan Fabric tenants at no extra cost and discover data assets across your Fabric landscape with data discovery, sensitive data classification, and end-to-end data lineage.



Search and explore Fabric data assets in the Purview Data Catalog including the technical properties, sensitive data classification, and data lineage



References



HITACHI
Inspire the Next



Persistent

TOPSOE

ABInBev

Lenovo

Icertis

Mahindra

PIAGGIO

Vancouver Coastal Health

A
THE ADECCO GROUP



INTESA SANPAOLO

MCKESSON

REPSOL

ZURICH

ambev

ETIHAD

JTI

NEW LOOK

SAMSUNG

CANADIAN TIRE

blue California

FERRING
PHARMACEUTICALS



NOKIA

Shell

T-Mobile

Bovemij



L'ORÉAL

OPTUS

Telefónica

Dow

Culligan

GRUPPO HERA

andel

Manulife

tomtom

Kroger

Lufthansa

Panasonic

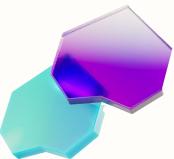
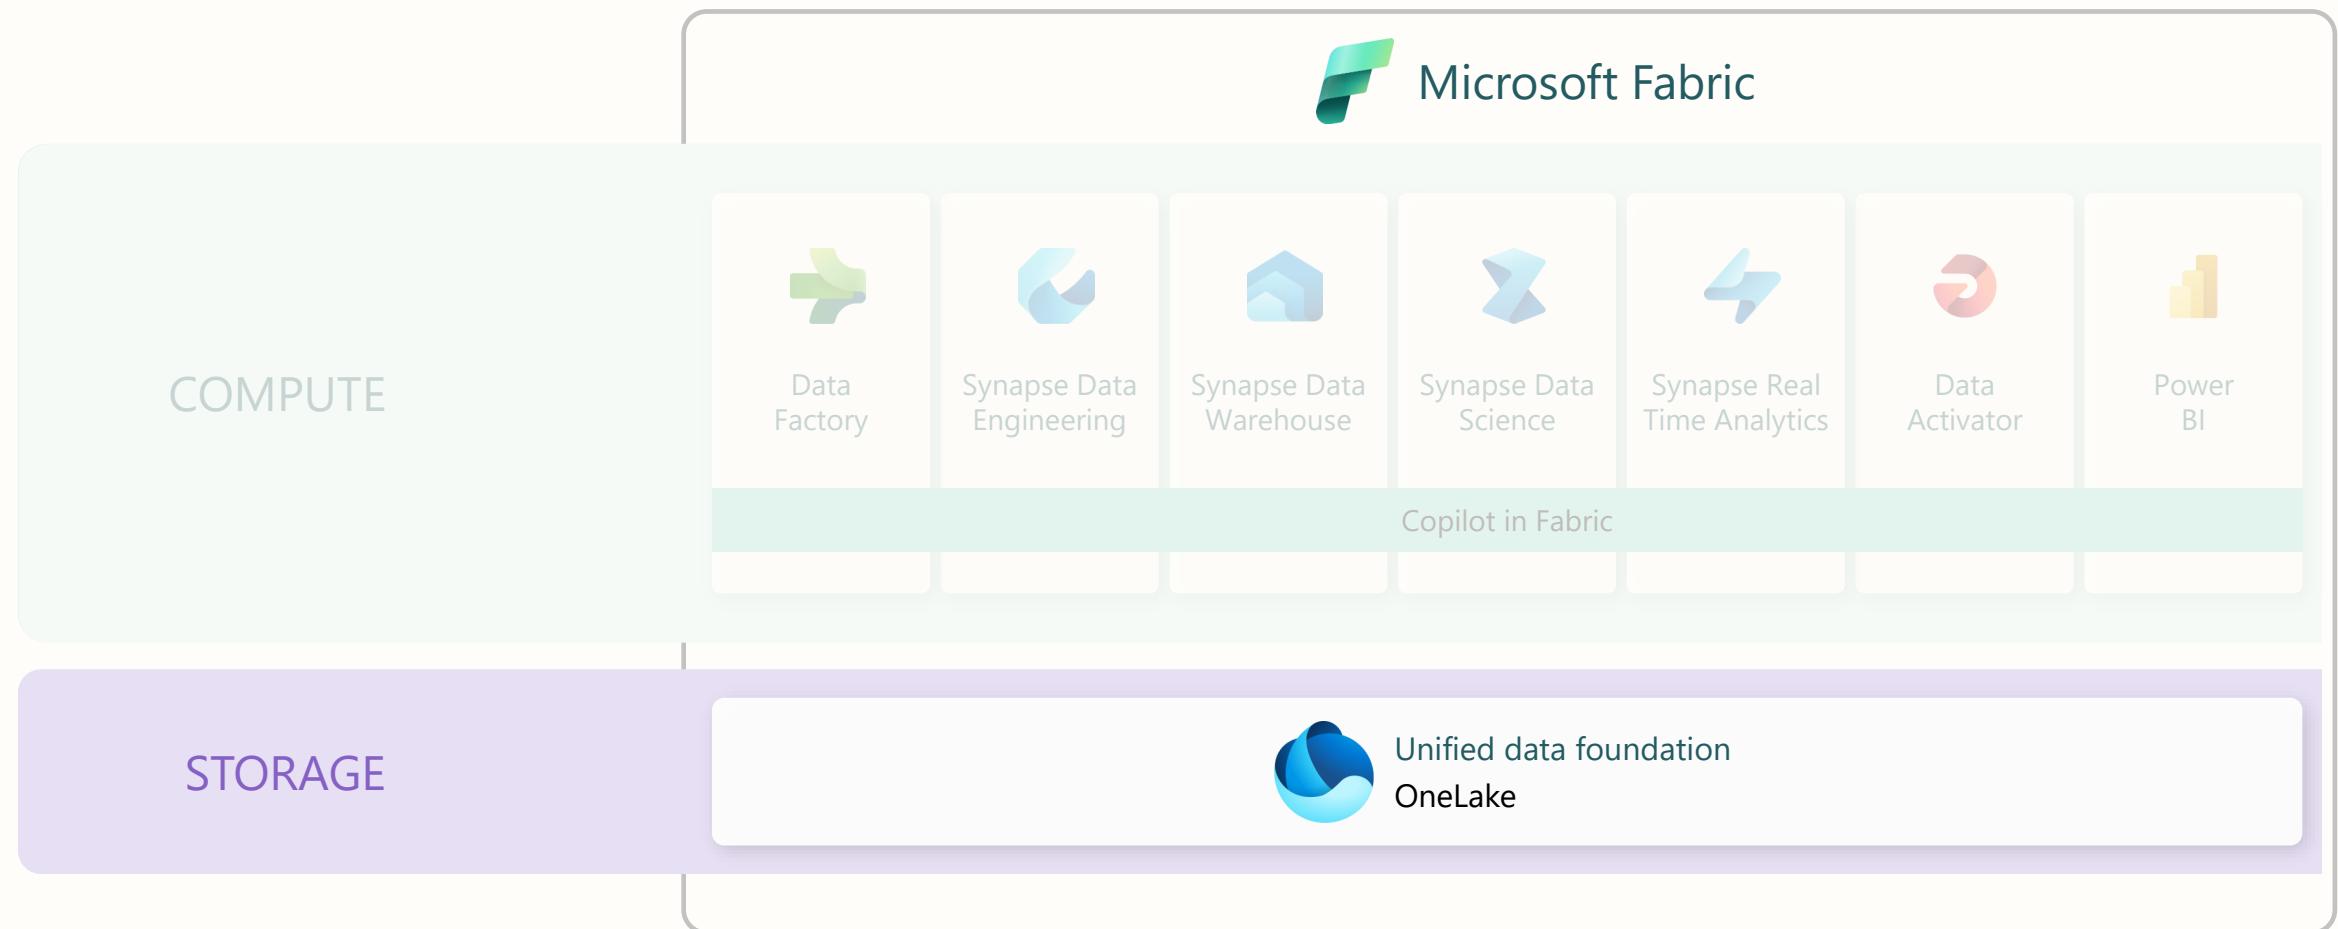


Microsoft

Microsoft Fabric

Microsoft Fabric Licensing

Microsoft Fabric business model overview



OneLake pricing

OneLake is a data lake built into Microsoft Fabric and provides a single place to store all organizational data. Data storage is charged at a rate of € per GB per month and priced uniquely across regions.



Data Storage

Type	Pay-as-you-go price at West Europe
OneLake storage	€0.0217 per GB/month
OneLake BCDR storage	€0.0415 per GB/month
OneLake cache	€0.2342 per GB/month

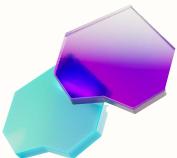
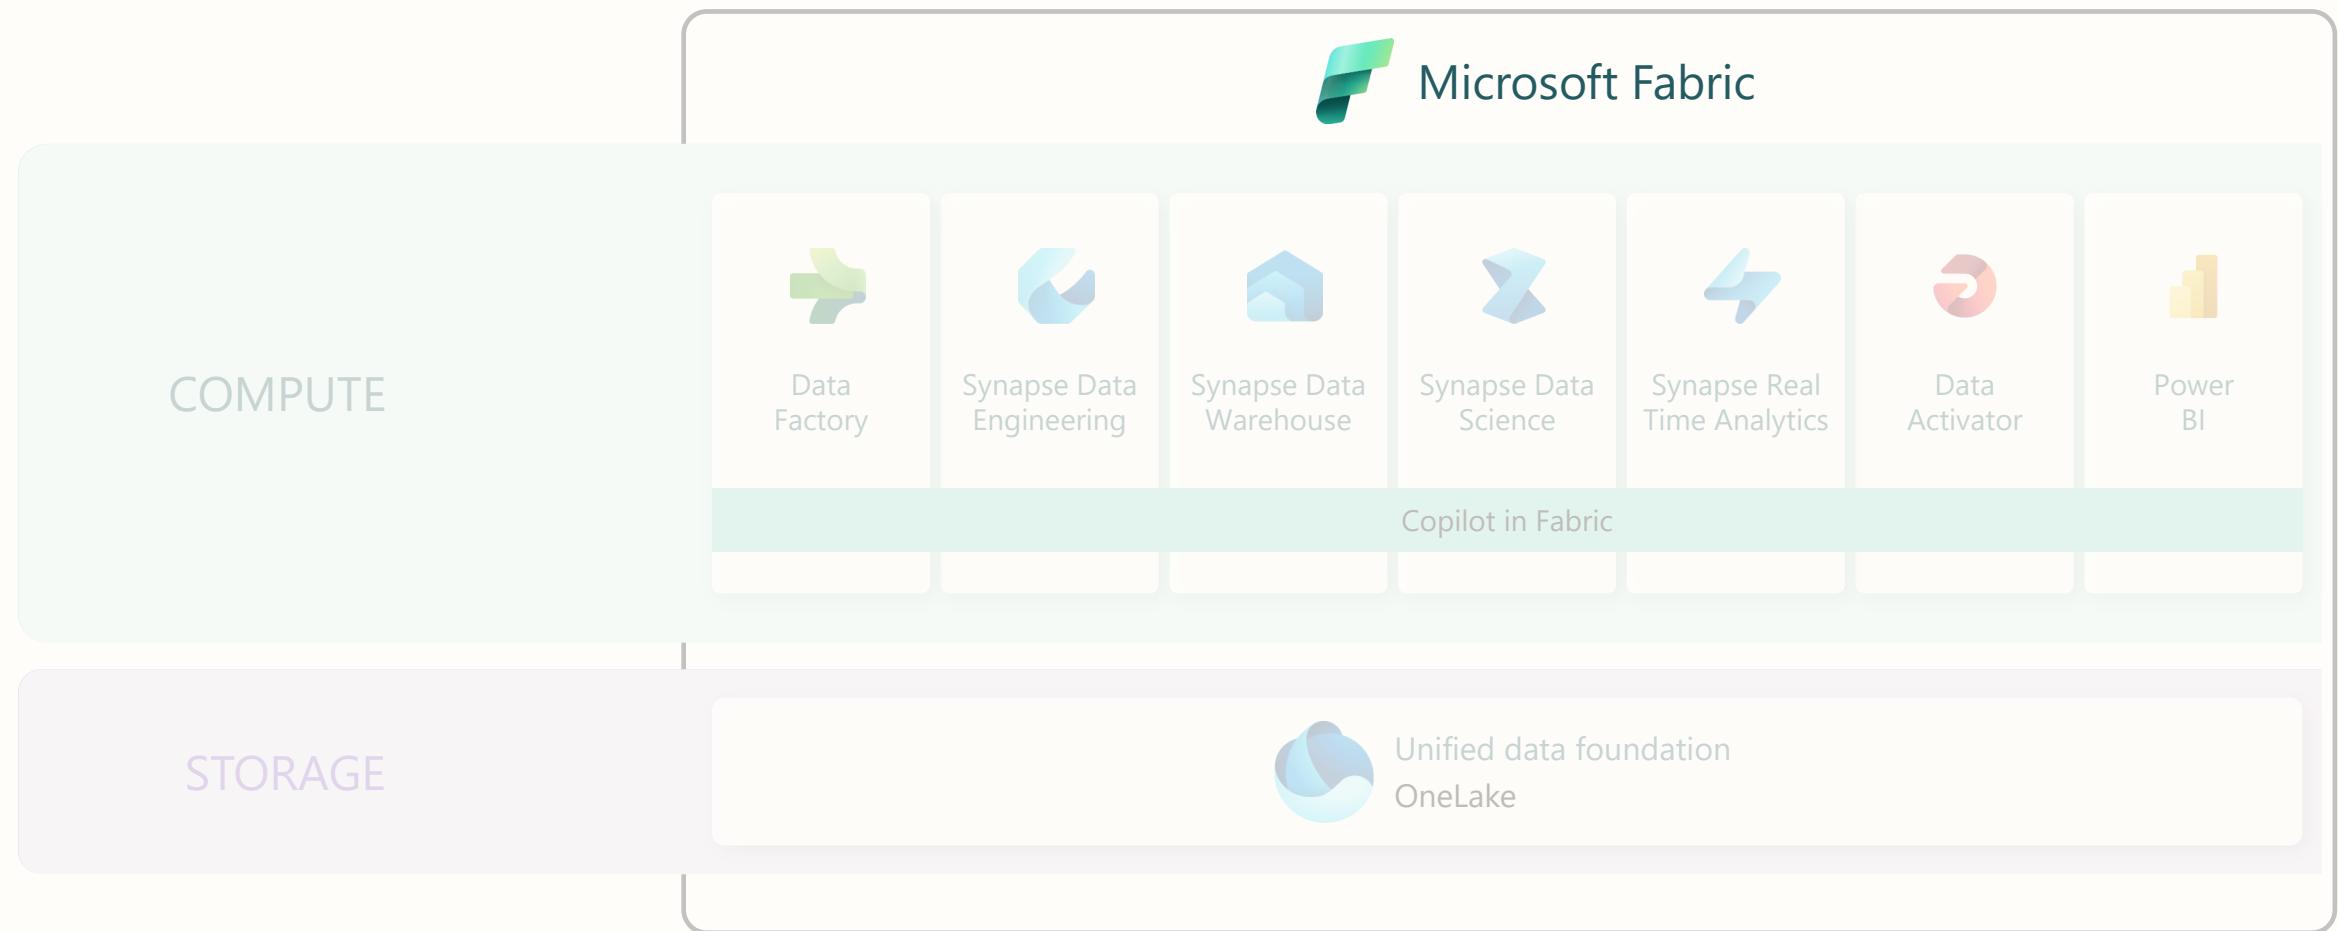


Data Transfer & Internet Egress

Cross-region data transfer network charges may apply based on source/destination of each storage access. Learn more at the [Bandwidth Pricing](#).

Note: OneLake cache is billed for KQL cache storage and Data Activator data retained.

Microsoft Fabric business model overview



How it works

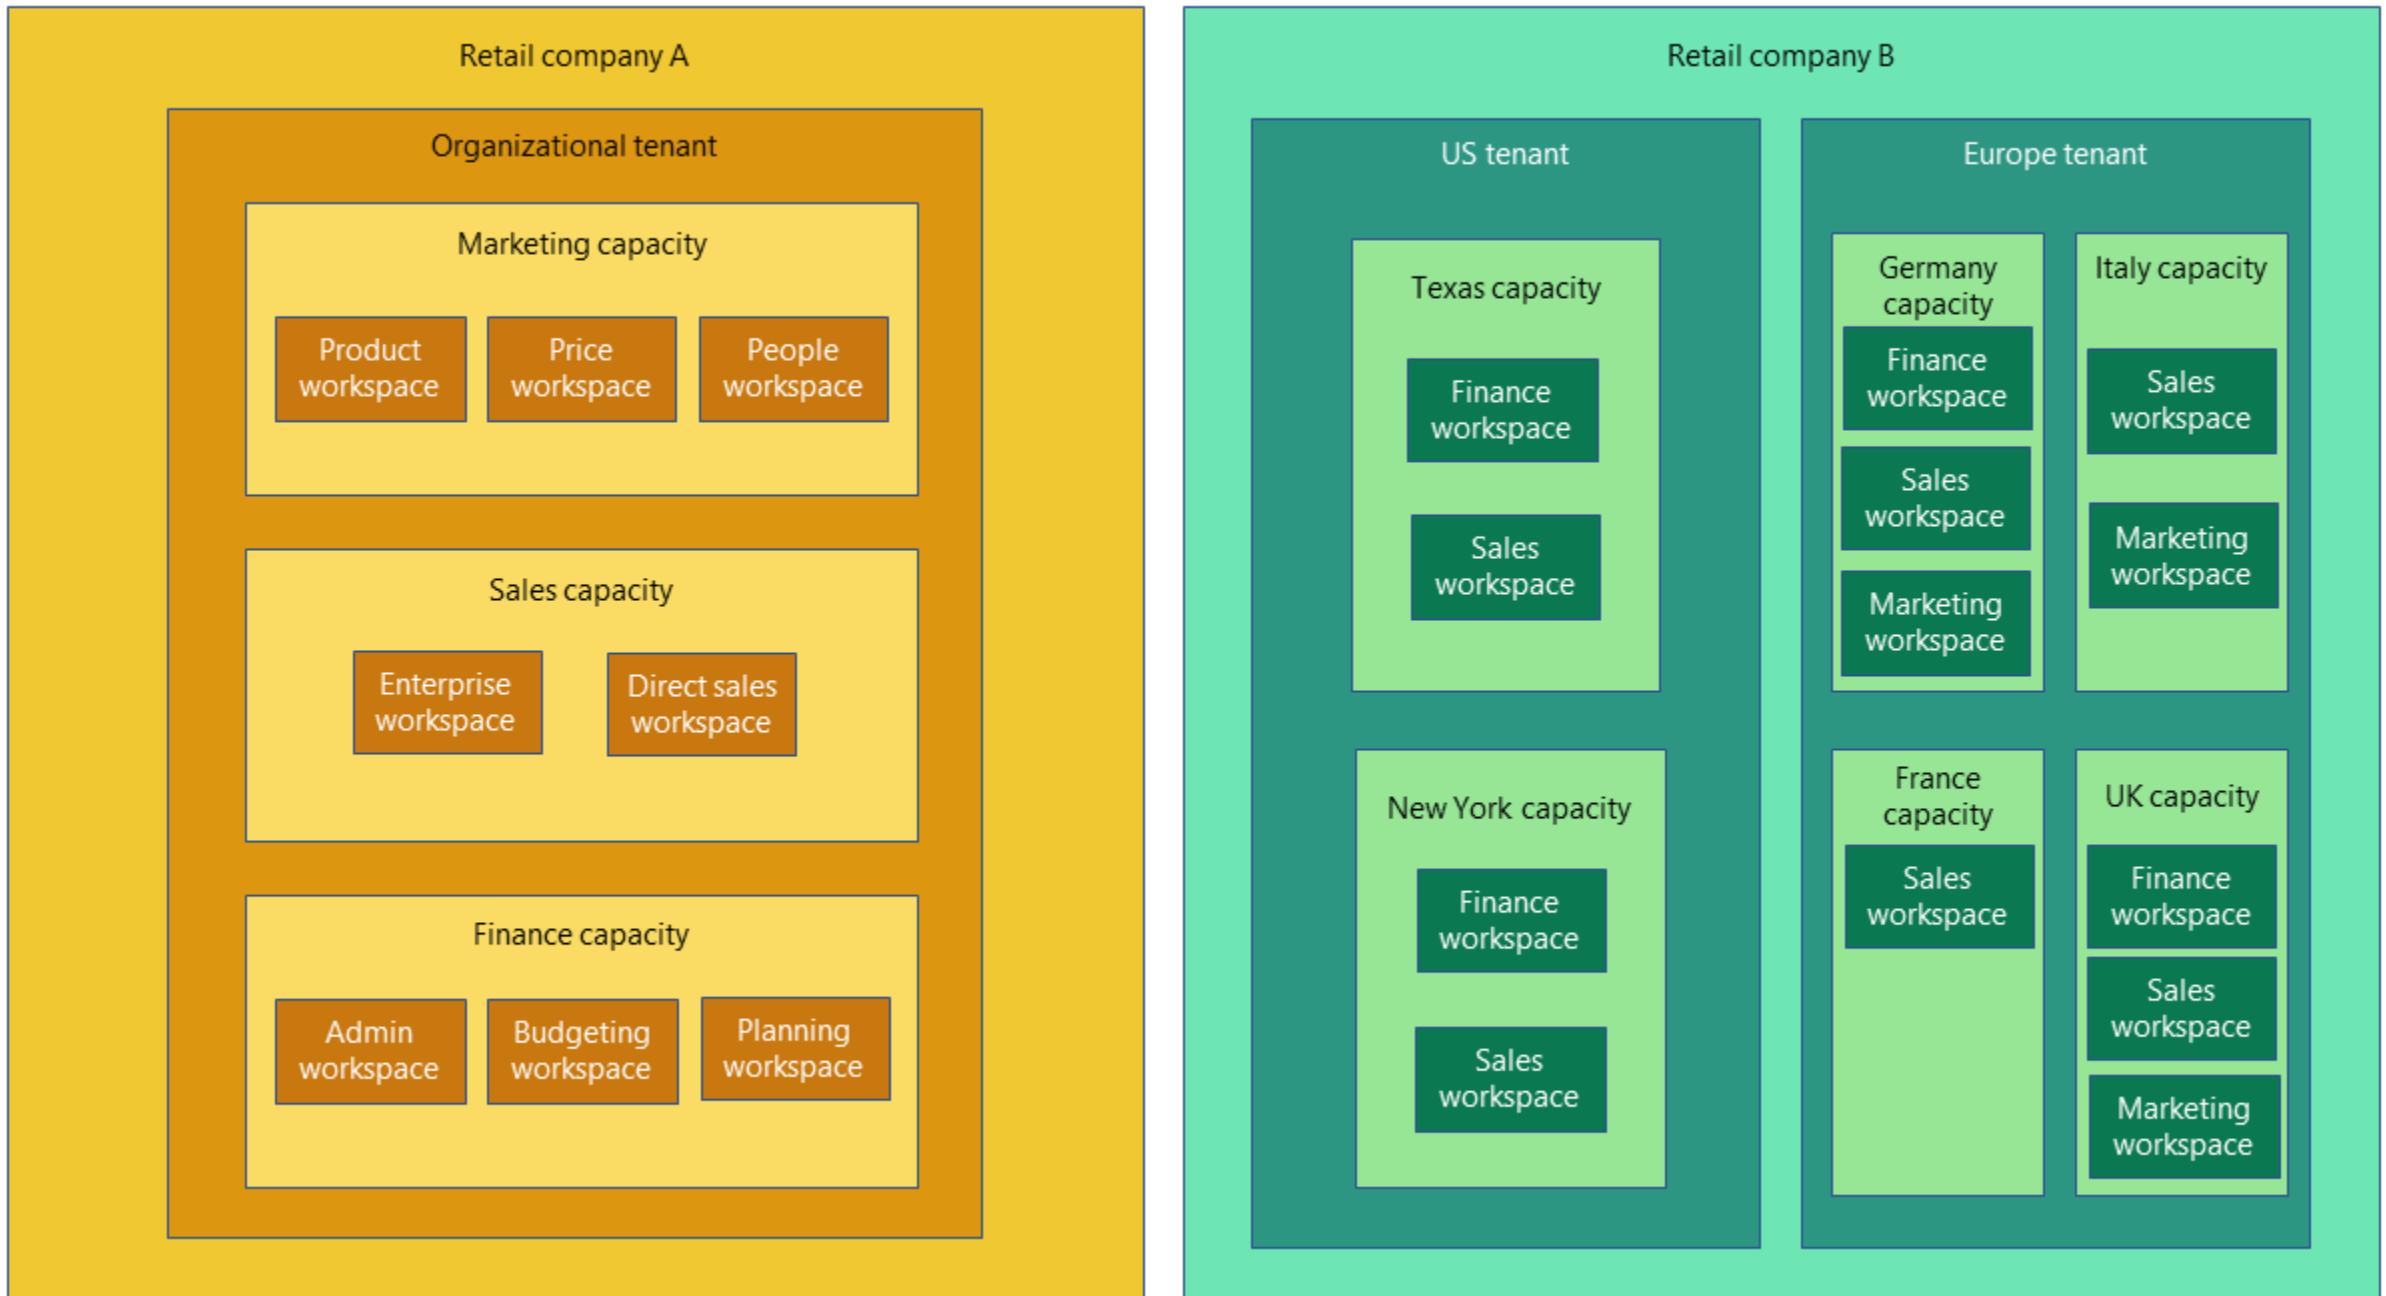
Capacity units (CUs) = Compute power

Capacity units (CUs) are units of measure that represent a pool of compute power needed. Compute power is required to run queries, jobs, or tasks.

CU Consumption

The CU consumption is highly correlated to the underlying compute effort needed for the tasks performed by the capability during the processing time.

Each capability, such as Power BI, Spark, Data Warehouse, with the associated queries, jobs, or tasks has a unique consumption rate.





Microsoft Fabric

How to License Compute in Fabric

Compute

Power BI Premium Capacity
P SKU: 1-5

Fabric Capacity
F SKU: 2-2048

- Purchased via M365 on the EA
 - P SKU powers all Fabric compute workloads with no additional action required
 - Monthly commitment
 - Global pricing
-
- Purchased via Azure Portal
 - Smaller entry-level starting at F2
 - Pay-as-you-go with pause/resume or 1yr reserved instance at ~40.5% discount
 - Purchase against MACC
 - Regional pricing
 - **Include all advanced security features**



Microsoft Fabric

Advanced Security Feature (F64 or Higher)

P vs F Advanced Security Features Offering	P SKU	F SKU
1. Trusted Workspace Access to ADLS Gen2	No	Yes
2. Managed Private Endpoint for Spark	No	Yes
3. Workspace Level Private Link	No	Yes



Microsoft Fabric

Transitioning Power BI Premium per capacity to Fabric

End of sales for new customers

Existing customers transition to Fabric

End of Sales

End of Renewals

End of Life (non-EA agreements)

End of Life (EA agreements)

July 1, 2024

Jan 1, 2025

Jan 1, 2026

Jan 1, 2028

For additional documentation on the Power BI Premium per capacity transition to Fabric, please see here: <https://aka.ms/P-SKU-Blog>

Translation for existing Power BI Premium per capacity Customers

Today	
SKU	PBI vCore
P1	8
P2	16
P3	32
P4	64
P5	128



Microsoft Fabric Generally Available	
SKU	CU
F 64	64
F 128	128
F 256	256
F 512	512
F 1024	1024

- Power BI Premium P SKUs get access to the full Fabric capabilities excluding network security
- 8 CUs are providing the same compute power of 1 Power BI Premium v-core for Power BI workloads. That is, F64 provides the equivalent compute power of Power BI Premium P1



Fabric capacity pricing

Fabric capacity is priced uniquely across regions. The following table shows the pricing at West Europe for reference. Fabric capacity can be purchased at Azure portal. Visit [Fabric pricing page](#) for more details.

1 CU pay-as-you-go price at West Europe €0,1981/hour

SKU	Capacity unit (CU)	Pay-as-you-go (monthly)	Reservation (monthly) ~40.5% saving over Pay-as-you-go
F 2	2	€289,24	€172,00
F 4	4	€578,48	€343,99
F 8	8	€1.156,96	€687,98
F 16	16	€2.313,91	€1.375,96
F 32	32	€4.627,83	€2.751,91
F 64	64	€9.255,65	€5.503,83
F 128	128	€18.511,30	€11.007,66
F 256	256	€37.022,60	€22.015,31
F 512	512	€74.045,21	€44.030,62
F 1024	1024	€148.090,41	€88.061,23
F 2048	2048	€296.180,820	€176.122,47

Note:

1. 1 CU PAYGO monthly rate calculation: €0,1981*730 =€144,62. F2 =€144,62*2=€289,24

2. Power BI Pro license is required for all Power BI Premium ("P") and Fabric Capacity ("F") SKUs to publish Power BI content to Microsoft Fabric. Enabling content consumers to review and interact with Power BI reports without additional paid per-user licenses is available at P1 and above (and F64 and above).

Copilot in Microsoft Fabric

Accelerate productivity

Unlock deeper insights

Deliver custom AI experiences for your data

We announced the public preview of Copilot in Fabric at Ignite, starting with the Power BI, Data Factory, Synapse Data Engineering, and Synapse Data Science experiences

Copilot in Fabric will be rolling out in stages with the goal that all customers with purchased capacity have access by the end of March 2024

Copilot in Fabric is limited to customers who have purchased Fabric capacity (F64 or higher) or Power BI Premium capacity (P1 and above) and is not included in the Fabric free account or trial

To use Copilot, you can simply count Copilot usage against their existing Fabric or Power BI Premium capacity

Copilot in Microsoft Fabric is in Public Preview

Free 60-day trial



Microsoft Fabric



aka.ms/try-fabric



Licensing FAQ (1/2)

1. Does P SKU include all Fabric experiences? No, it does not include the newly released network connectivity capabilities
2. Does F SKU include all Fabric experiences? Yes, including Power BI Embedded ([link](#)). But Copilot in Fabric and unlimited sharing of Power BI items requires F64 or above ([link](#)).
3. Does Copilot for Fabric require *additional* licensing? No. It consumes resources from capacities but requires the size to be P1/F64 or above ([link](#)).
4. Does Copilot in Fabric include Copilot for Power BI? Yes ([link](#)). Microsoft 365 Copilot does NOT include Copilot for Power BI.



Licensing FAQ (2/2)

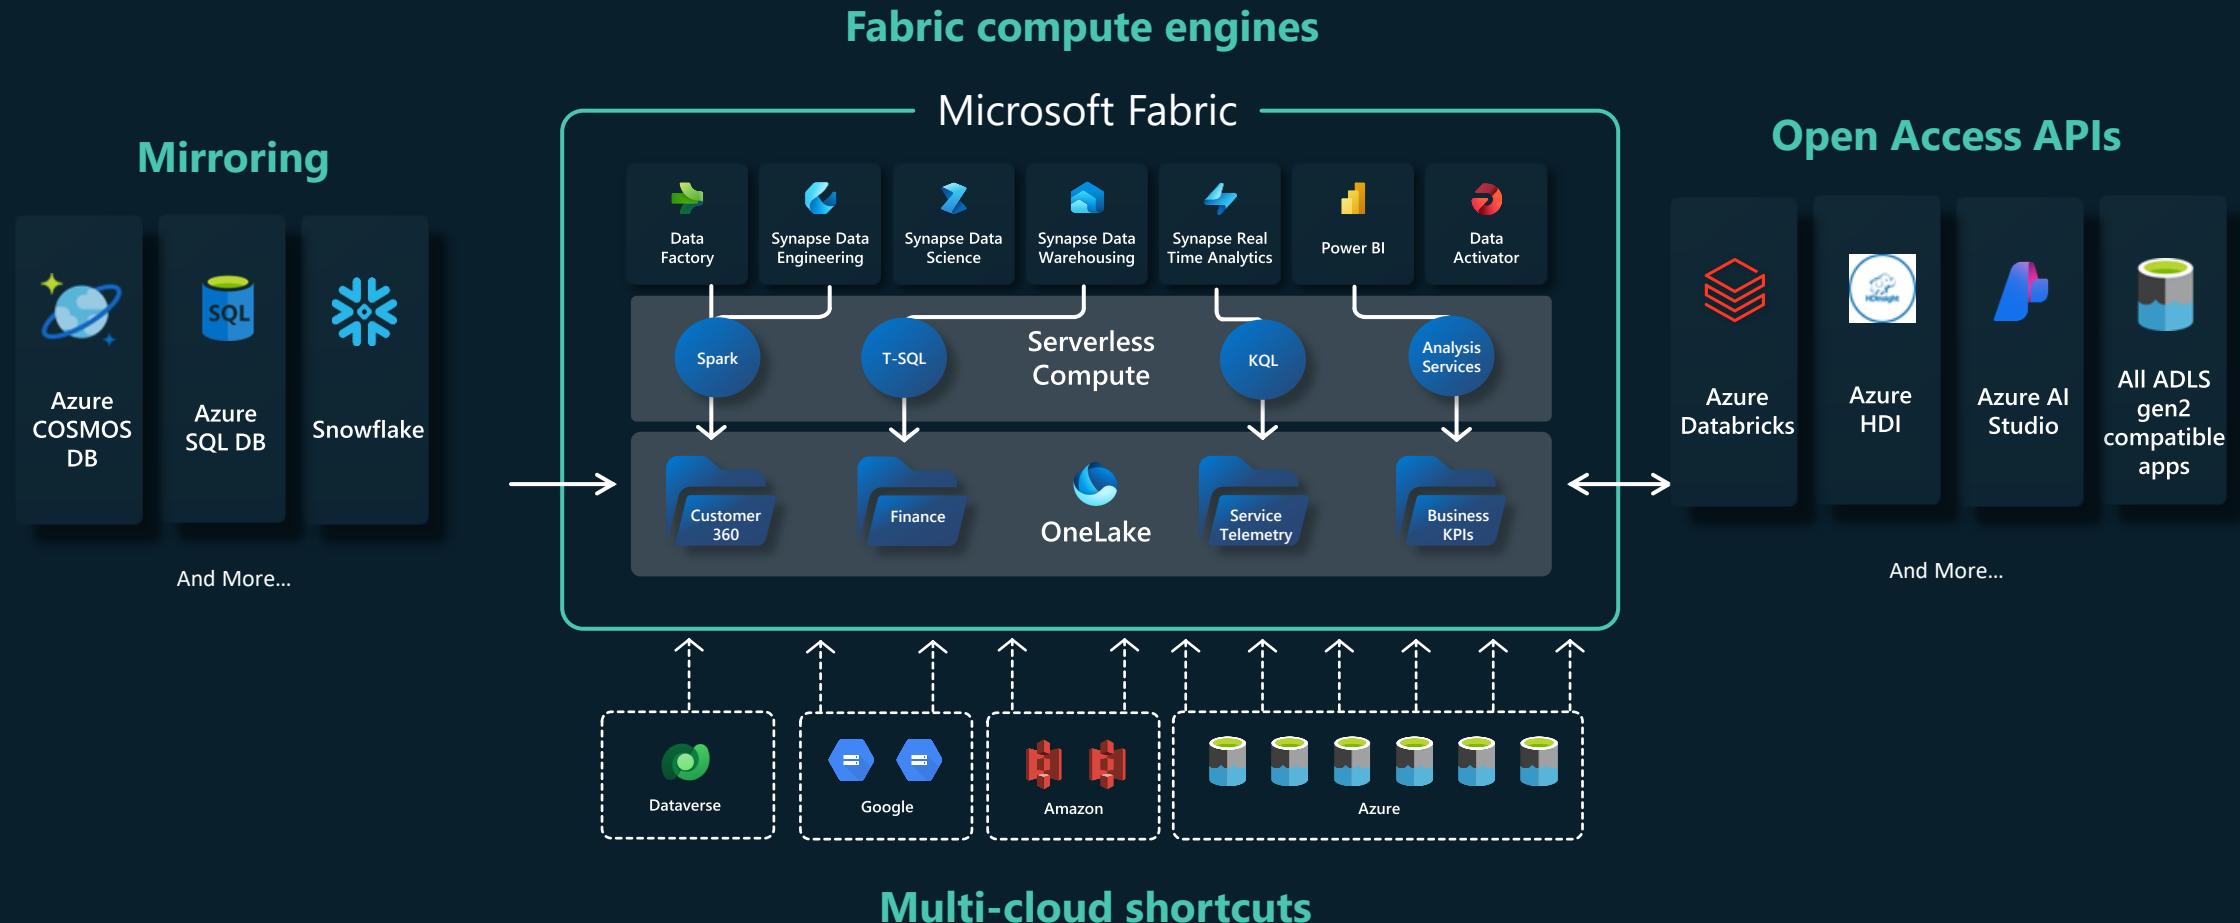
5. Is F SKU MACC eligible? Yes, but P SKU is not ([link](#)).
6. What capacity size do I need? Test your workload for free using a trial capacity (aka.ms/try-fabric) and monitor the capacity usage with the Fabric Capacity Metrics App ([link](#)) for an accurate indication.
7. Does Power BI Premium Per User license include Microsoft Fabric? No ([link](#)).
8. Do I still need a Power BI Pro license? Yes, to share Power BI content with other users. Consumers of Power BI content do NOT need a Power BI Pro license if the workspace is backed by capacity size F64/P1 or larger ([link](#)).



Transitioning to Microsoft Fabric

All roads lead to OneLake

Creating Data Gravity in OneLake



Releasing task flows soon!

What Architecture should I use?

Select a task flow

General
A high-level data processing task flow that guides you through typical tasks and the items assigned to them.

Basic data analytics
A basic, step-by-step task flow for batch data analytics.

Data analytics using a SQL analytics endpoint
Select unstructured, semi-structured, or structured data from lakehouse files, and then create reports.

Warehouse to datamart
Centralize data in warehouse and tailor it to datamarts.

Medallion
Organize and improve data progressively as it moves through each layers.

Event analytics
Process and analyze real-time data as it is generated to extract insights quickly.

Lambda
Process batch and real-time data in one data process flow.

Sensitive data insights
Process and analyze your sensitive data.

Basic machine learning models
Train machine learning models and generate insights to drive decisions.

End-to-end datamart
Create rich visualizations based on the data in your datamart using little or no code.

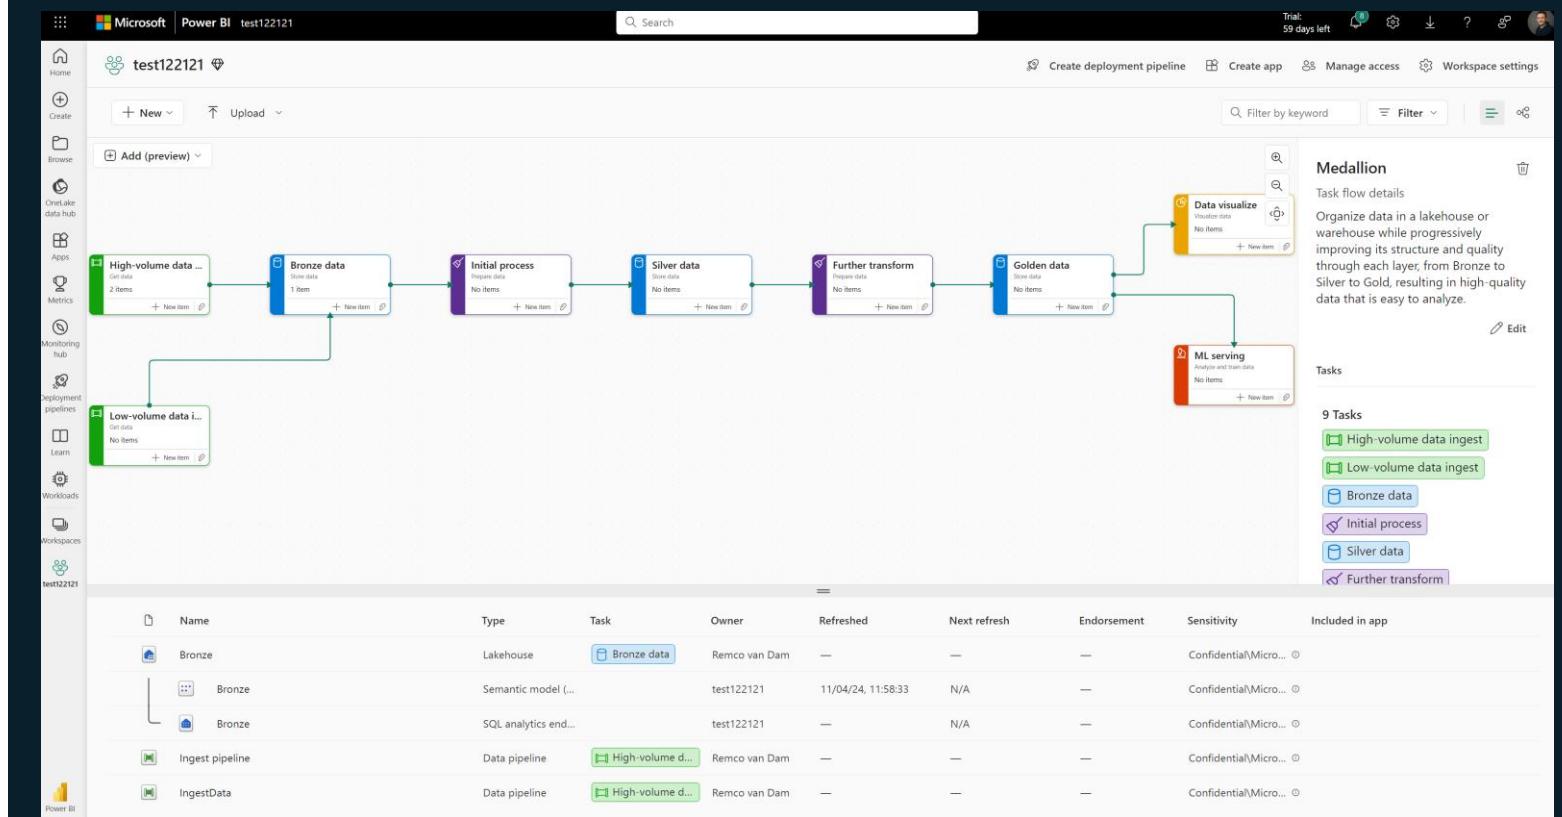
Required Workloads

- Data Engineering
- Data Factory
- Data Science
- Data Warehouse
- Real-Time Analytics
- Power BI

Required item types

- Data pipeline
- Eventstream
- Dataflow Gen2
- Notebook
- Spark Job Definition
- Lakehouse
- Warehouse
- Report
- Dashboard
- ML model
- Experiment
- Environment

Select **Cancel**



Medallion architecture

Although the 3-layered design is common and well-known, there are many discussions on the scope, purpose, and best practices on each of these layers.



Bronze layer

Typically raw, "as-is"

- Maintains the raw state in the structure "as-is"
- Data is immutable (read-only)
- Delivery-based partitioned tables, i.e., YYYYMMDD
- Mostly Delta Parquet. Sometimes other formats
- Can be any combination of streaming and batch transactions
- May include extra metadata (schema)
- May be fed from a "mediation layer"
- Used for debugging, testing



Silver layer

Cleaned, filtered

- Uses data quality rules for validation
- Usually only functional data
- Historization is merged (SCD2)
- Efficient storage format; Delta Parquet
- Versioning for rolling back
- Handles missing or incorrect data
- Usually enriched with reference data
- Source-oriented, although queryable and cluttered around subject areas
- Usually used by operational analytical teams

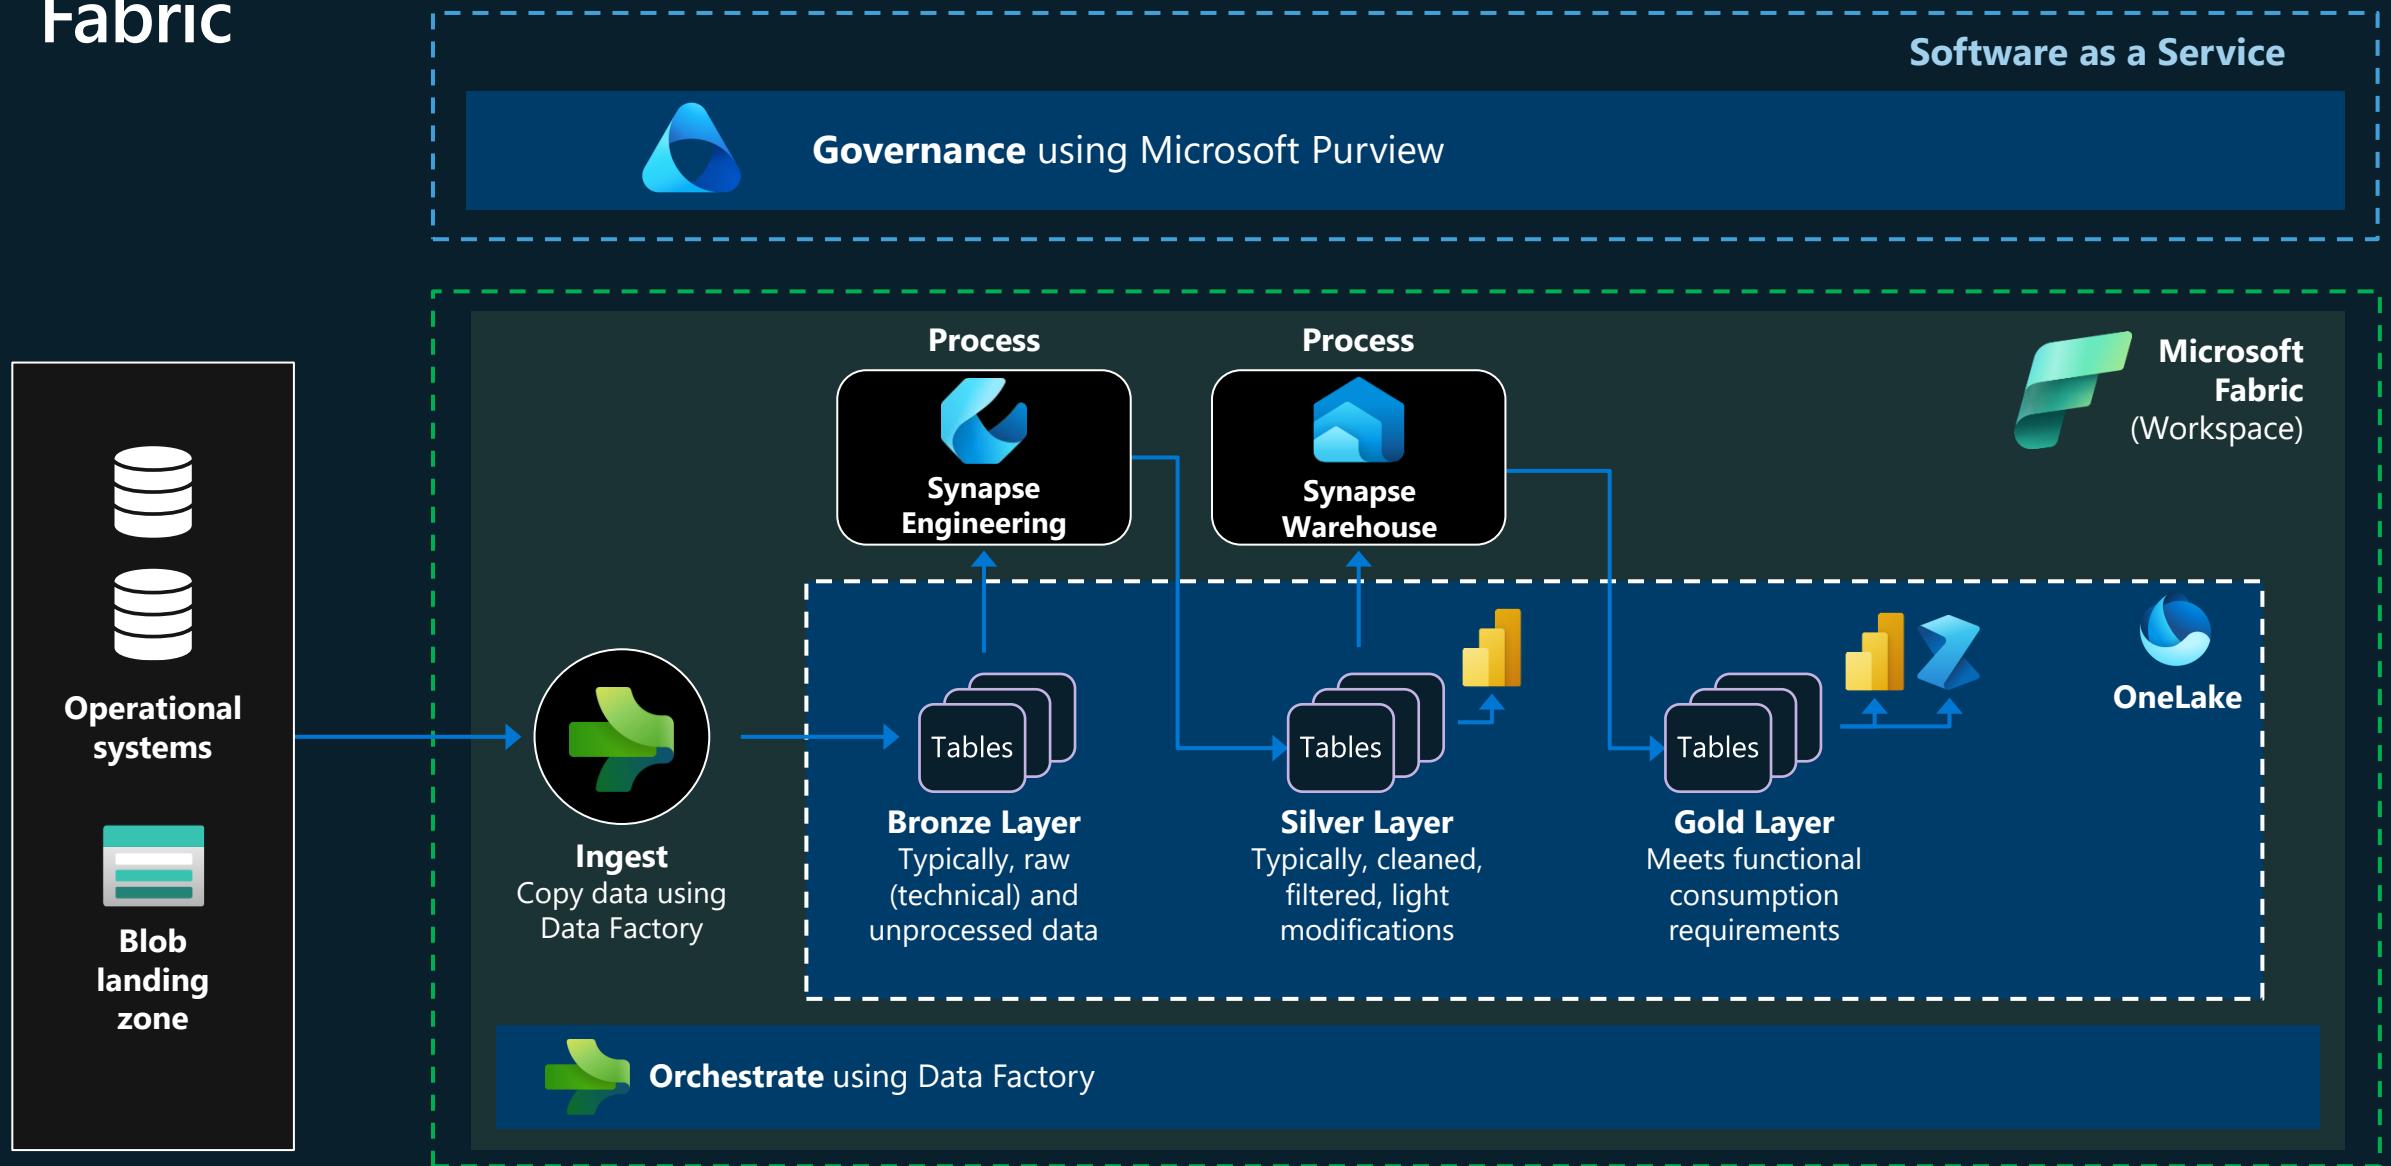


Gold layer

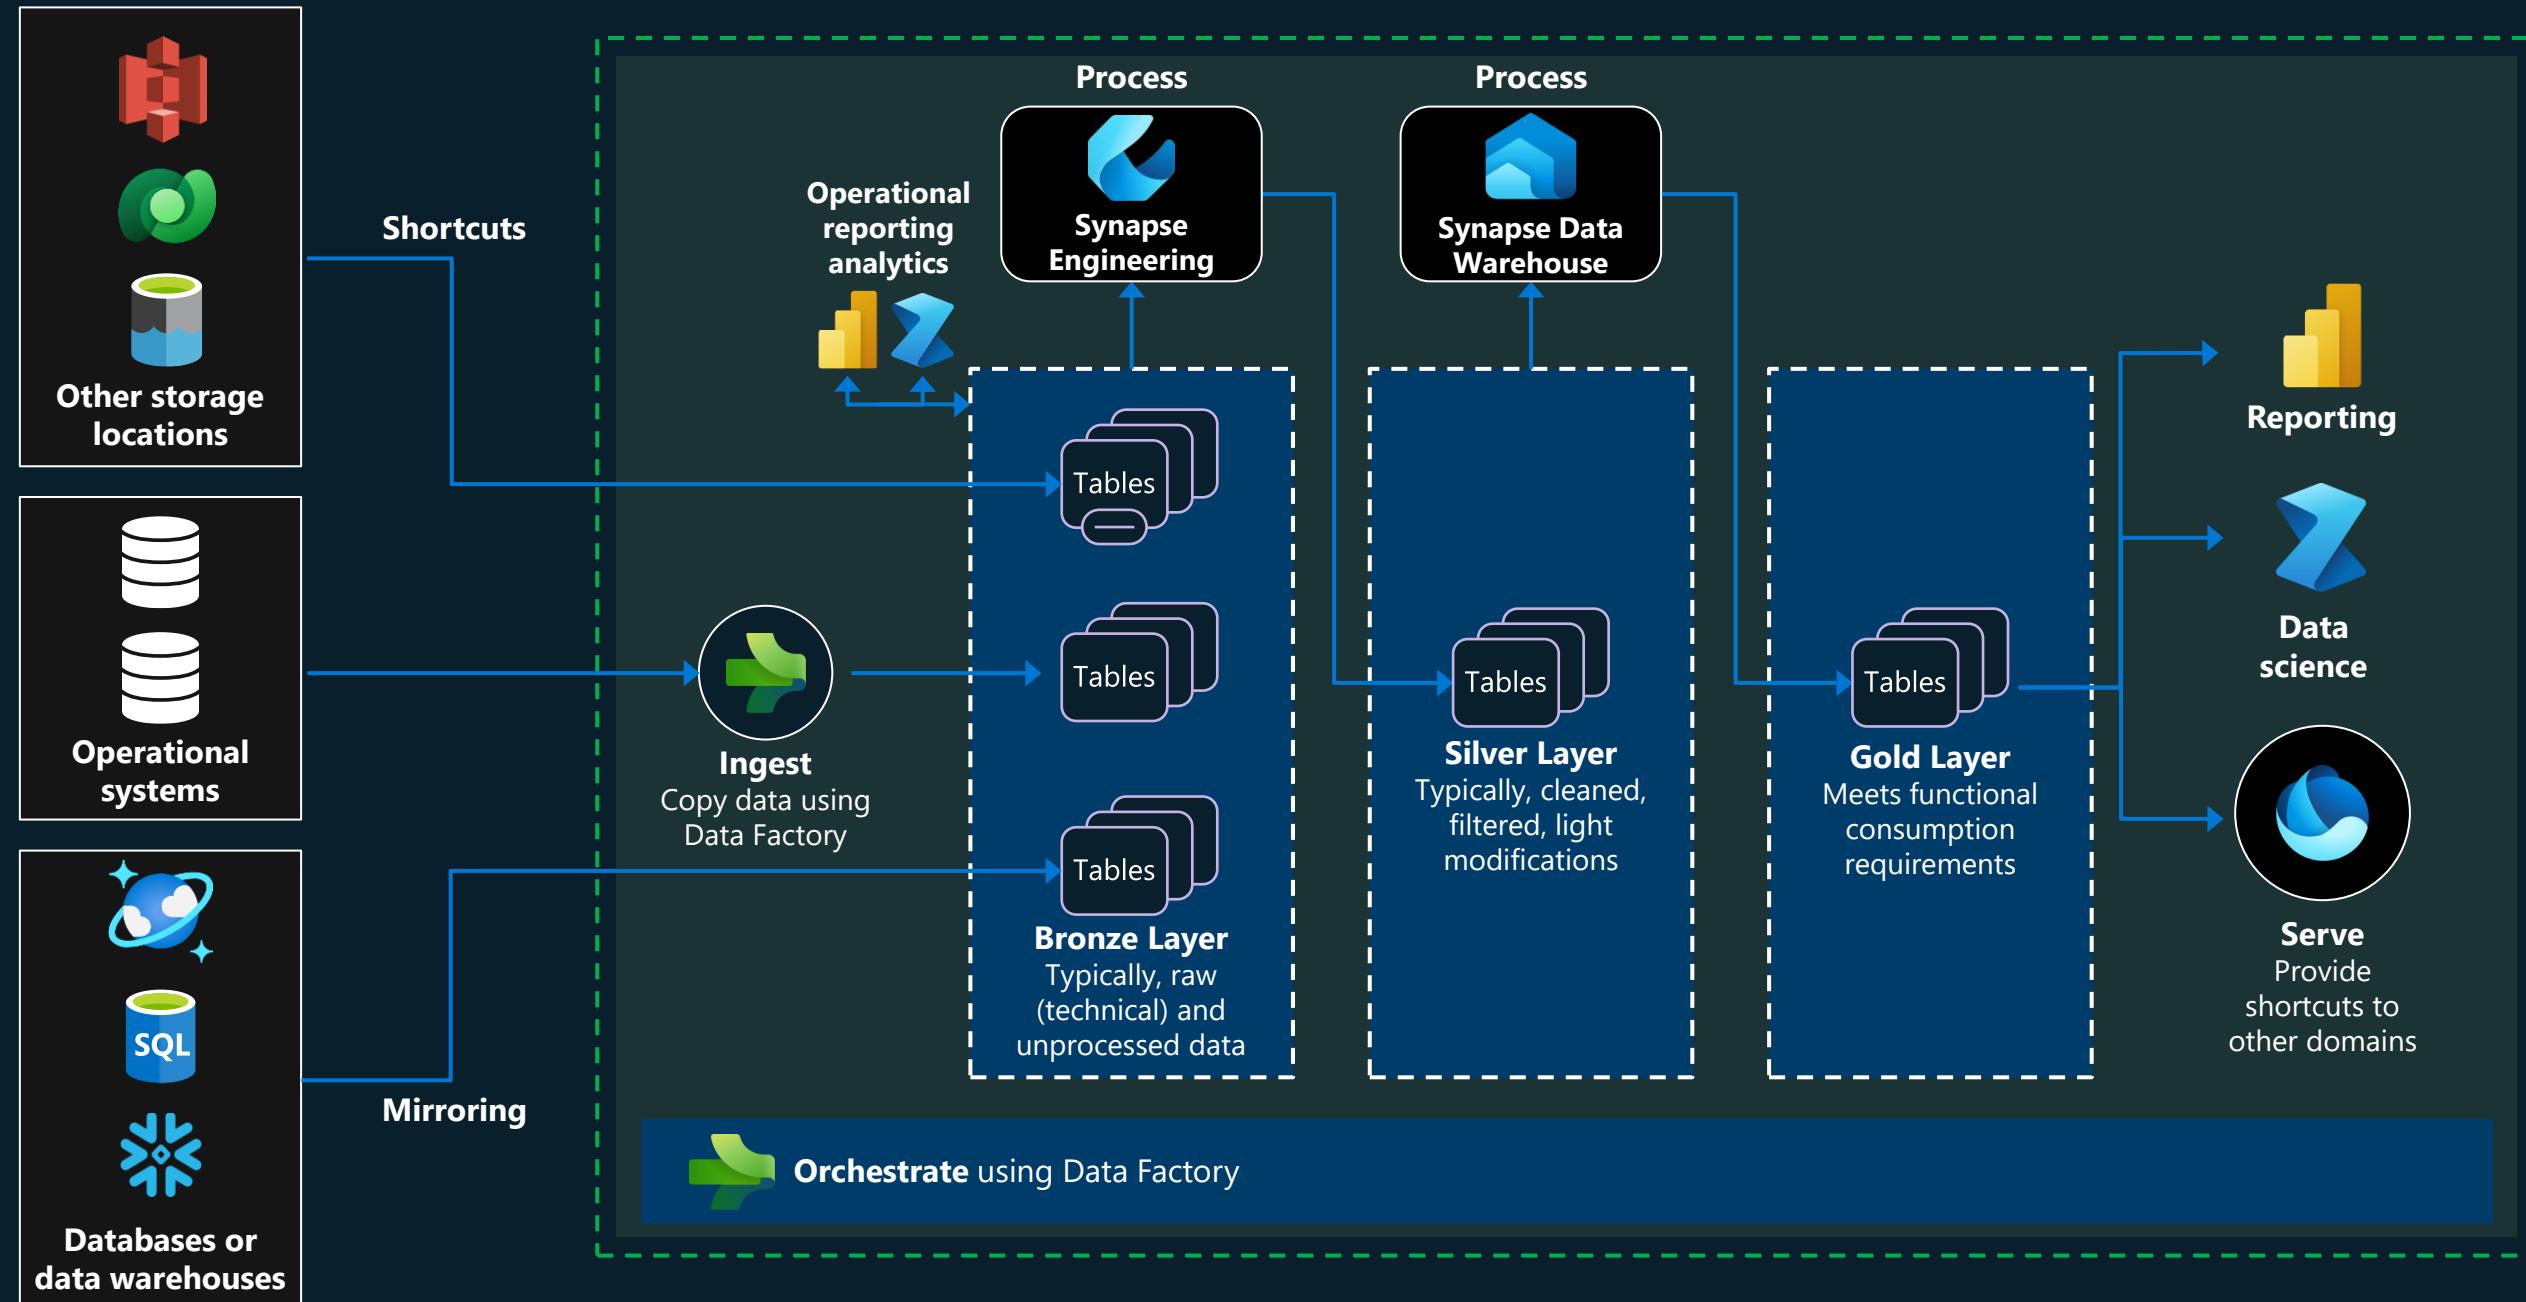
Refined business-level

- What enterprises call data products: consumer-ready / user-friendly data
- Data is highly governed and well-documented
- Historization is applied only for the set of use cases or consumers
- Contains complex business rules, such as calculations and enrichments
- Efficient storage format; Delta Parquet
- Versioning for rolling back
- Might contain additional sub layers for sharing or distributing data

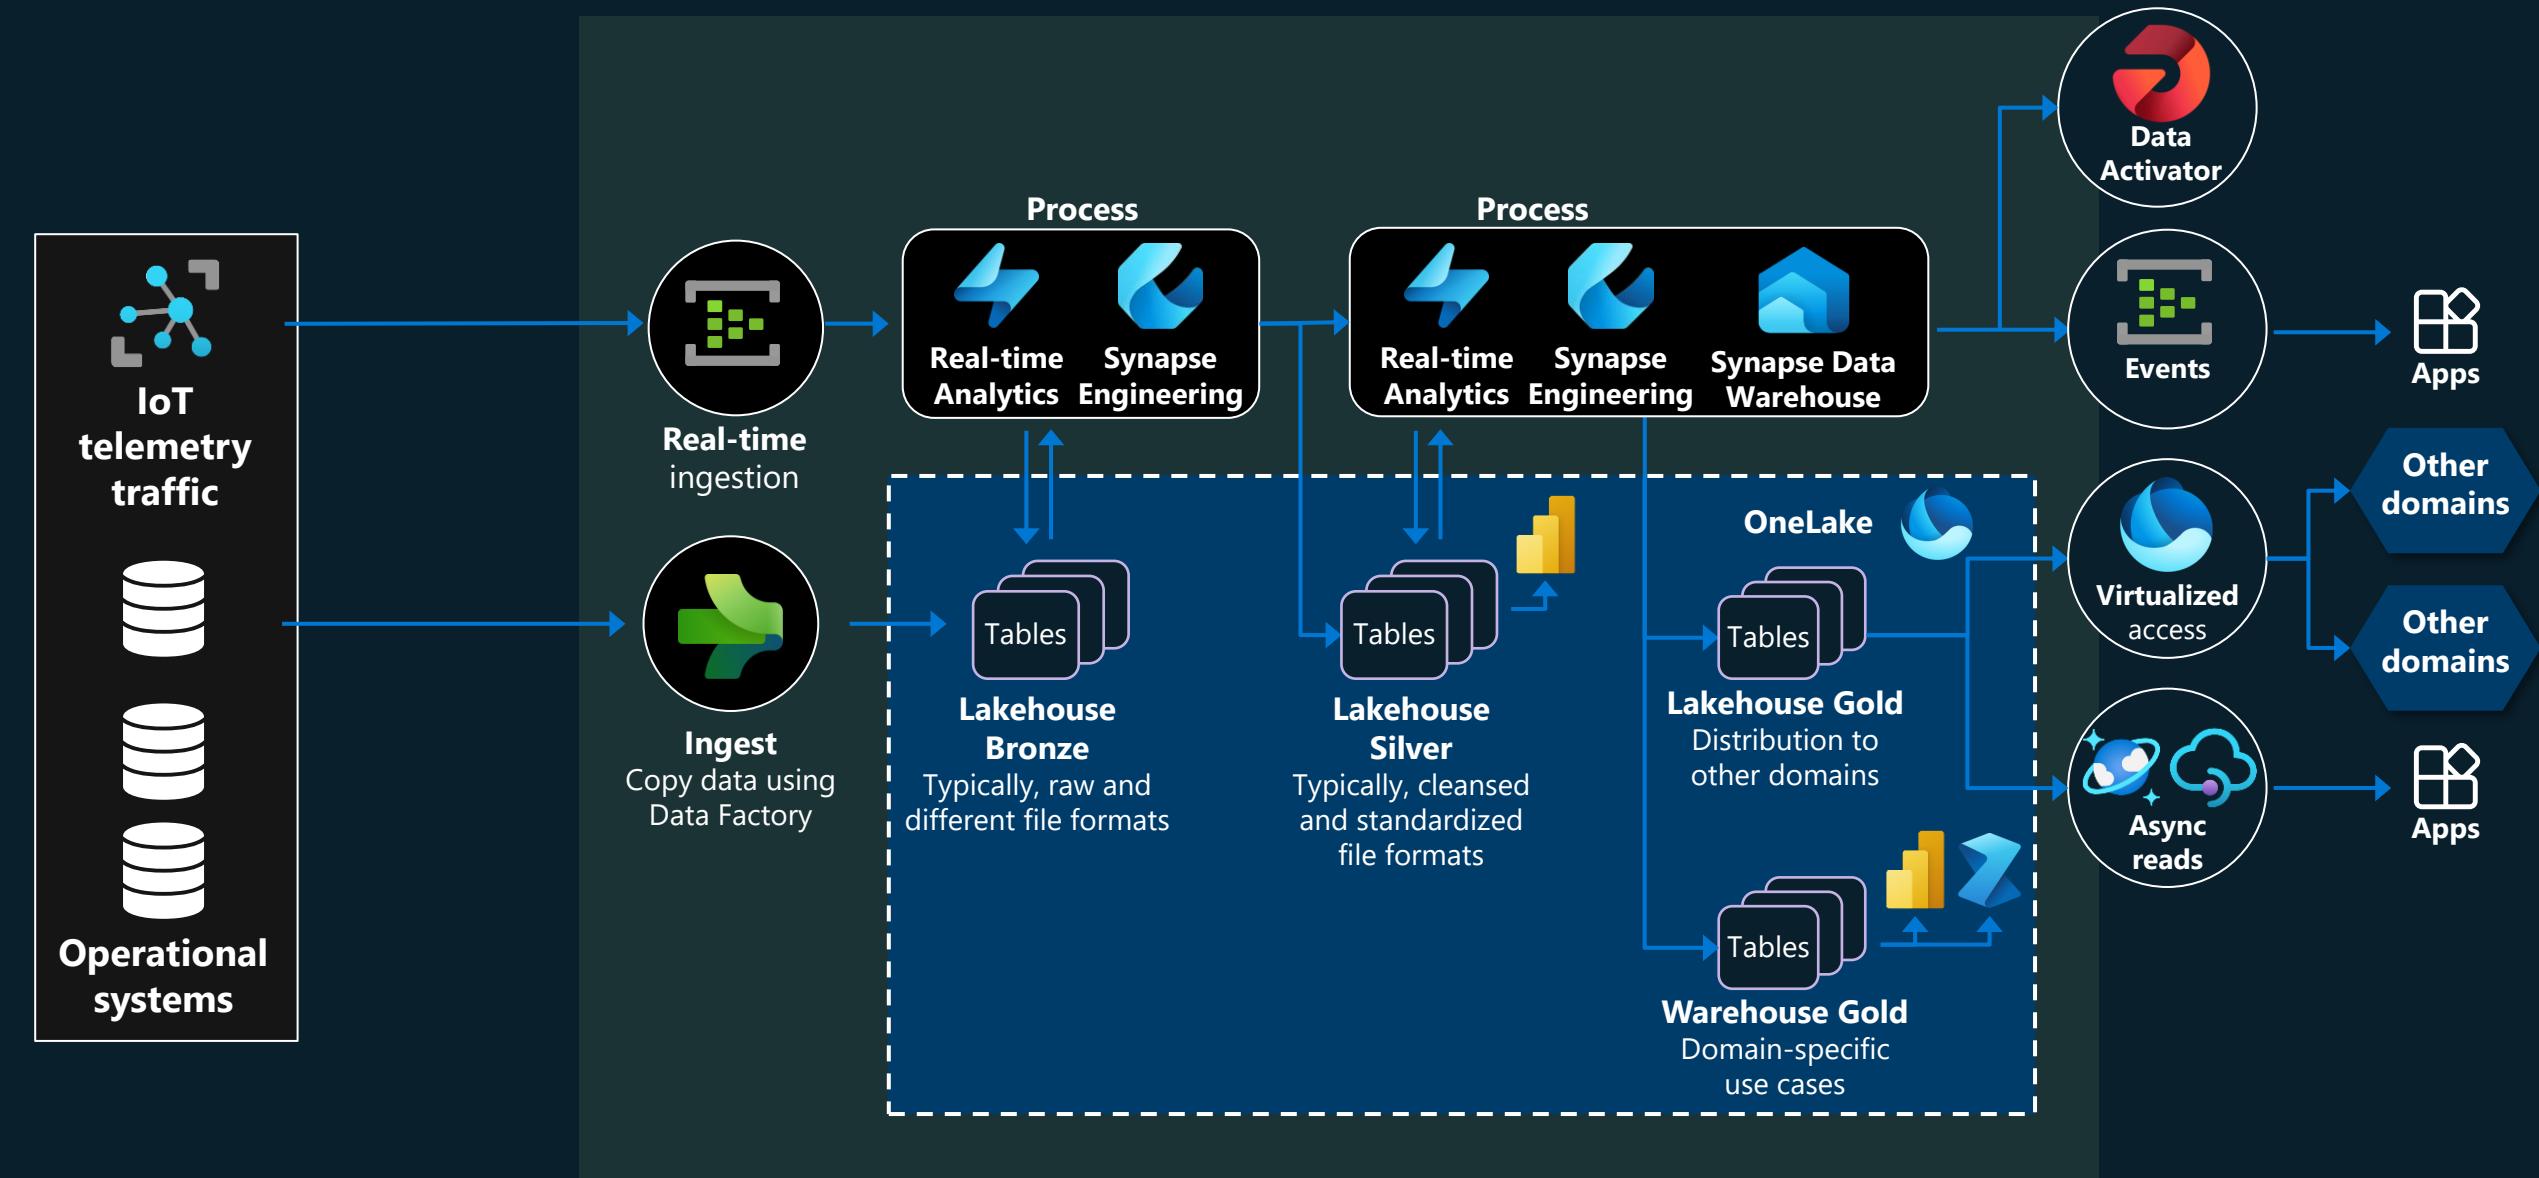
Logical overview of how the Lakehouse Architecture is presented for Fabric



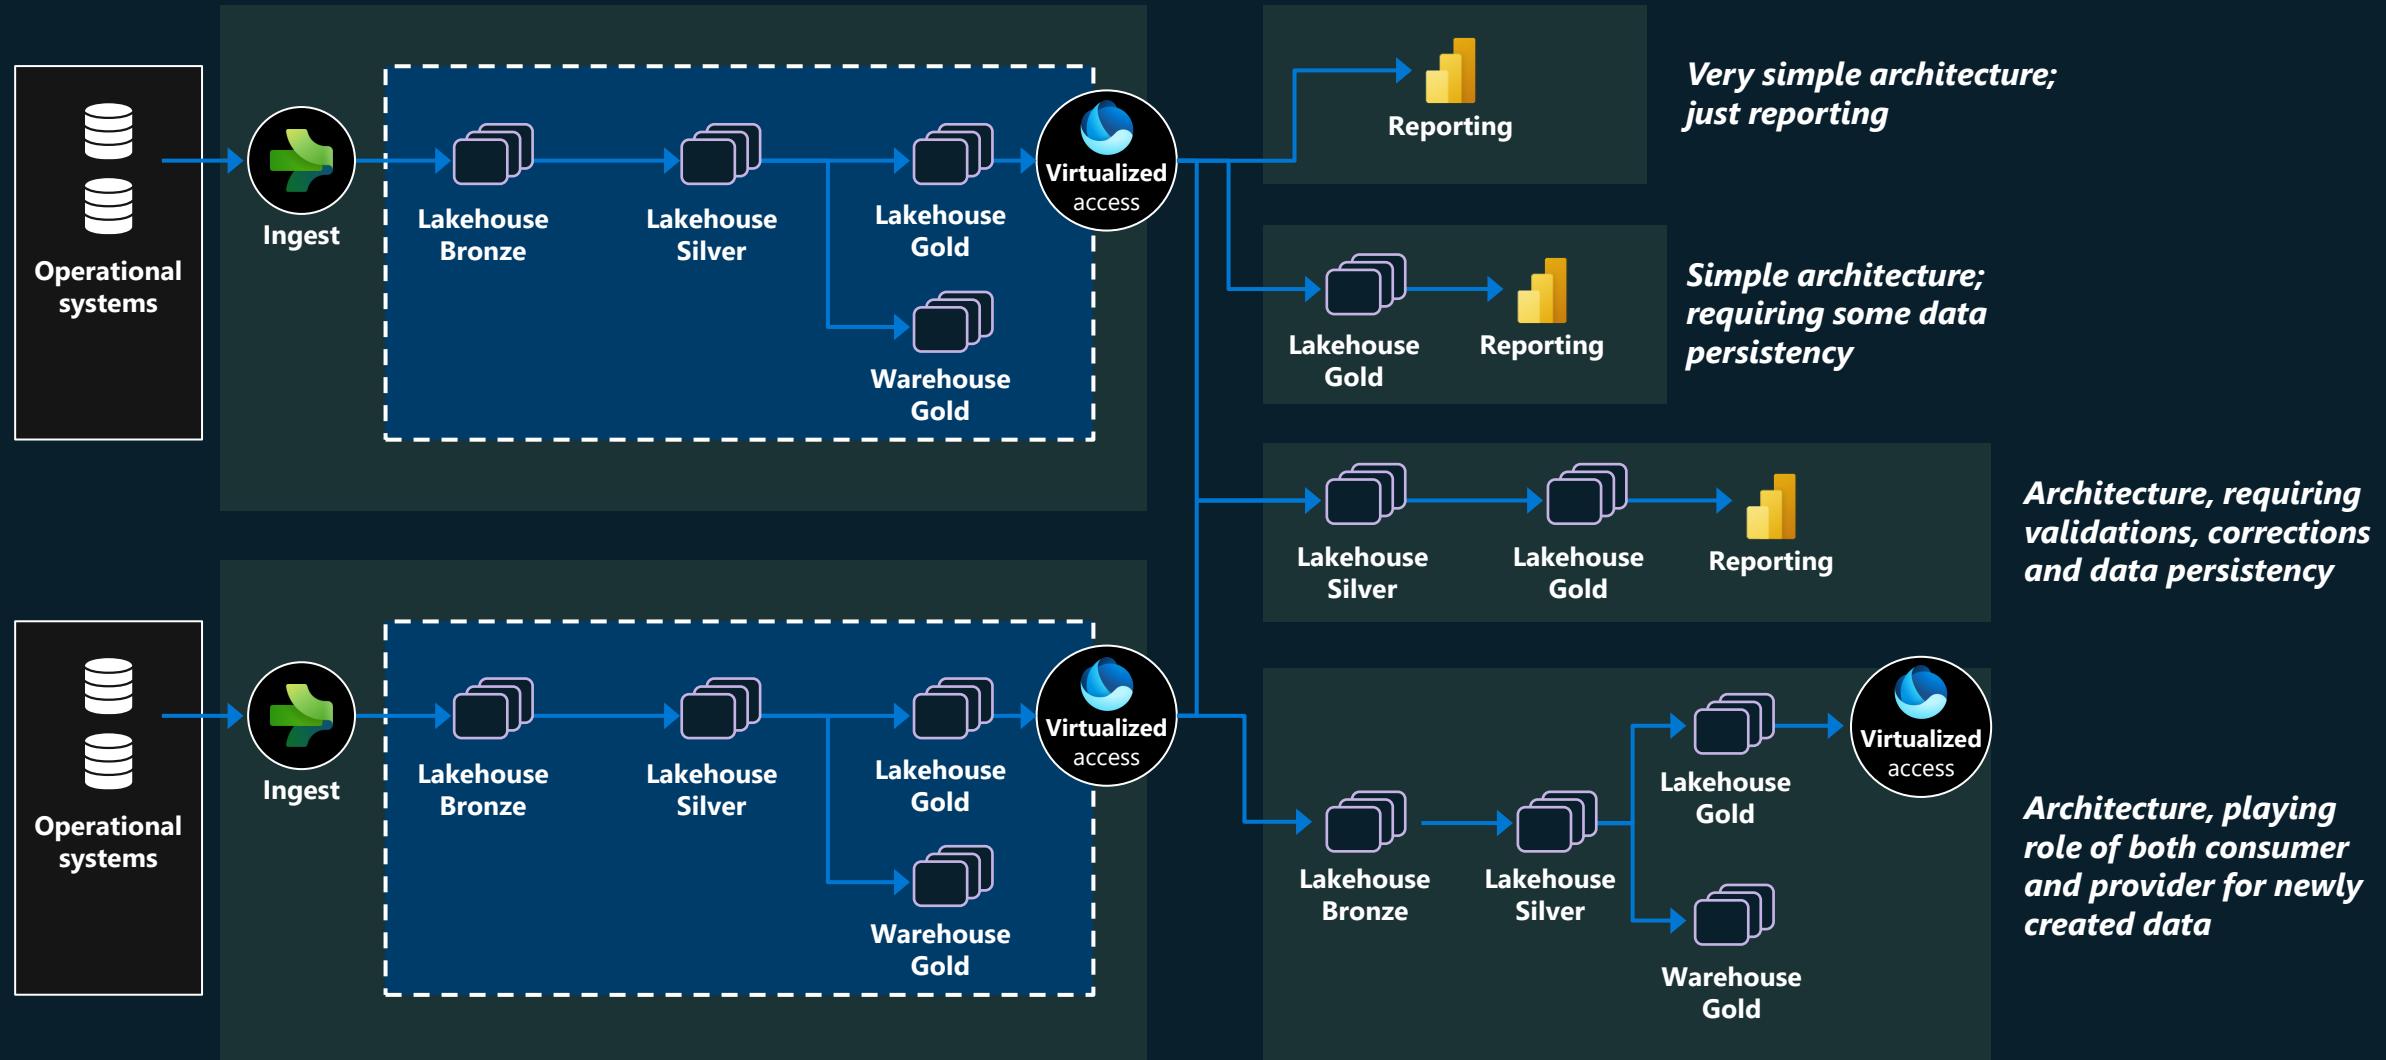
Bronze operating as a conglomerate of physical copies, shortcuts and mirroring



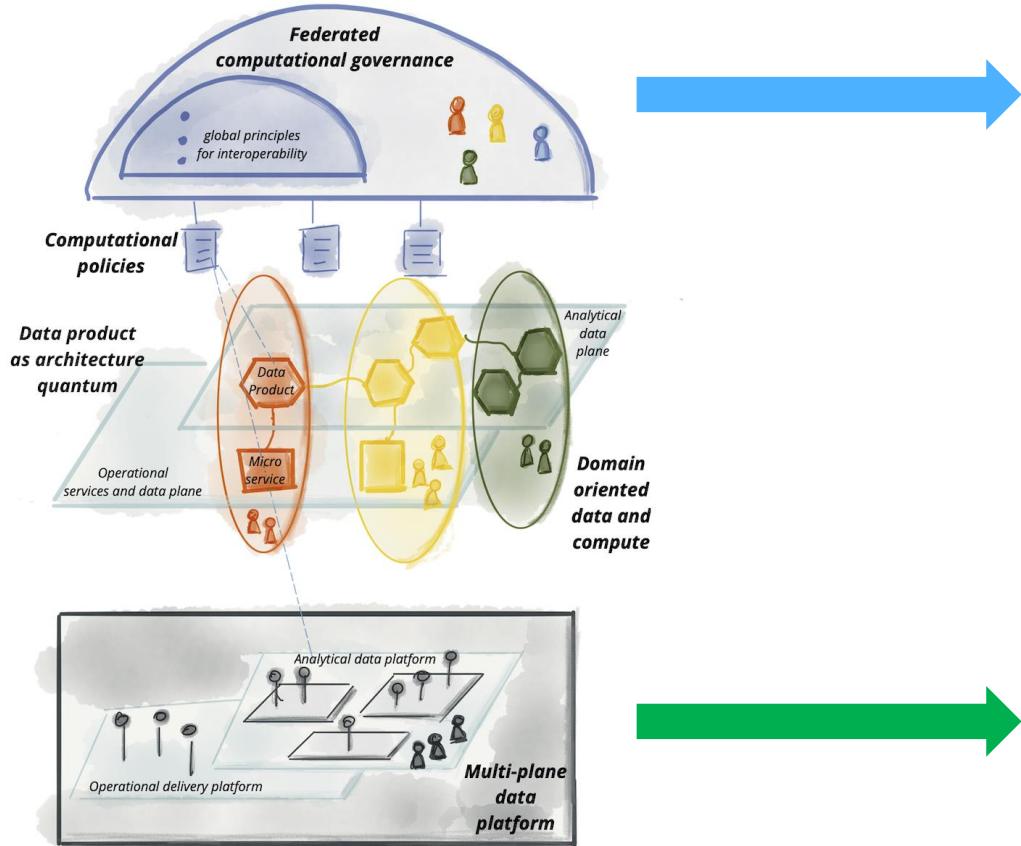
Adding API- and event-based integration patterns to the Lakehouse Architecture



Lakehouse architectures on the consuming side will also look quite different depending on consumer's archetype



Mapping concepts to a solution architecture

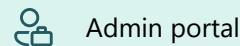


Data Mesh Principles by Zhamak Dehghani

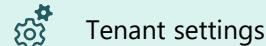
Gain full visibility and control with industry-leading features

Governance and security features built into Fabric

Manage your data estate



Admin portal



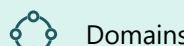
Tenant settings



Workspace settings



Capacities

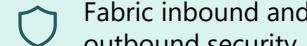


Domains

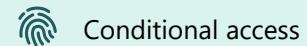


Workspaces

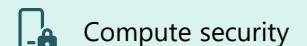
Secure, protect, and comply



Fabric inbound and outbound security



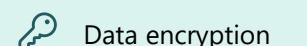
Conditional access



Compute security



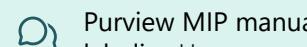
Workspace & item security



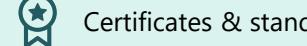
Data encryption



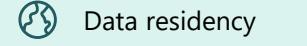
Universal security model*



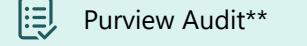
Purview MIP manual labeling**



Certificates & standards

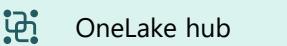


Data residency



Purview Audit**

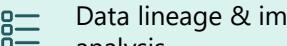
Encourage data discovery and use



OneLake hub



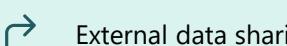
Endorsements



Data lineage & impact analysis



Metadata scanning



External data sharing*



Metadata curation



Discover & Reuse

Uncover insights and act



Admin monitoring



Capacity metrics



Monitoring hub



Purview Hub

Fabric metadata



Microsoft Purview



Comprehensive data security¹



Federated data governance¹



Risk & compliance posture¹



Fabric auto-attaches to Microsoft Purview

*Features coming soon



**Additional Microsoft Purview purchase required



Your Analytics upgrade path to Microsoft Fabric

Applicable for existing customers

Investment	Pattern	Solution	Availability
	#1 Azure Databricks only	Using Azure Databricks for entire implementation.*	Continue relying on Azure Databricks Lakehouse. Enable Microsoft Fabric with ADLS G2 Shortcuts to use the best of both worlds between ADB + Fabric.* Ready to use When using unity: <ul style="list-style-type: none">• Use notebook• Native integration is being worked on
	#2 Azure Databricks and Synapse	1. Using Azure Databricks and Spark and data engineering. 2. Synapse SQL Dedicated Pools for data warehouse and serving	1. ADLS G2 Shortcuts for OneLake. 2. Microsoft 1st Party solution for Synapse Gen2 to Fabric SQL code migration. 1. ADLS G2 Shortcuts ready to use. 2. Microsoft 1st Party solutions coming in 2024.
	#3 Synapse Dedicated Pools	Using Synapse SQL Dedicated Pools for most of implementation with some Spark.	Microsoft 1st Party solutions for DW replication and Synapse G2-to-Fabric SQL code migration. Microsoft 1st Party solutions coming in 2024. Accelerator available today
	#4 Power BI	Using Power BI Dataflow for data preparation or other legacy reporting services.	1. Test drive Dataflow G2 if currently using G1 Dataflow 2. Consider migrate to Power BI Premium if currently using on-prem reporting servers Ready to use

*Solution applicable for investments in Synapse Spark.

Better together through OneLake

Open and Governed Lakehouse



Microsoft Fabric



Azure Databricks

All Fabric engines have been redesigned and optimized for delta parquet as their storage format for tabular data

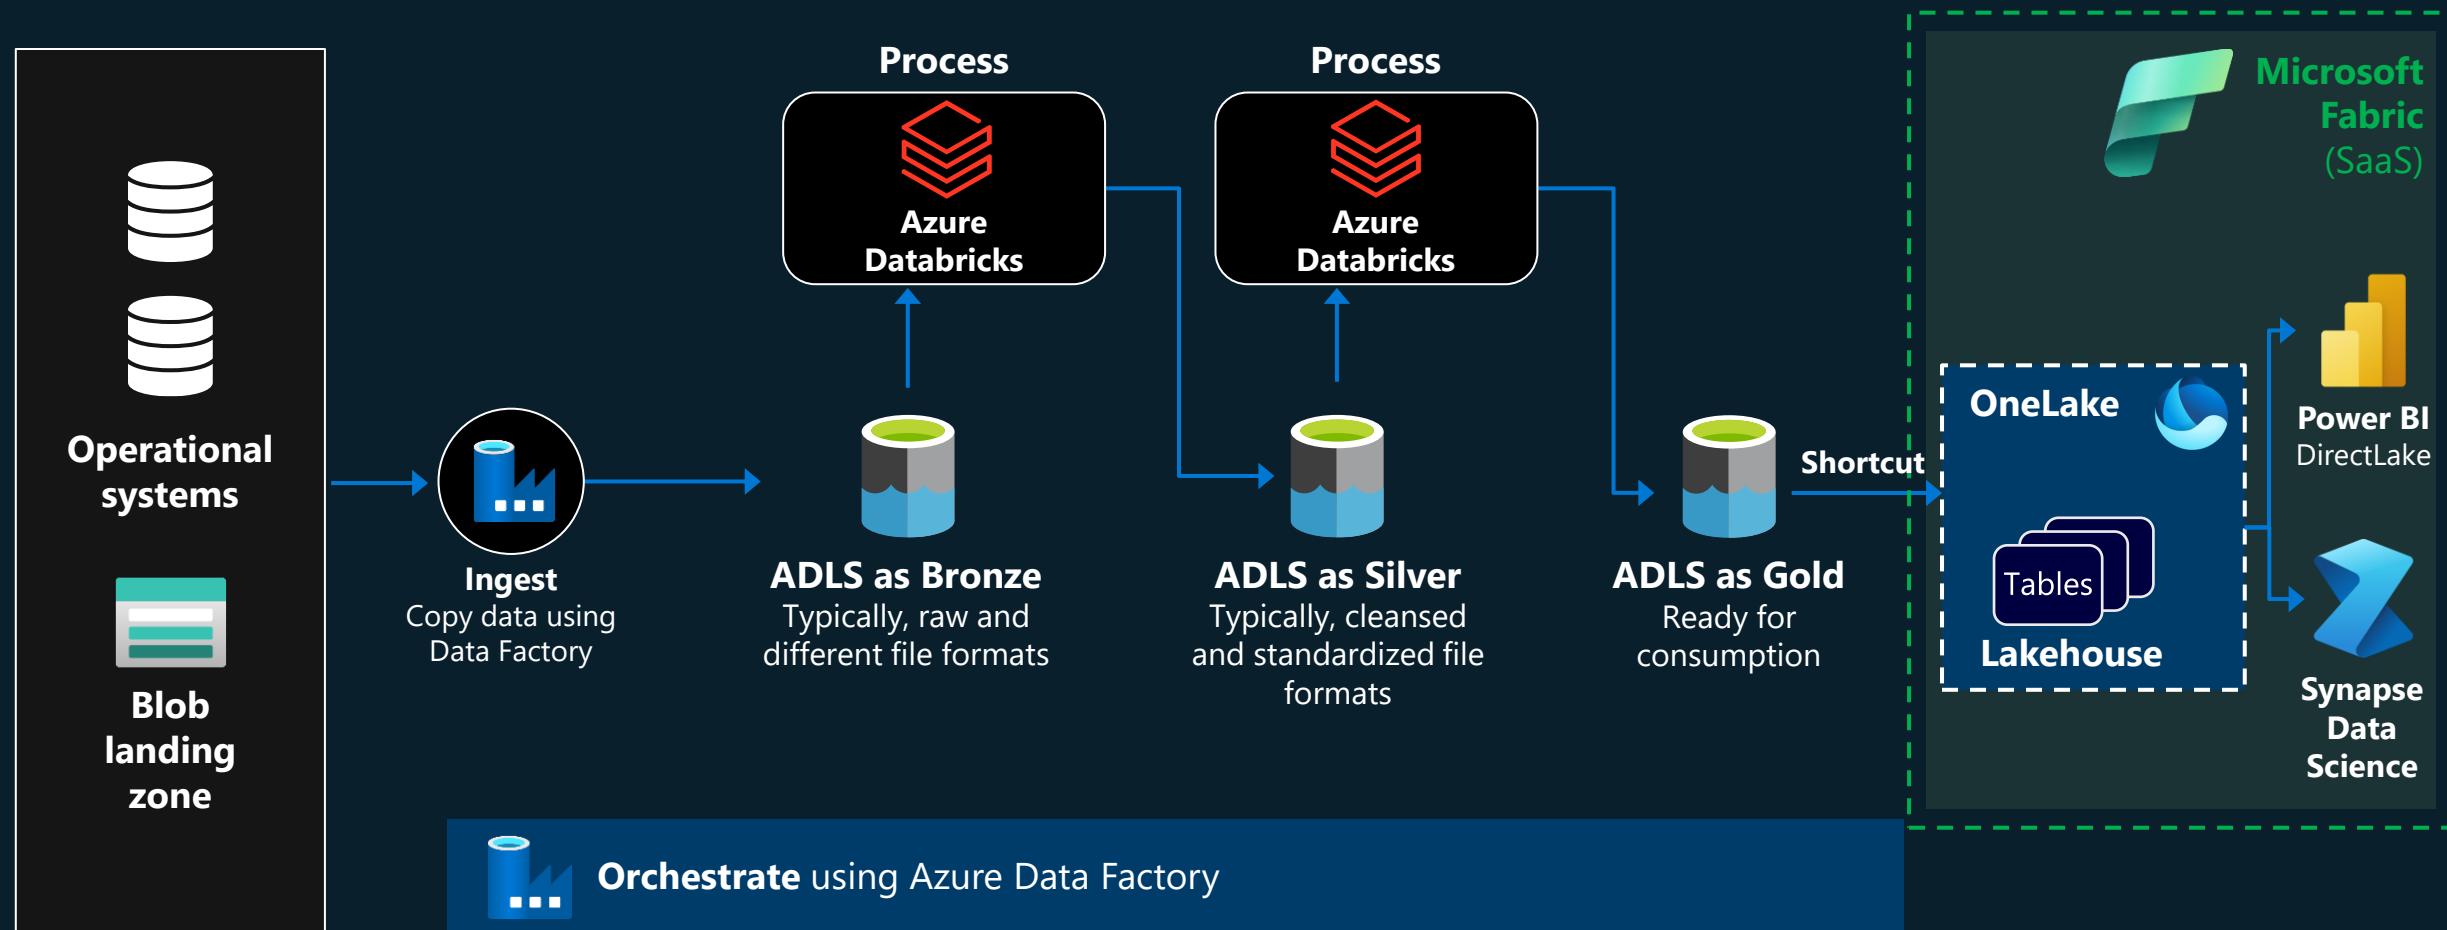
Azure Databricks can directly read/write data from any Fabric data item stored in OneLake through external link.

Existing non-Unity Catalog managed data in ADLS gen2 storage can be a shortcut into OneLake.

Works with all Fabric engines including Power BI with direct lake mode.

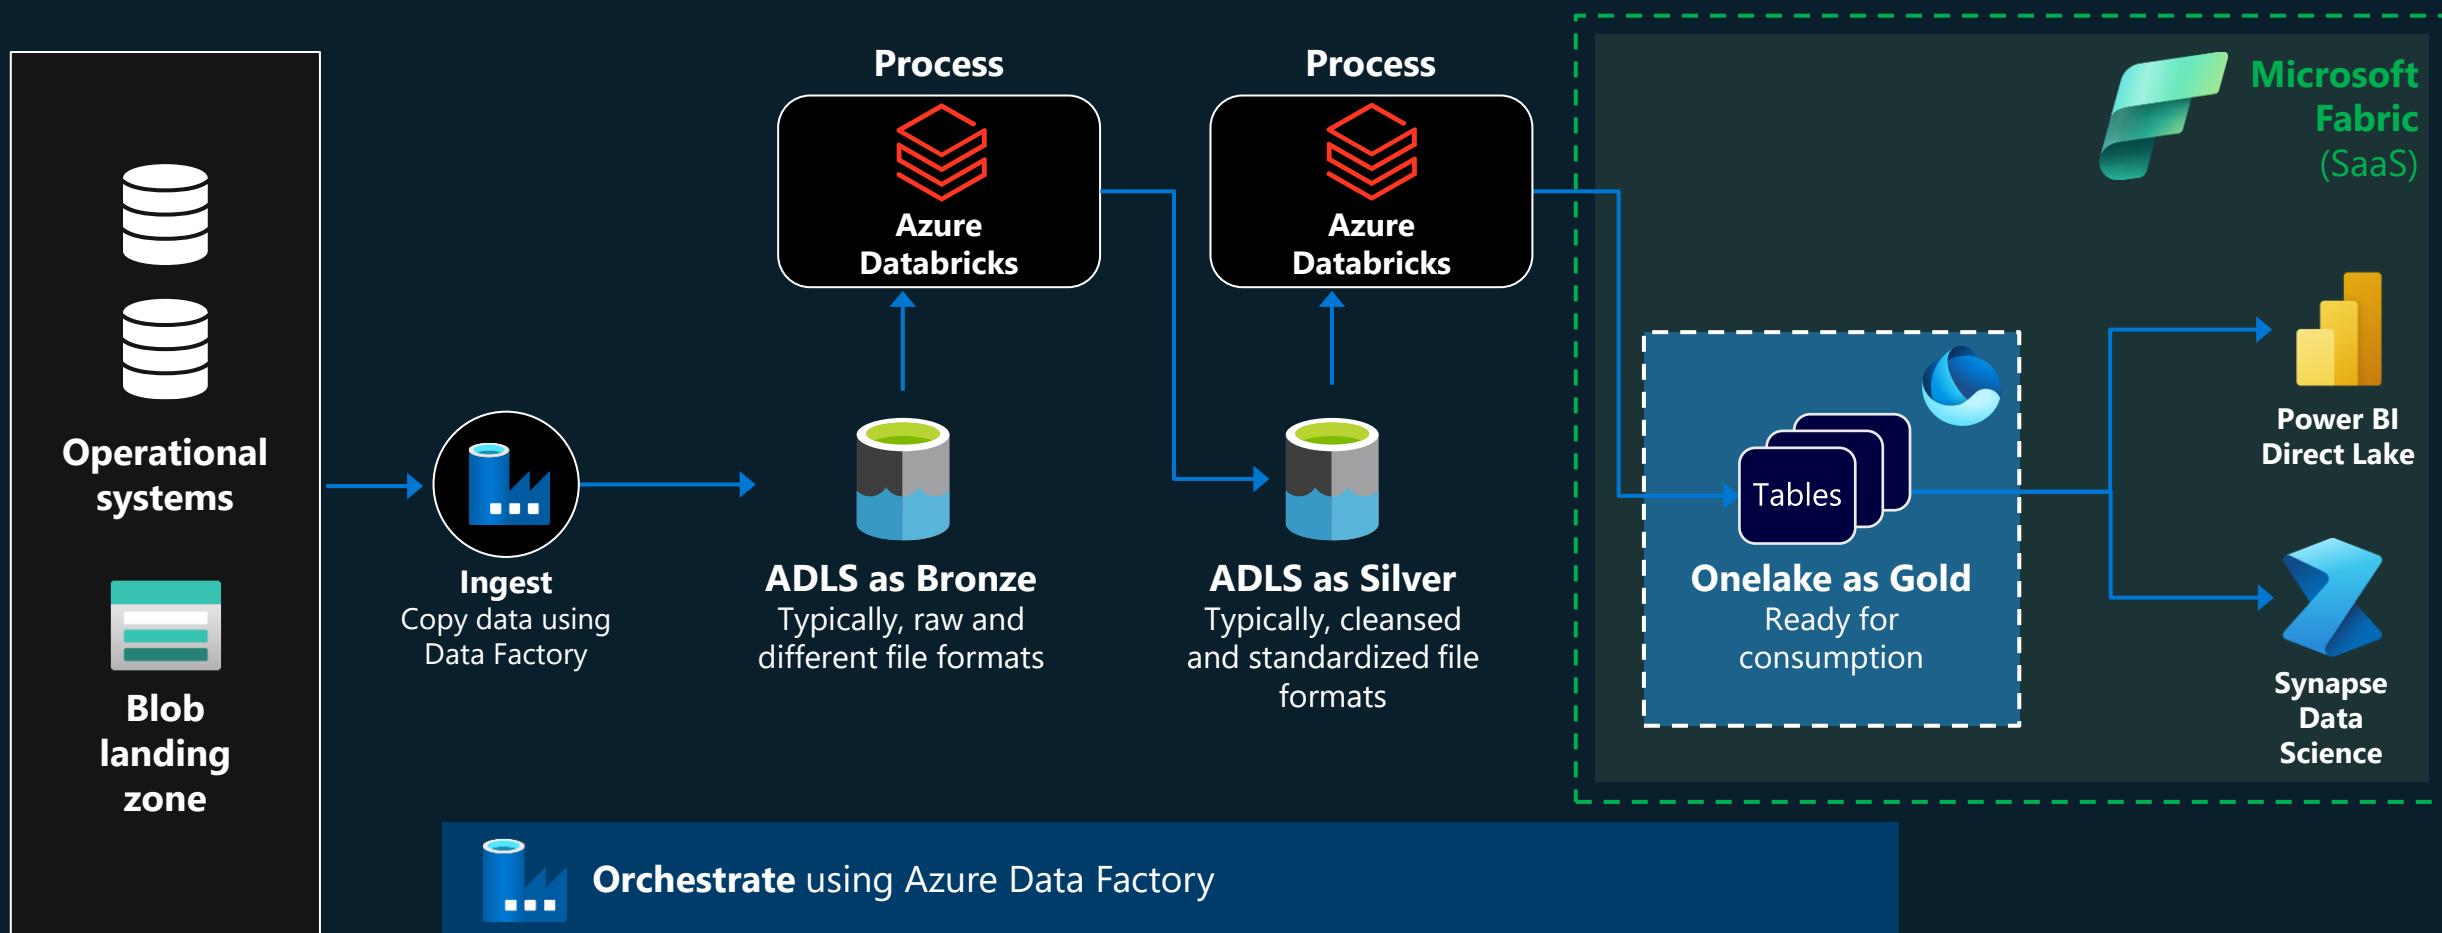
Lakehouse architecture – ADLS as Gold Layer

<https://learn.microsoft.com/en-us/fabric/onelake/create-adls-shortcut>



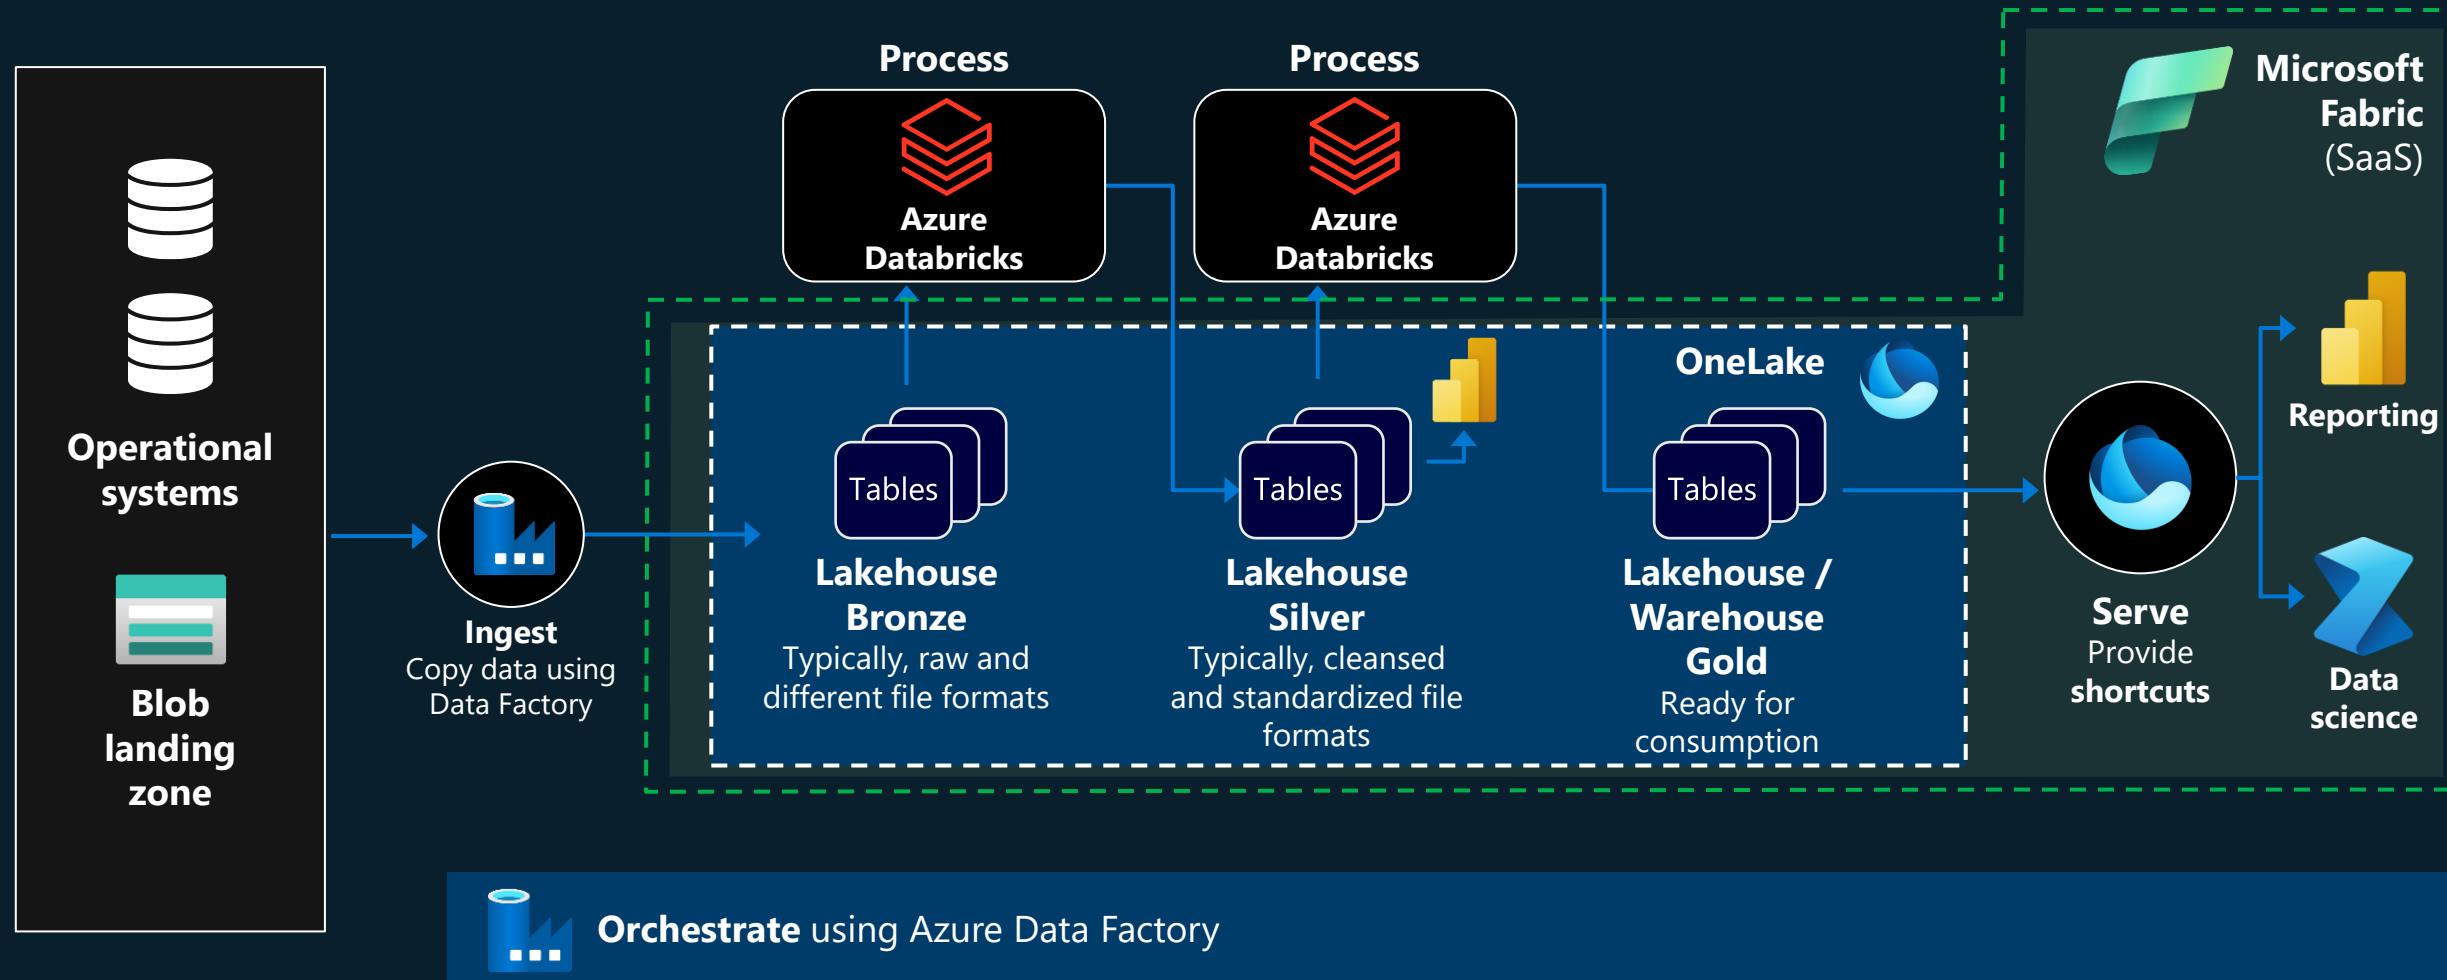
Lakehouse architecture – OneLake as Gold Layer

<https://learn.microsoft.com/en-us/fabric/onelake/onelake-azure-databricks>



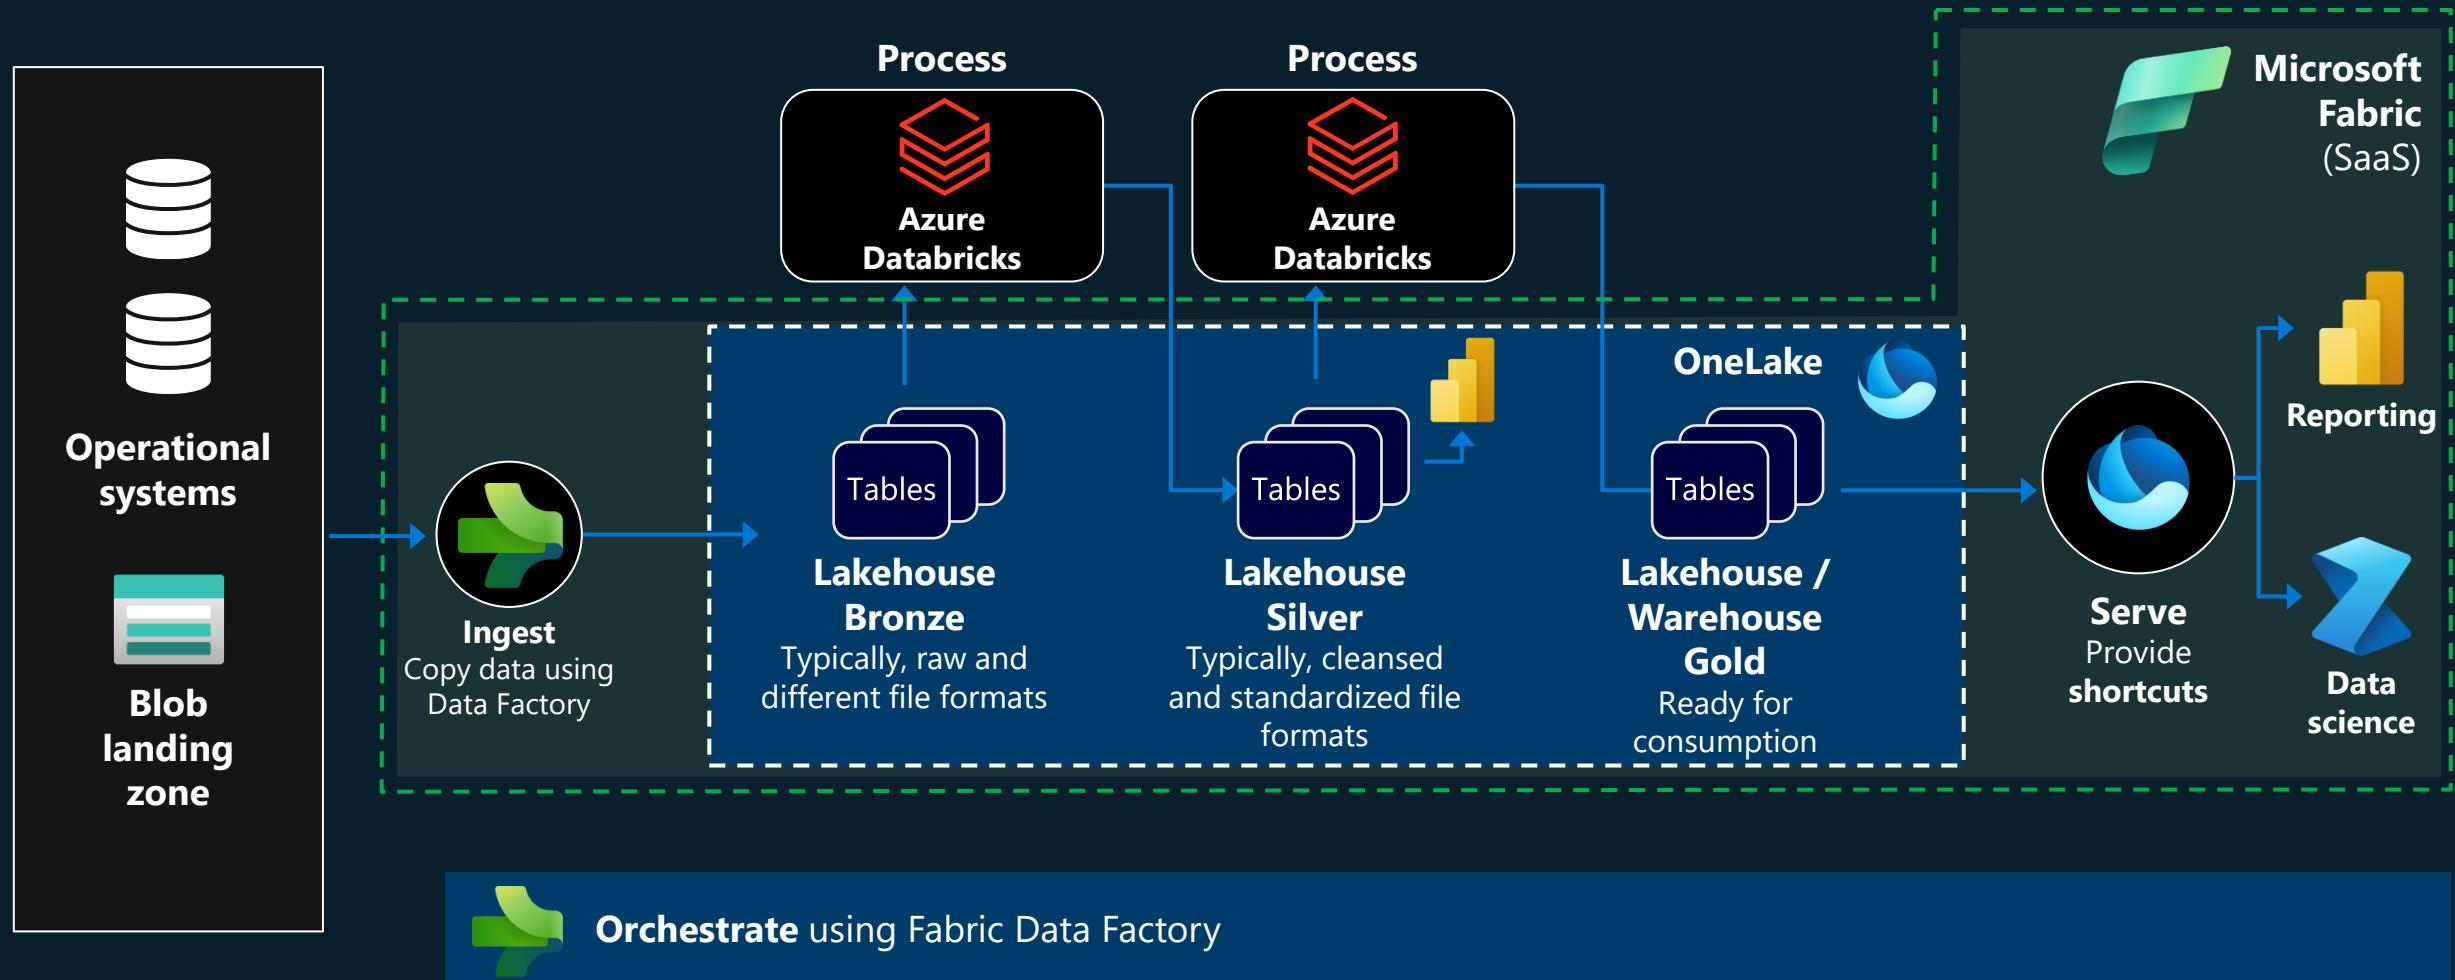
Store all Delta tables in OneLake

<https://learn.microsoft.com/en-us/fabric/onelake/onelake-azure-databricks>



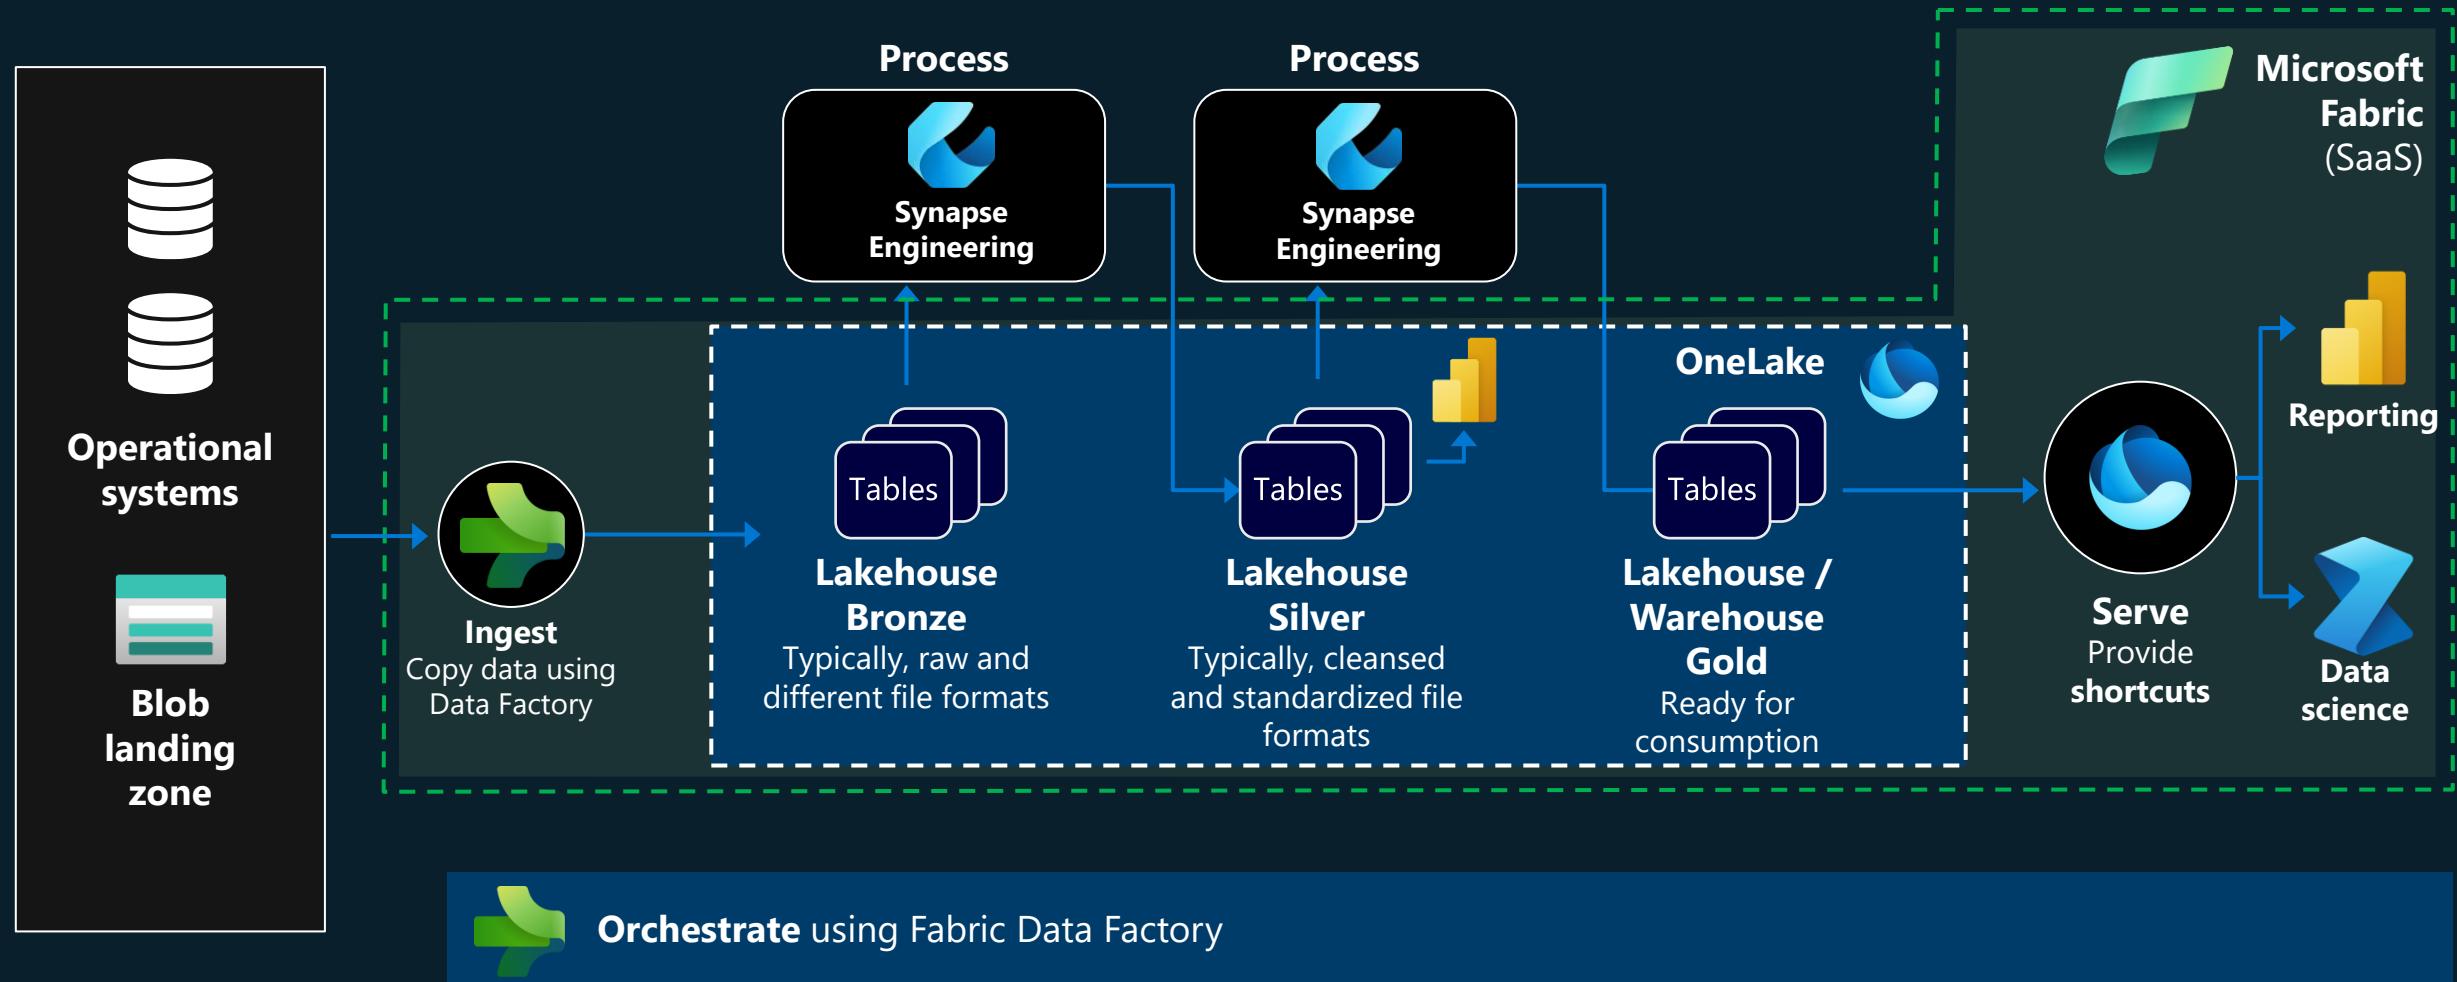
Store all Delta tables in OneLake – Orchestrate with FDF

<https://learn.microsoft.com/en-us/fabric/onelake/onelake-azure-databricks>



Store all Delta tables in OneLake

<https://learn.microsoft.com/en-us/fabric/onelake/onelake-azure-databricks>



What's Next for Azure Synapse Analytics?

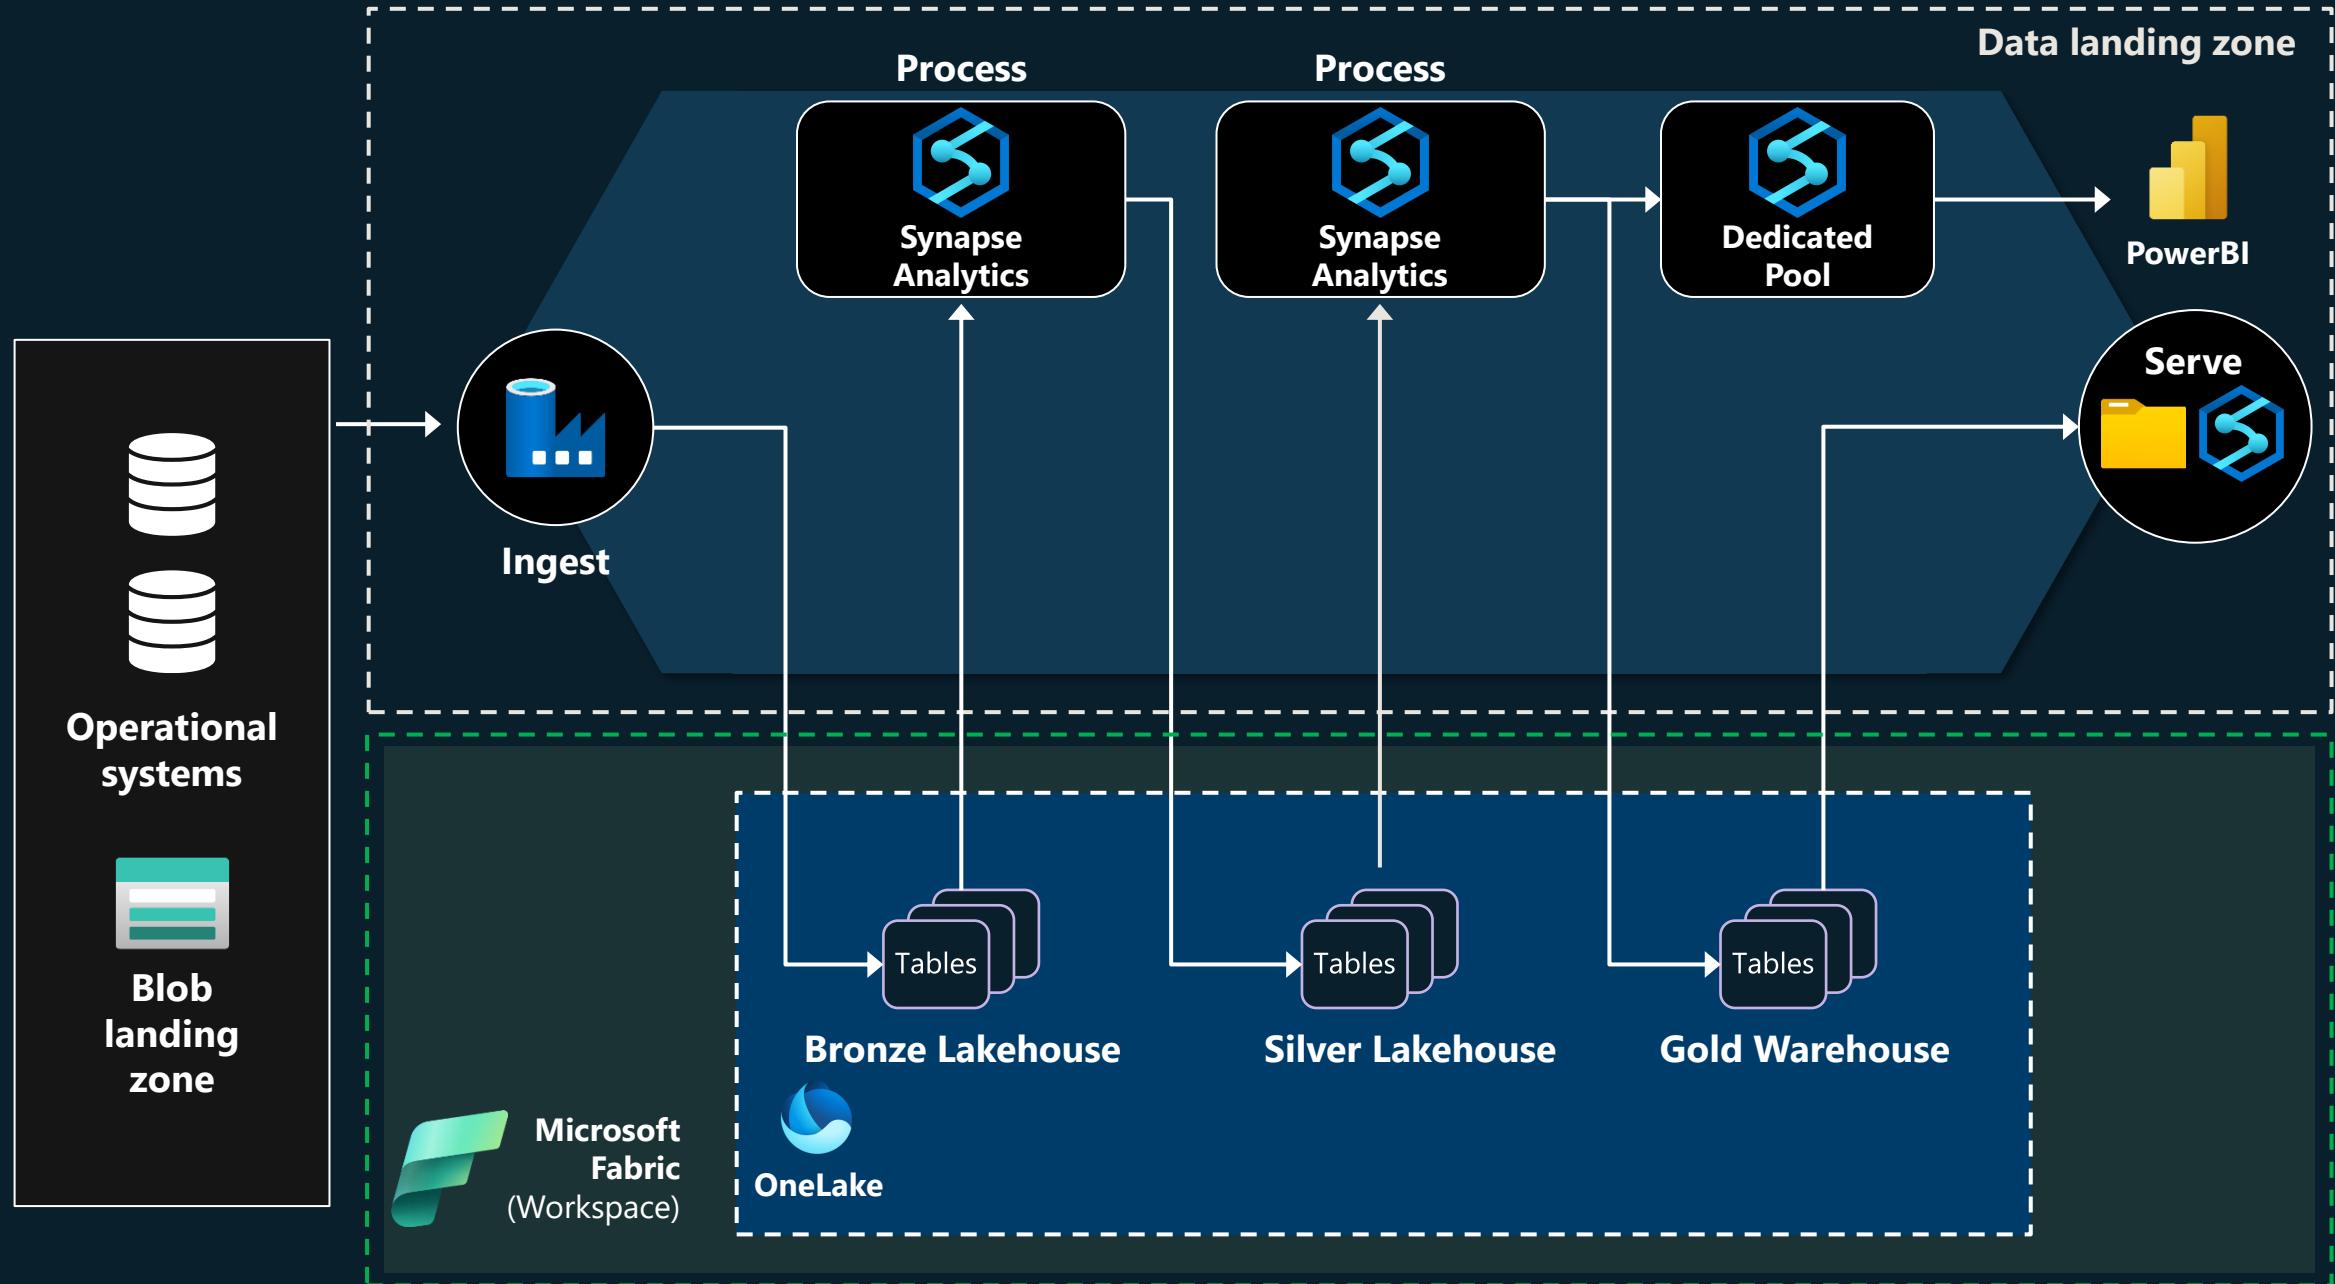


Microsoft Fabric

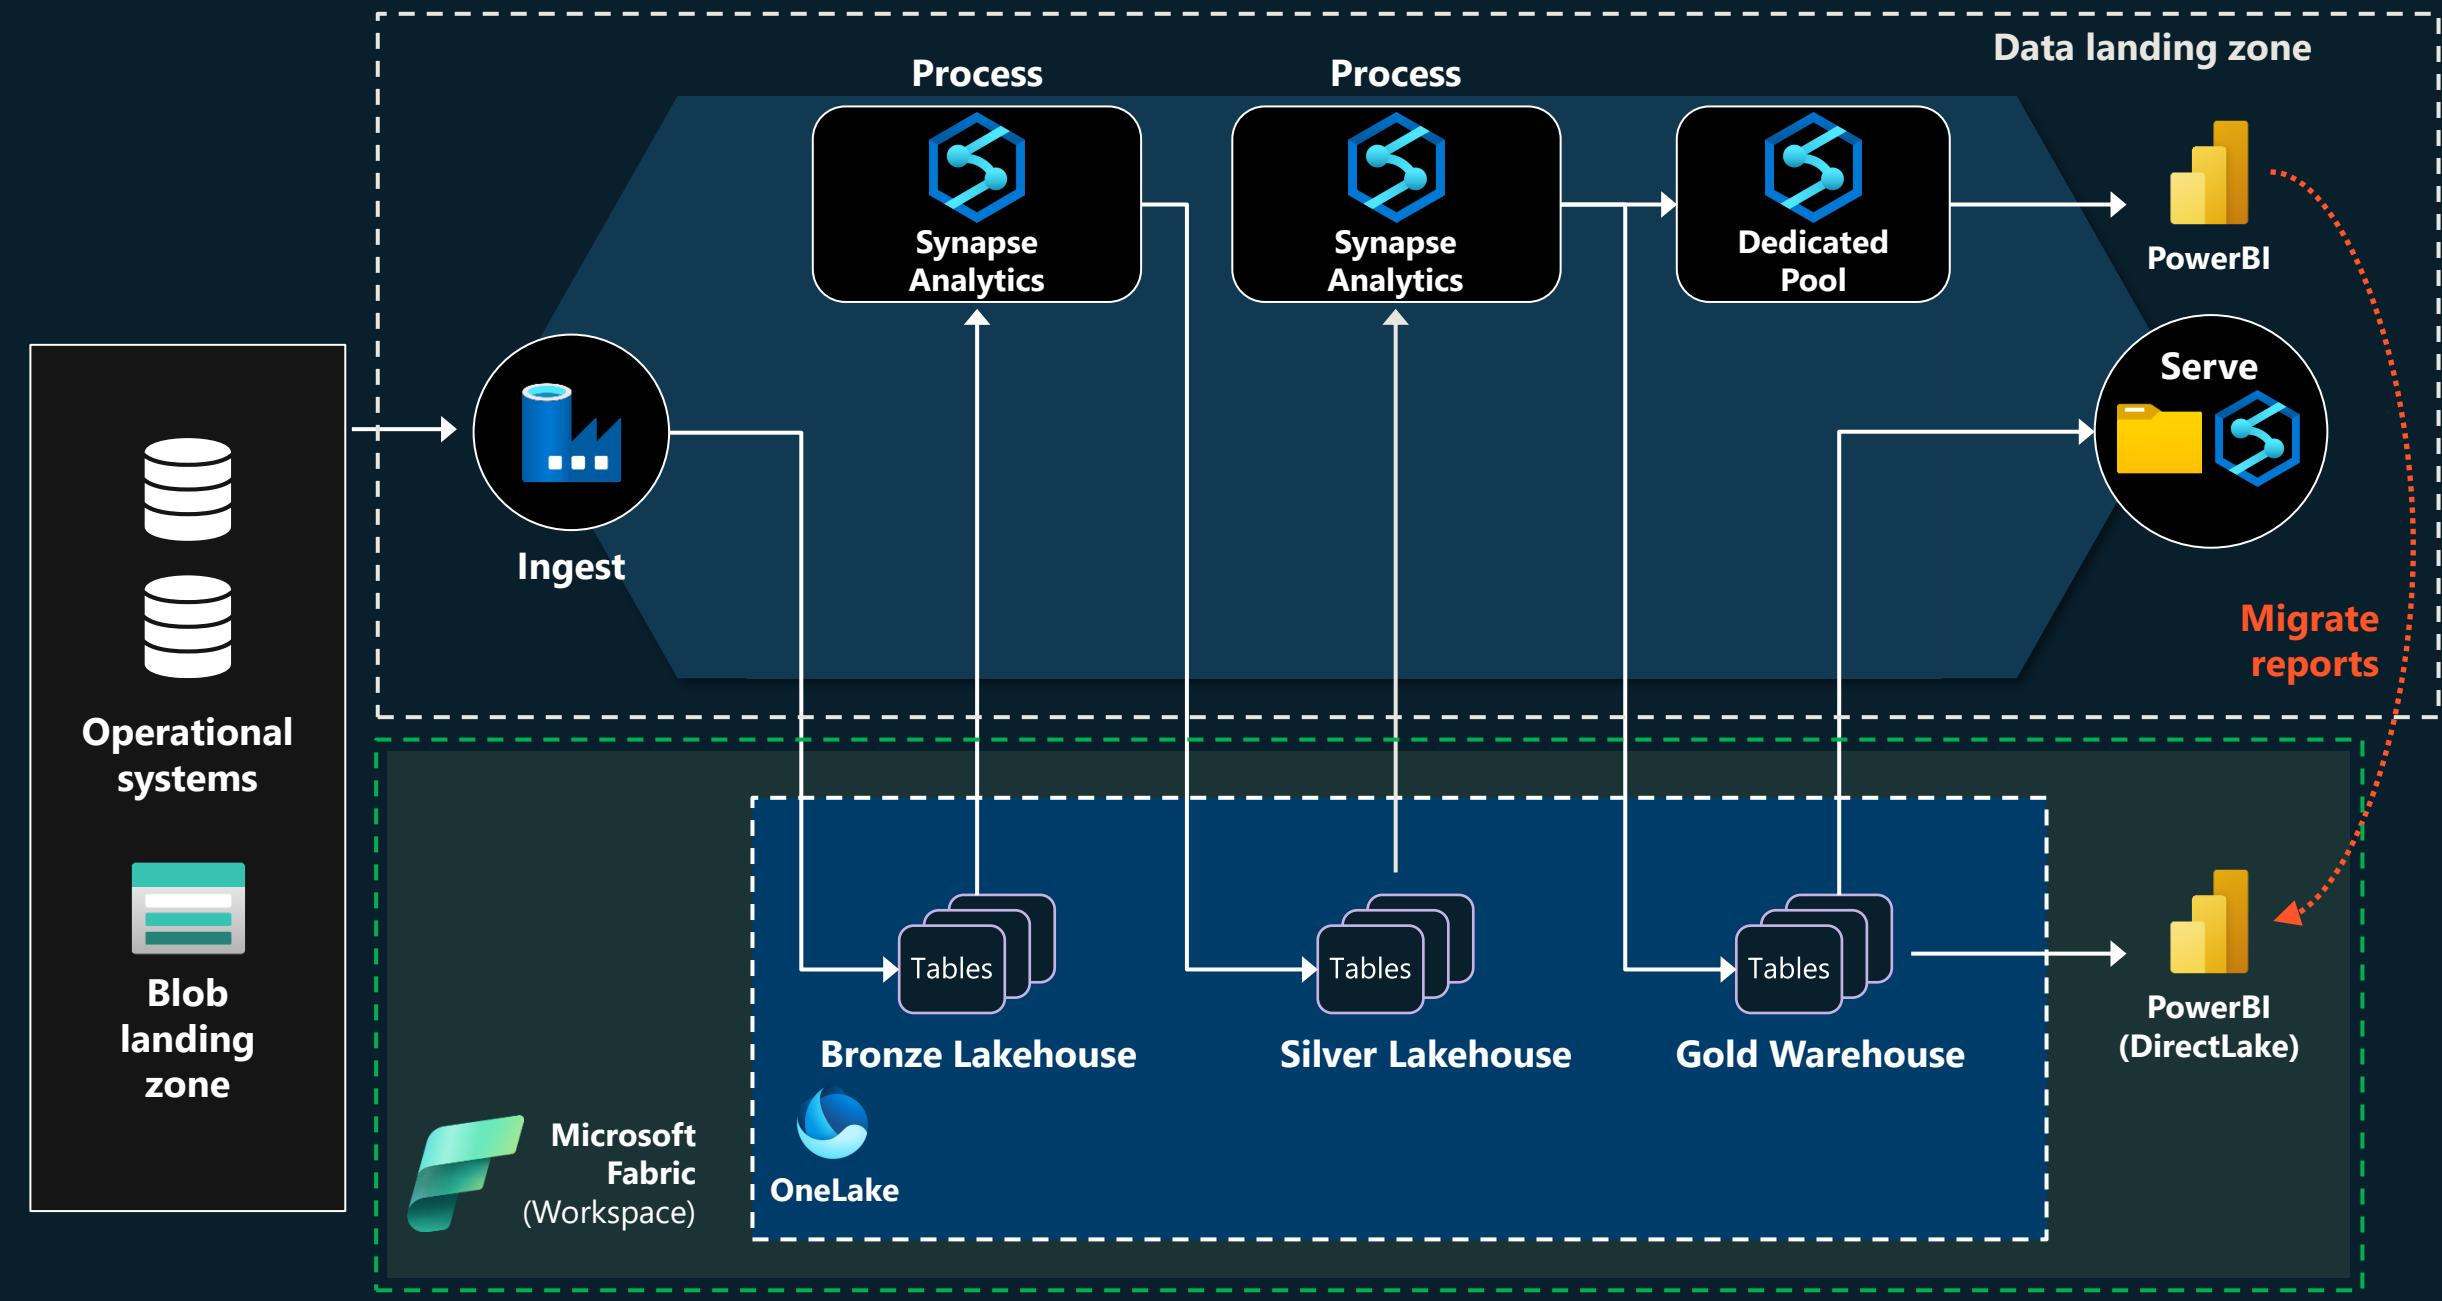


Azure Synapse Analytics

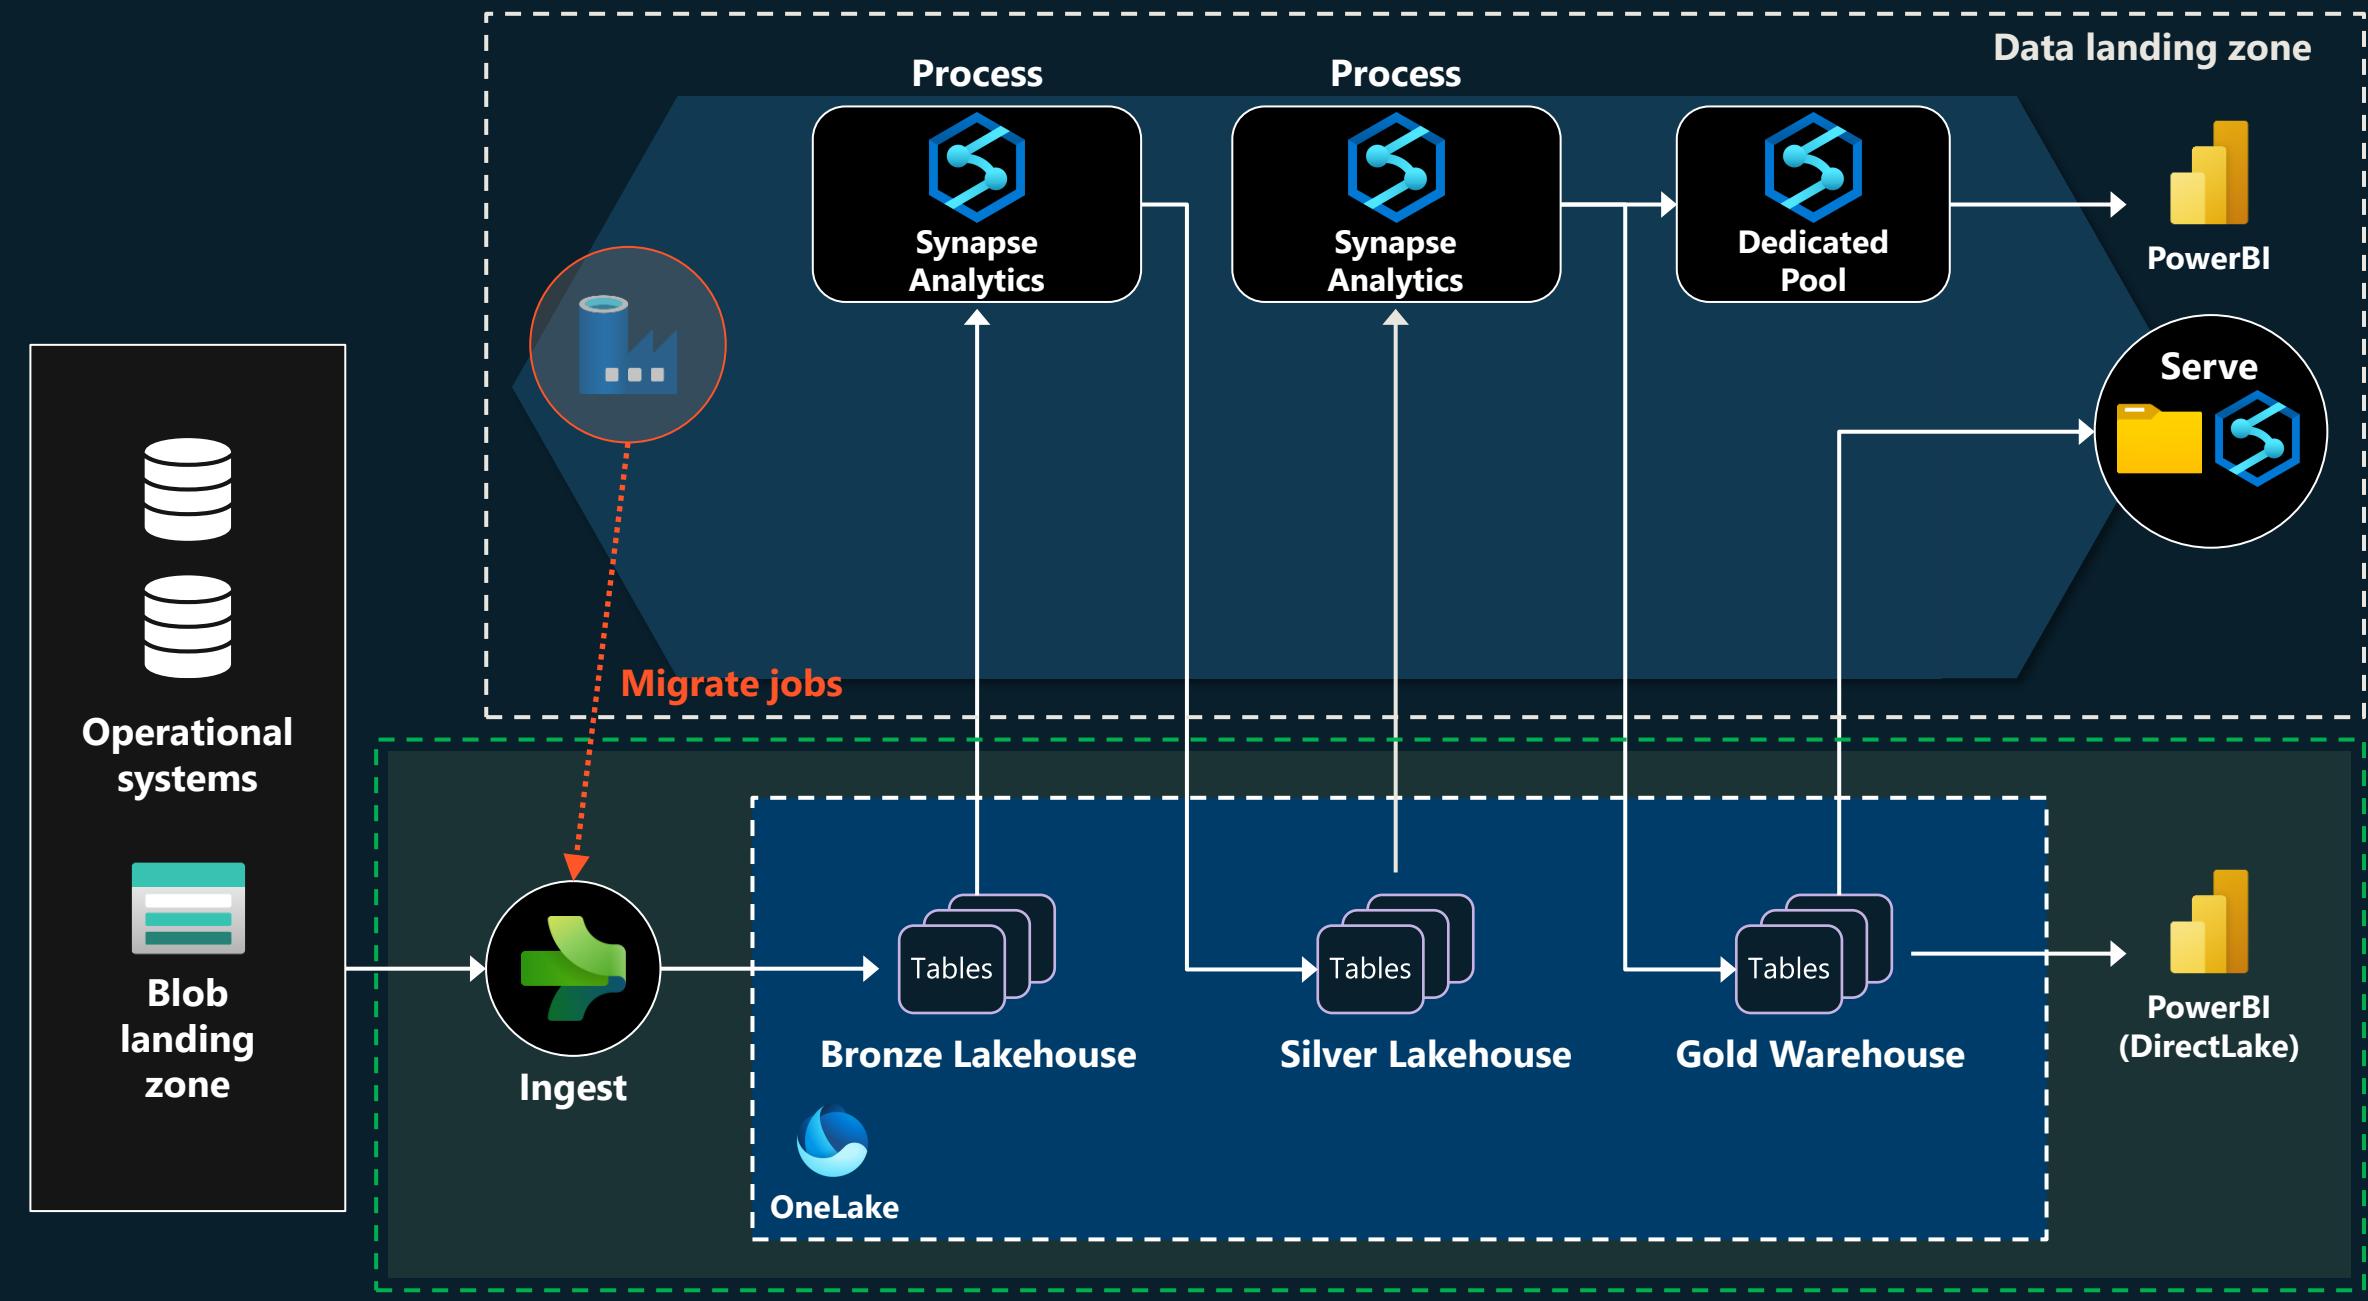
Step 1: Integrate Synapse Analytics with OneLake (migrate data)



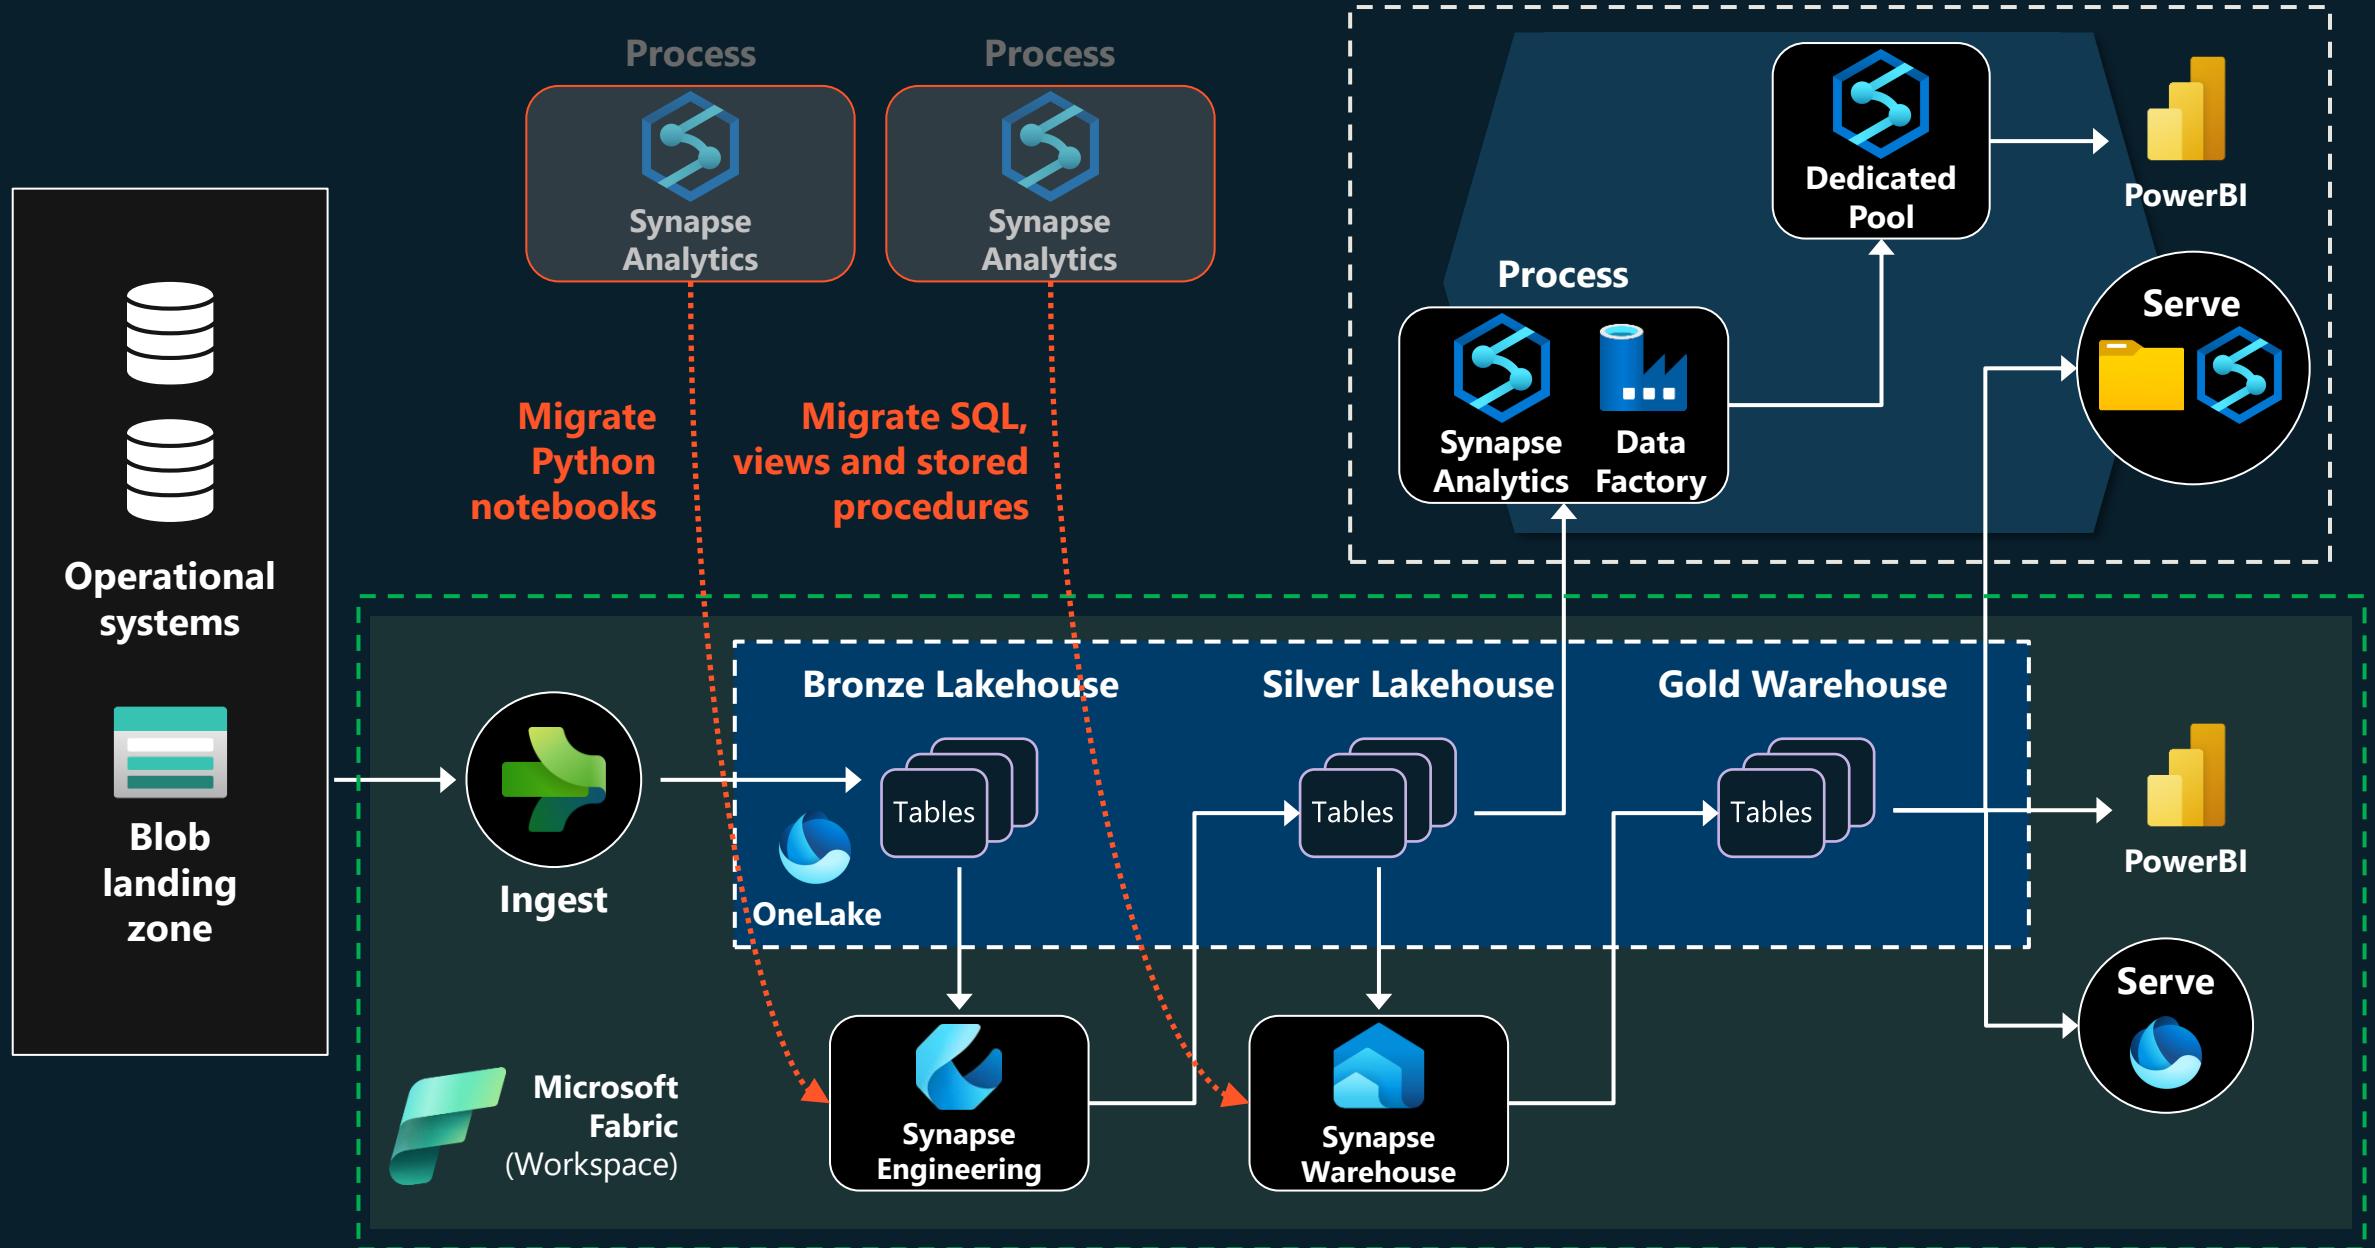
Step 2: Migrate first set of PowerBI reports



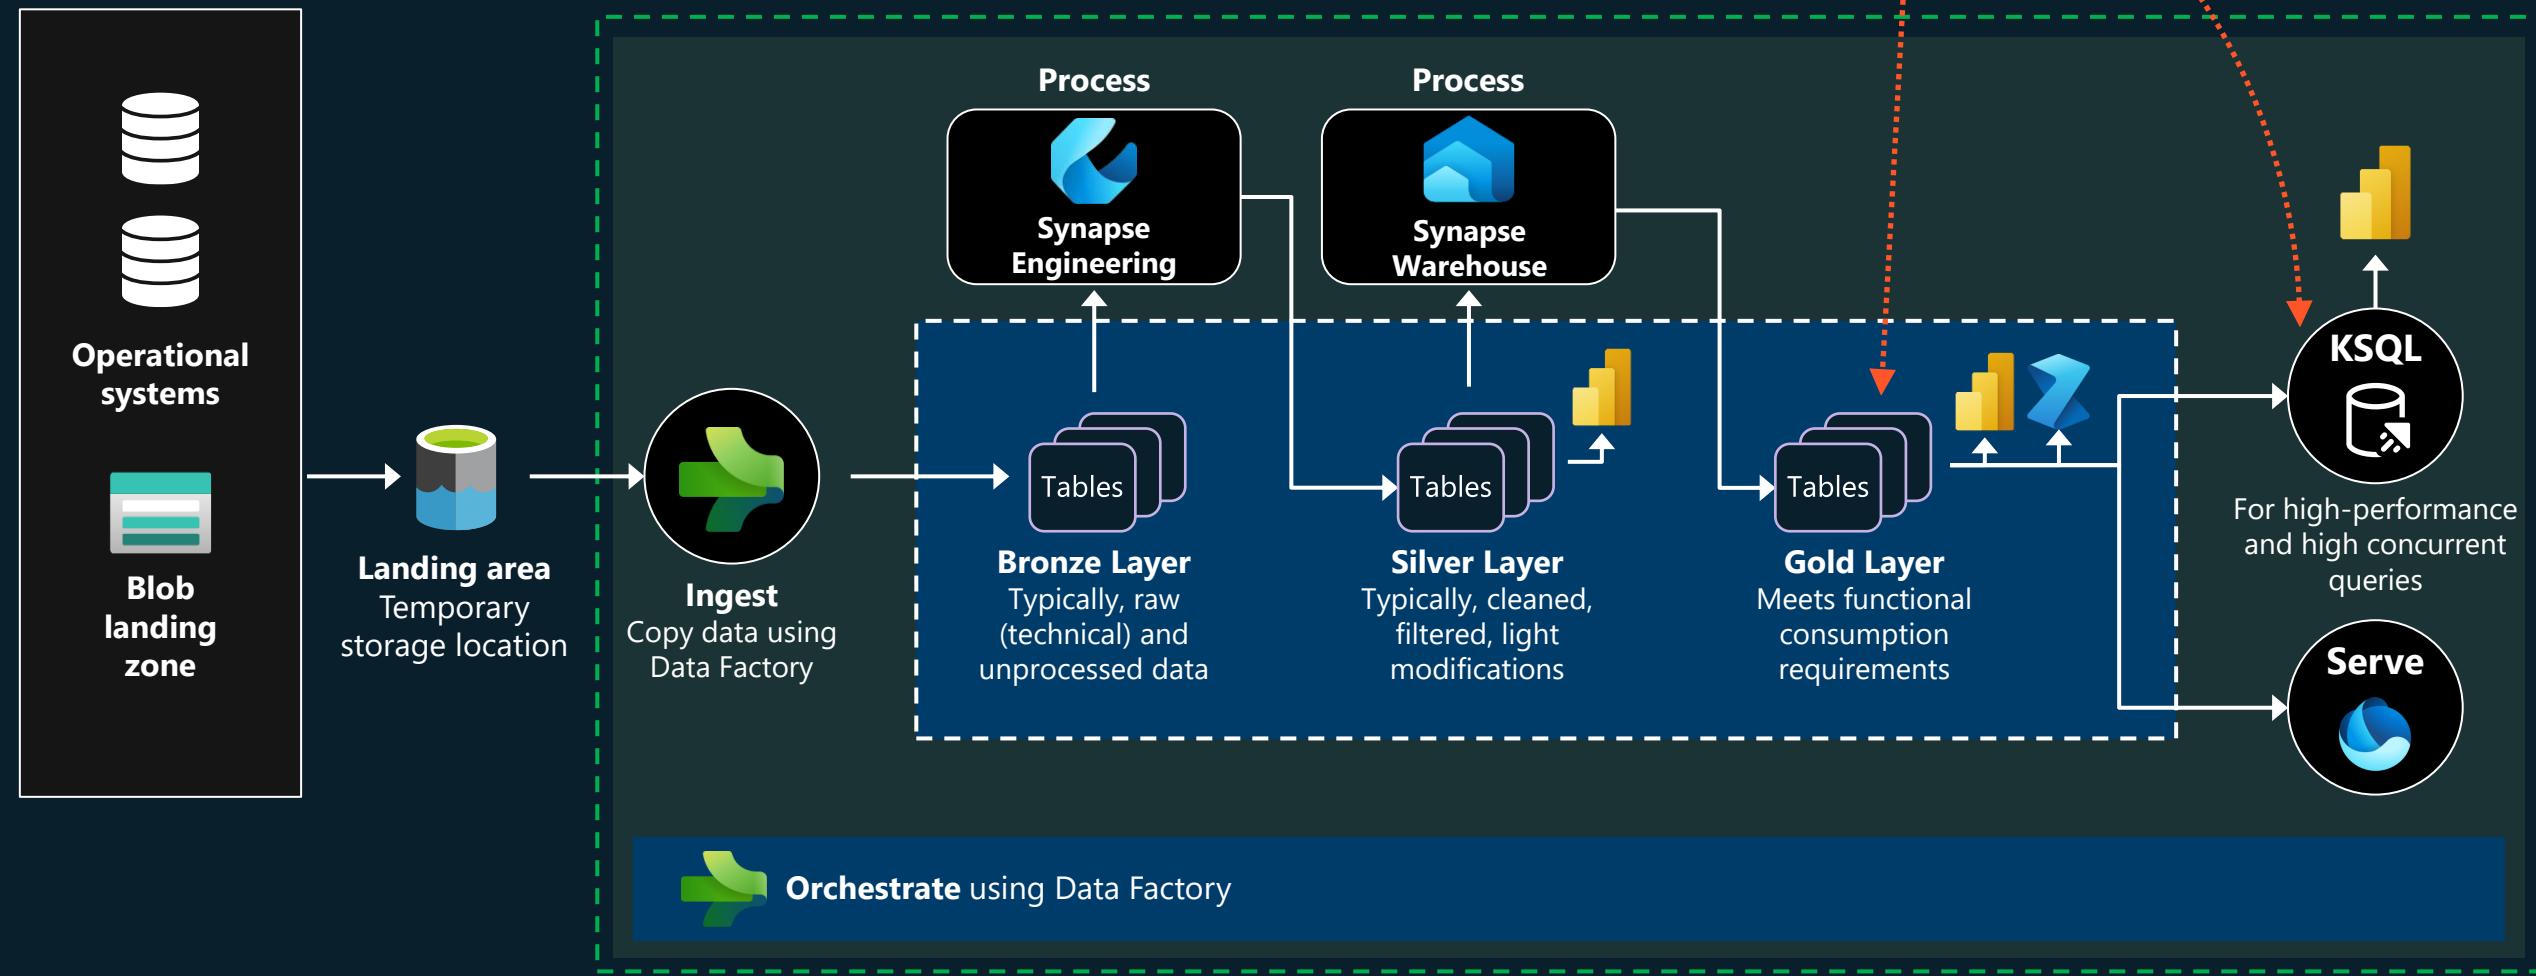
Step 3: Ingest (and some orchestration) with Fabric Data Factory



Step 4: Migrate Spark processing to Fabric



Step 5: Migrate Dedicated Pool workloads to Fabric



Considerations for migrating from Synapse

- Don't implement any new dedicated pools.
 - Migration scenarios and tooling will come. Migration will be done by attaching your existing Synapse workspace. All dedicated pool data then is replicated OneLake. Both will be kept in sync. Reporting and analytics can be moved over. You can scale down the DW units because Fabric is used. Move writes over to Fabric. Finally, redirect all queries.
- Check the Azure Data Factory/Synapse Pipeline feature parity. Mapping Dataflows (ADF/Synapse) is replaced with Data Flows Gen2 in Fabric. Migration scenarios will come (e.g., mounting, tooling).
- Serverless SQL: Works, however, Fabric does not support OPENROWSET. Instead use SQL Endpoints of Fabric Lakehouse (only Delta support).
- On Prem Data Gateways now support PBI and Data Factory. Different gateway options are now possible (on-prem data gateway, Vnet data gateway).

Explore the Fabric Roadmap

Rapid pace of innovation

The image displays three side-by-side screenshots of Microsoft Fabric blog posts, each featuring a large green circular overlay containing the text "Jul 2023", "Aug 2023", and "Sep 2023" respectively. The first screenshot shows the "Microsoft Fabric July 2023 Update" post, the second shows the "Microsoft Fabric August 2023 update", and the third shows the "Microsoft Fabric September 2023 update". Each post includes a video thumbnail titled "Fabric Monthly Update - July/August/September 2023" and a "Watch on YouTube" button.

aka.ms/FabricRoadmap

July 2023

Microsoft Fabric Blog

Blog > Announcements > Microsoft Fabric July 2023 Update

Microsoft Fabric July 2023 Update

Announcements Apache Spark Core Data Activator Data Factory Data Science Data Warehouse Monthly Update OneLake Real-time Analytics Synapse

August 9, 2023 by Ryan Majidimehr

Welcome to the July 2023 update. We have features in Core, Synapse, Data Factory, Data Warehouse, and Power BI.

Fabric Monthly Update - July 2023

Watch later Share

Watch on YouTube

Contents

Microsoft Fabric Blog

Blog > June 2023 > Microsoft Fabric August 2023 update

Microsoft Fabric August 2023 update

Announcements Apache Spark Core Data Activator Data Factory Data Science Data Warehouse Monthly Update OneLake Real-time Analytics Synapse

August 9, 2023 by Ryan Majidimehr

Welcome to the August 2023 update. We have lots of exciting new features including the recently announced switcher for Power BI, SSD caching in Synapse Data Warehouse, in-line Python support for KQL in Synapse Real-time Analytics, lookup activity for Data Factory Dataflows, and much more. Continue reading for more details on our new features!

Fabric Monthly Update - August 2023

Watch later Share

Watch on YouTube

Contents

Microsoft Fabric Blog

Blog > All > Microsoft Fabric September 2023 update

Microsoft Fabric September 2023 update

Announcements Apache Spark Community Core Data Factory Data Science Data Warehouse Monthly Update OneLake Real-time Analytics Synapse

September 12, 2023 by Ryan Majidimehr

Welcome to the September 2023 update.

We have lots of features this month including updates to the monitoring hub, Fabric Metrics app, VS code integration for Data Engineering, Real-time data sharing and many more. Continue reading for more details on our new features!

Contents

- Core
 - Monitoring hub – column options
- OneLake
 - OneLake file explorer (v.1.0.10) – Menu Option to View Workspaces and Items Online
- Power BI
 - General
 - Upgrade Power BI Desktop .NET Framework to 4.7.2 or newer
 - Reporting



Unlock more value from your data with Microsoft Fabric

Q/A

Discovery Days – April 17th 2024